

SMOKE TEST DOCUMENT

Confluent Kafka Smoke Test Cases

Date Prepared: July 2019

Document Information

Project Name	Confluent Kafka Smoke Test Document		
Project Owner		Document Version No	1.0
Quality Review Method	By email/HP SharePoint		
Prepared By		Preparation Date	July 2019
Reviewed By	Refer to version history	Review Date	

Table of Contents

1	PREREQUISITES	4
2	LAUNCH THE KSQL CLI	5
3	INSPECT KAFKA TOPICS BY USING SHOW AND PRINT STATEMENTS.....	6
4	CREATE A STREAM AND TABLE	8
5	WRITE QUERIES:	10
6	WORD COUNT LAMBDA INTEGRATION TEST:	15
7	KSQL REST API TESTING	16
8	TESTING CONFLUENT KAFKA BROKER AND KAFKA ZOOKEEPER SERVICES.....	17
9	TESTING SCHEMA REGISTRY	19
10	MONITOR AND MANAGE – CONTROL CENTER	22
10.1.1	System Health.....	22
10.1.2	Kafka Connect	26
10.1.3	Clusters	27
10.1.4	Topics	29
10.1.5	Alerts	31
11	TESTING KAFKA CONNECT	33
12	CONFIGURING CONFLUENT KAFKA WITH HDFS AND DTAP USING CONNECTORS	38

1 PREREQUISITES

- This installation includes a Kafka broker, KSQL, Control Center, Zookeeper, Schema Registry, REST Proxy, and Kafka Connect.
- If you installed Confluent Platform via TAR or ZIP, navigate into the installation directory. The paths and commands used throughout this tutorial assume that you are in this installation directory.
- Java: Minimum version 1.8. Install Oracle Java JRE or JDK ≥ 1.8 on your local machine.

Create and produce data to the Kafka topics `pageviews` and `users`. These steps use the KSQL datagen that is included Confluent Platform.

- Create the `pageviews` topic and produce data using the data generator. The following example continuously generates data with a value in DELIMITED format.
 - `$ <path-to-confluent>/bin/ksql-datagen quickstart=pageviews format=delimited topic=pageviews maxInterval=500`
- Produce Kafka data to the user's topic using the data generator. The following example continuously generates data with a value in JSON format.
 - `$ <path-to-confluent>/bin/ksql-datagen quickstart=users format=json topic=users maxInterval=100`

2 LAUNCH THE KSQL CLI

After KSQL is started, your terminal should resemble this.

```
=====
=
=      _ _ _ _ _
=      | | / / __ | / _ \ |
=      | ' / | ( _ | | | |
=      | < \_ \ | | | |
=      | . \ __ ) | | _ | | __
=      | _ | \ \ __ / \_ \ \ __ |
=
=
= Streaming SQL Engine for Apache Kafka® =
=====
```

Copyright 2018 Confluent Inc.

CLI v5.1.1, Server v5.1.1 located at <http://localhost:8088>

Having trouble? Type 'help' (case-insensitive) for a rundown of how things work!

ksql>

3 INSPECT KAFKA TOPICS BY USING SHOW AND PRINT STATEMENTS

KSQL enables inspecting Kafka topics and messages in real time.

KSQL enables inspecting Kafka topics and messages in real time.

- Use the SHOW TOPICS statement to list the available topics in the Kafka cluster.
- Use the PRINT statement to see a topic's messages as they arrive.

In the KSQL CLI, run the following statement:

- SHOW TOPICS;

Output:

Kafka Topic	Registered	Partitions	Partition Replicas	Consumers	Consumer Groups
-------------	------------	------------	--------------------	-----------	-----------------

_confluent-metrics	false	12	1	0	0
_schemas	false	1	1	0	0
pageviews	false	1	1	0	0
users	false	1	1	0	0

- Inspect the users topic by using the PRINT statement:

PRINT 'users';

Output:

Format:JSON

```
{ "ROWTIME":1540254230041,"ROWKEY":"User_1","registertime":1516754966866,"userid":"
User_1","regionid":"Region_9","gender":"MALE"}
{"ROWTIME":1540254230081,"ROWKEY":"User_3","registertime":1491558386780,"userid":"
User_3","regionid":"Region_2","gender":"MALE"}
{"ROWTIME":1540254230091,"ROWKEY":"User_7","registertime":1514374073235,"userid":"
User_7","regionid":"Region_2","gender":"OTHER"}
^C{"ROWTIME":1540254232442,"ROWKEY":"User_4","registertime":1510034151376,"userid"
```

```
":"User_4","regionid":"Region_8","gender":"FEMALE"}  
Topic printing ceased
```

- Inspect the pageviews topic by using the PRINT statement:


```
PRINT 'pageviews';
```

Output:

```
Format:STRING  
10/23/18 12:24:03 AM UTC , 9461 , 1540254243183,User_9,Page_20  
10/23/18 12:24:03 AM UTC , 9471 , 1540254243617,User_7,Page_47  
10/23/18 12:24:03 AM UTC , 9481 , 1540254243888,User_4,Page_27  
^C10/23/18 12:24:05 AM UTC , 9521 , 1540254245161,User_9,Page_62  
Topic printing ceased  
ksql>
```

4 CREATE A STREAM AND TABLE

These examples query messages from Kafka topics called pageviews and users, using the following schemas:



- Create a stream pageviews_original from the Kafka topic pageviews, specifying the value_format of DELIMITED.
- Describe the new STREAM. Notice that KSQL created additional columns called ROWTIME, which corresponds to the Kafka message timestamp, and ROWKEY, which corresponds to the Kafka message key.
 - CREATE STREAM pageviews_original (viewtime bigint, userid varchar, pageid varchar) WITH \ (kafka_topic='pageviews', value_format='DELIMITED');

Output:

Message

Stream created

- Create a table users_original from the Kafka topic users, specifying the value_format of JSON.
 - CREATE TABLE users_original (registertime BIGINT, gender VARCHAR, regionid VARCHAR, userid VARCHAR) WITH \ (kafka_topic='users', value_format='JSON', key = 'userid');

Output:

Message

Table created

- Show all streams and tables.

```
ksql> SHOW STREAMS;
```

Stream Name	Kafka Topic	Format
-------------	-------------	--------

PAGEVIEWS_ORIGINAL	pageviews	DELIMITED
--------------------	-----------	-----------

```
ksql> SHOW TABLES;
```

Table Name	Kafka Topic	Format	Windowed
------------	-------------	--------	----------

USERS_ORIGINAL	users	JSON	false
----------------	-------	------	-------

5 WRITE QUERIES:

Below are example queries using KSQL.

Note: By default KSQL reads the topics for streams and tables from the latest offset.

- Use SELECTS to create a query that returns data from a STREAM. This query includes the LIMIT keyword to limit the number of rows returned in the query result. Note that exact data output may vary because of the randomness of the data generation.

- SELECT pageid FROM pageviews_original LIMIT 3;

Output:

```
Page_24
Page_73
Page_78
LIMIT reached
Query terminated
```

- Create a persistent query by using the CREATE STREAM keywords to precede the SELECT statement. The continual results from this query are written to the PAGEVIEWS_ENRICHED Kafka topic. The following query enriches the pageviews STREAM by doing a LEFT JOIN with the users_original TABLE on the user ID.

- CREATE STREAM pageviews_enriched AS \
 - SELECT users_original.userid AS userid, pageid, regionid, gender \
 - FROM pageviews_original \
 - LEFT JOIN users_original \
 - ON pageviews_original.userid = users_original.userid;

Output:

```
Message
-----
Stream created and running
-----
```

- Use SELECT to view query results as they come in. To stop viewing the query results, press <ctrl-c>. This stops printing to the console but it does not terminate the actual query. The query continues to run in the underlying KSQL application.

- `SELECT * FROM pageviews_enriched;`

Output:

```
1519746861328 | User_4 | User_4 | Page_58 | Region_5 | OTHER
1519746861794 | User_9 | User_9 | Page_94 | Region_9 | MALE
1519746862164 | User_1 | User_1 | Page_90 | Region_7 | FEMALE
^CQuery terminated
```

- Create a new persistent query where a condition limits the streams content, using WHERE.

Results from this query are written to a Kafka topic called PAGEVIEWS_FEMALE.

- `CREATE STREAM pageviews_female AS \`
`SELECT * FROM pageviews_enriched \`
`WHERE gender = 'FEMALE';`

Output:

Message

Stream created and running

- Create a new persistent query where another condition is met, using LIKE. Results from this query are written to the pageviews_enriched_r8_r9 Kafka topic.

- `CREATE STREAM pageviews_female_like_89 \`
`WITH (kafka_topic='pageviews_enriched_r8_r9') AS \`
`SELECT * FROM pageviews_female \`
`WHERE regionid LIKE '%_8' OR regionid LIKE '%_9';`

Output:

Message

Stream created and running

- Create a new persistent query that counts the pageviews for each region and gender combination in a tumbling window of 30 seconds when the count is greater than one. Results from this query are written to the PAGEVIEWS_REGIONS Kafka topic in the Avro format. KSQL will register the Avro schema with the configured Schema Registry when it writes the first message to the PAGEVIEWS_REGIONS topic.

- `CREATE TABLE pageviews_regions \`
`WITH (VALUE_FORMAT='avro') AS \`
`SELECT gender, regionid , COUNT(*) AS numusers \`
`FROM pageviews_enriched \`

```
WINDOW TUMBLING (size 30 second) \
GROUP BY gender, regionid \
HAVING COUNT(*) > 1;
```

Output:

Message

Table created and running

- Create a new persistent query that counts the pageviews for each region and gender combination in a tumbling window of 30 seconds when the count is greater than one. Results from this query are written to the PAGEVIEWS_REGIONS Kafka topic in the Avro format. KSQL will register the Avro schema with the configured Schema Registry when it writes the first message to the PAGEVIEWS_REGIONS topic.

- ```
CREATE TABLE pageviews_regions \
WITH (VALUE_FORMAT='avro') AS \
SELECT gender, regionid , COUNT(*) AS numusers \
FROM pageviews_enriched \
WINDOW TUMBLING (size 30 second) \
GROUP BY gender, regionid \
HAVING COUNT(*) > 1;
```

Output:

Message

-----

Table created and running

-----

- Optional: View results from the above queries using SELECT.
  - ```
SELECT gender, regionid, numusers FROM pageviews_regions LIMIT 5;
```

Output:

FEMALE | Region_6 | 3

FEMALE | Region_1 | 4

FEMALE | Region_9 | 6

MALE | Region_8 | 2

OTHER | Region_5 | 4

LIMIT reached

Query terminated

ksql>

- Optional: Show all persistent queries.

- SHOW QUERIES;

Output:

Query ID	Kafka Topic	Query String
CSAS_PAGEVIEWS_FEMALE_1	PAGEVIEWS_FEMALE	CREATE STREAM pageviews_female AS SELECT * FROM pageviews_enriched WHERE gender = 'FEMALE';
CTAS_PAGEVIEWS_REGIONS_3	PAGEVIEWS_REGIONS	CREATE TABLE pageviews_regions WITH (VALUE_FORMAT='avro') AS SELECT gender, regionid, COUNT(*) AS numusers FROM pageviews_enriched WINDOW TUMBLING (size 30 second) GROUP BY gender, regionid HAVING COUNT(*) > 1;
CSAS_PAGEVIEWS_FEMALE_LIKE_89_2	PAGEVIEWS_FEMALE_LIKE_89	CREATE STREAM pageviews_female_like_89 WITH (kafka_topic='pageviews_enriched_r8_r9') AS SELECT * FROM pageviews_female WHERE regionid LIKE '%_8' OR regionid LIKE '%_9';
CSAS_PAGEVIEWS_ENRICHED_0	PAGEVIEWS_ENRICHED	CREATE STREAM pageviews_enriched AS SELECT users_original.userid AS userid, pageid, regionid, gender FROM pageviews_original LEFT JOIN users_original ON pageviews_original.userid = users_original.userid;

For detailed information on a Query run: EXPLAIN <Query ID>;

- Optional: Examine query run-time metrics and details. Observe that information including the target Kafka topic is available, as well as throughput figures for the messages being processed.

- DESCRIBE EXTENDED PAGEVIEWS_REGIONS;

Output:

Name : PAGEVIEWS_REGIONS
 Type : TABLE
 Key field : KSQL_INTERNAL_COL_0|+|KSQL_INTERNAL_COL_1
 Key format : STRING
 Timestamp field : Not set - using <ROWTIME>
 Value format : AVRO
 Kafka topic : PAGEVIEWS_REGIONS (partitions: 4, replication: 1)

Field | Type

ROWTIME | BIGINT (system)
ROWKEY | VARCHAR(STRING) (system)
GENDER | VARCHAR(STRING)
REGIONID | VARCHAR(STRING)
NUMUSERS | BIGINT

Queries that write into this TABLE

CTAS_PAGEVIEWS_REGIONS_3 : CREATE TABLE pageviews_regions WITH
(value_format='avro') AS SELECT gender, regionid , COUNT(*) AS numusers
FROM pageviews_enriched WINDOW TUMBLING (size 30 second) GROUP
BY gender, regionid HAVING COUNT(*) > 1;

For query topology and execution plan please run: EXPLAIN <QueryId>

Local runtime statistics

messages-per-sec: 3.06 total-messages: 1827 last-message: 7/19/18
4:17:55 PM UTC

failed-messages: 0 failed-messages-per-sec: 0 last-failed: n/a

(Statistics of the local KSQL server interaction with the Kafka topic
PAGEVIEWS_REGIONS)

ksql>

6 WORD COUNT LAMBDA INTEGRATION TEST:

<https://github.com/confluentinc/kafka-streams-examples/blob/5.1.1-post/src/test/java/io/confluent/examples/streams/WordCountLambdaIntegrationTest.java>

7 KSQL REST API TESTING

- GET the status of KSQL Server

```
curl -sX GET "http://<IP_Address>:8088/info";
```

Output:

```
{"KsqlServerInfo":{"version":"5.0.1","kafkaClusterId":"sNmrikrMSGGxZO3elbvQig","ksqlServiceId":"default_"}}
```

- Show all streams and tables using rest API

```
curl -X "POST" "<IP_Address>:8088/ksql" \  
-H "Content-Type: application/vnd.ksql.v1+json; charset=utf-8" \  
-d $'{  
  "ksql": "LIST STREAMS;",  
  "streamsProperties": {}  
'
```

Output:

```
[  
  {"@type":"streams","statementText":"LIST STREAMS;","streams":[  
    {"type":"STREAM","name":"PAGEVIEWS","topic":"_confluent-  
metrics","format":"DELIMITED"},  
    {"type":"STREAM","name":"TEST_STREAM","topic":"_confluent-  
metrics","format":"DELIMITED"}}]  
]
```


8 TESTING CONFLUENT KAFKA BROKER AND KAFKA ZOOKEEPER SERVICES

- Create a Kafka topic

To create a topic called Test in running cluster, use the following command:

```
/bin/kafka-topics --create --zookeeper <Zookeeper_1_IP_Address>:  
:2181<Zookeeper_2_IP_Address>:2181, <Zookeeper_3_IP_Address>  
:2181 --replication-factor 1 --partitions 1 --topic Test
```

Output:

Created topic "Test".

Note: kafka_zookeeper service will be used to create Kafka topic

- List the Kafka topic

To list the topics in running cluster, use the following command:

```
/bin/kafka-topics --list --zookeeper <Zookeeper_1_IP_Address>  
: 2181, <Zookeeper_2_IP_Address>:2181<Zookeeper_3_IP_Address>:2181
```

Output:

TEST-STREAMING

Test

__confluent.support.metrics

__consumer_offsets

_confluent-command

_confluent-ksql-default__command_topic

_confluent-metrics

_confluent-monitoring

_schemas

connect-configs

connect-offsets

connect-status

Note: This command returns the list with the names of all of the running topics in the

Cluster.

- Produce the data in the Kafka topic

To create a Kafka Producer, run the following:

```
/bin/kafka-console-producer --broker-list <Broker_1_IP_Address>:9092,  
<Broker_2_IP_Address>:9092, <Broker_3_IP_Address>:9092 --topic Test
```

Now we can start sending the messages to the Kafka cluster from the console. The Messages will be published to the Kafka topic, "test".

Output:

```
>1, First Entry  
>2, Second Entry  
>3, Third Entry  
>4, Fourth Entry  
>5, Fifth Entry
```

Note: The output is just an example you can send message according to your requirement.

- Read data from a Kafka topic

To read data from Kafka topic, run the following command:

```
/bin/kafka-console-consumer --bootstrap-server list <Broker_1_IP_Address>:9092,  
<Broker_2_IP_Address>:9092, <Broker_3_IP_Address>:9092 --topic Test --from-  
beginning
```

Start typing messages in the producer. Consumer would get the messages via Kafka topic

Output:

```
1, First Entry  
2, Second Entry  
3, Third Entry  
4, Fourth Entry  
5, Fifth Entry  
Processed a total of 5 messages
```

9 TESTING SCHEMA REGISTRY

- Registering a new version of a Schema under the subject “Kafka-key”

```
curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" \  
  --data '{"schema": "{\"type\": \"string\"}"}' \  
  http://<Schema_registry_IP_address>:8081/subjects/Kafka-key/versions
```

Output:

```
{"id":1}
```

- Registering a new version of a Schema under the subject "Kafka-value"

```
curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" \  
  --data '{"schema": "{\"type\": \"string\"}"}' \  
  http://<Schema_registry_IP_address>:8081/subjects/Kafka-value/versions
```

Output:

```
{"id":1}
```

- Listing all Subjects

```
curl -X GET http://<Schema_Registry_IP_address>:8081/subjects
```

Output:

```
["Kafka-value", "Kafka-key"]
```

- Fetching a Schema by Globally Unique ID

```
curl -X GET http://<Schema_Registry_IP_address>:8081/schemas/ids/1
```

Output:

```
{"schema": "\"string\""}
```

- Listing all Schema versions registered under the subject “Kafka-value”

```
curl -X GET http://<Schema_registry_IP_address>:8081/subjects/Kafka-value/versions
```

Output:

```
[1]
```

- Fetch Version 1 of the Schema registered under subject “Kafka-value”

```
curl -X GET http://<Schema_Registry_IP_address>:8081/subjects/Kafka-value/versions/1
```

Output:

```
{"subject":"Kafka-value","version":1,"id":1,"schema":"\string\""}

```

- Checking if a Schema is registered under subject “Kafka-key”

```
curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" \
--data '{"schema": {"type": "string"}}' \
http://<Schema_Registry_IP_address>:8081/subjects/Kafka-key
```

Output:

```
{"subject":"Kafka-key","version":1,"id":1,"schema":"\string\""}

```

- Testing compatibility of a Schema with the latest schema under subject “Kafka-value”

```
curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" \
--data '{"schema": {"type": "string"}}' \
http://<Schema_Registry_IP_address>:8081/compatibility/subjects/Kafka-value/versions/latest
```

Output:

```
{"is_compatible":true}

```

- Getting the top level config

```
curl -X GET http://<Schema_Registry_IP_address>:8081/config
```

Output:

```
{"compatibilityLevel":"BACKWARD"}

```

- Updating compatibility requirements globally

```
curl -X PUT -H "Content-Type: application/vnd.schemaregistry.v1+json" \  
  --data '{"compatibility": "NONE"}' \  
  http://<Schema_Registry_IP_address>:8081/config
```

Output:

```
{"compatibility": "NONE"}
```

- Deleting all Schema versions registered under subject “Kafka-value”

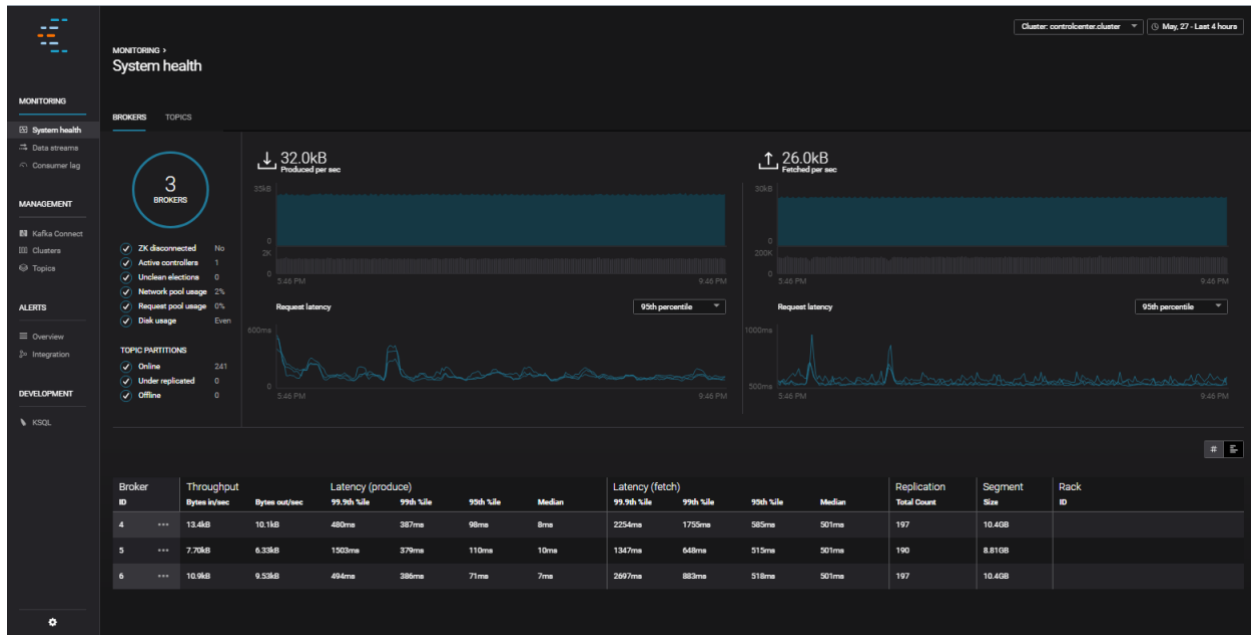
```
curl -X DELETE http://<Schema_registry_IP_address>:8081/subjects/Kafka-value
```

Output:

```
[1]
```

10 MONITOR AND MANAGE – CONTROL CENTER

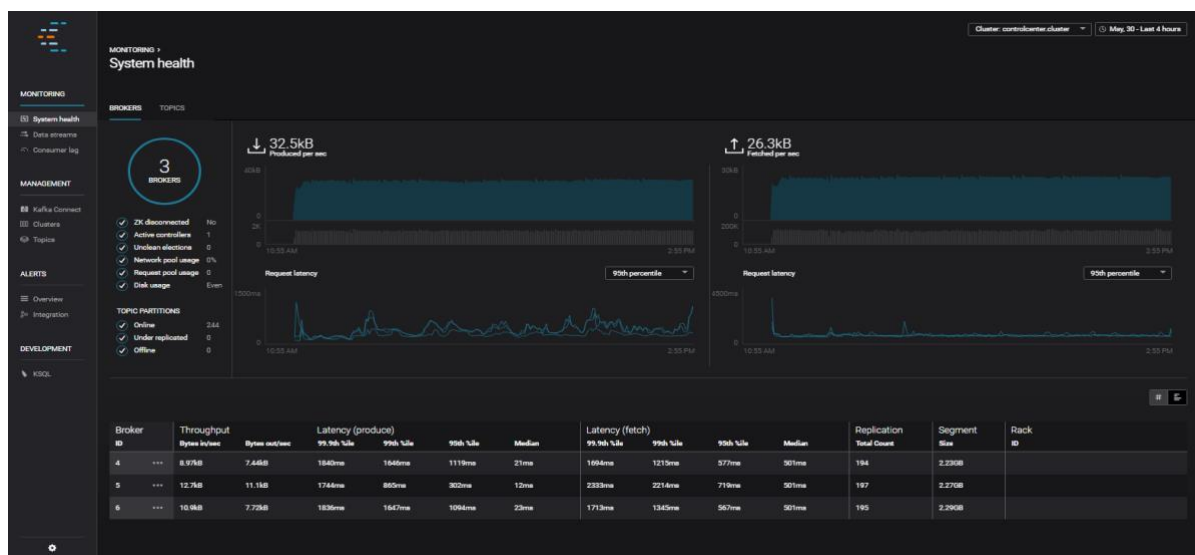
Open your web browser and navigate to http://<IP_Address>:9021 (by default). As you will enter you will see the Confluent dashboard, as shown below:



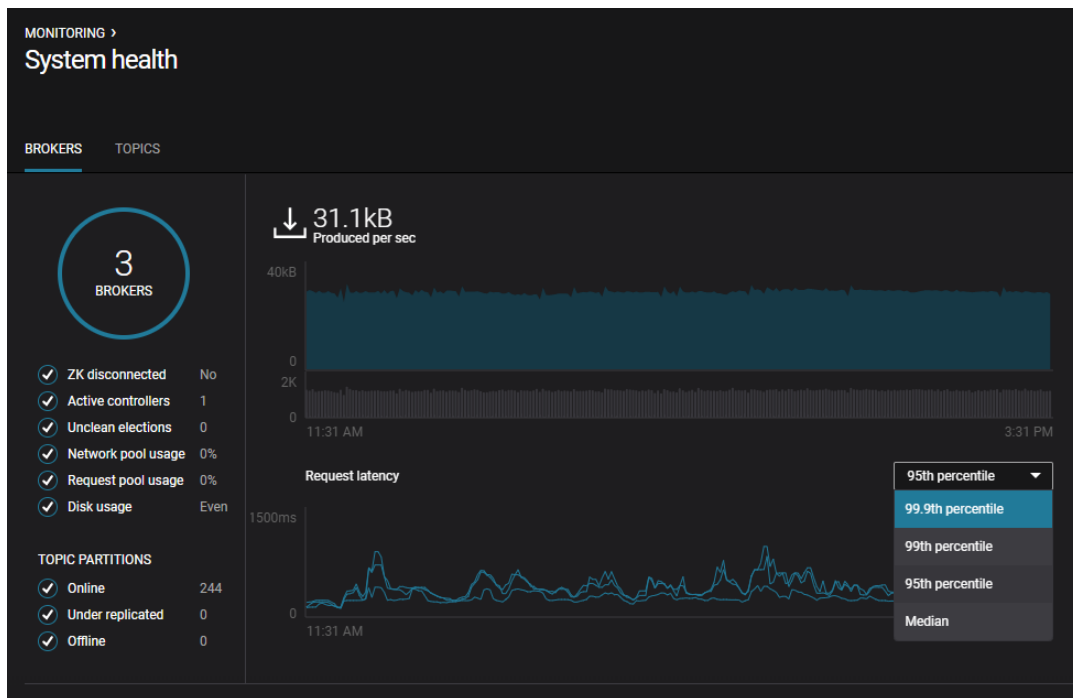
10.1.1 System health

From the left-hand menu, click on System health. System Health page will appear and it will display the health of broker and topic.

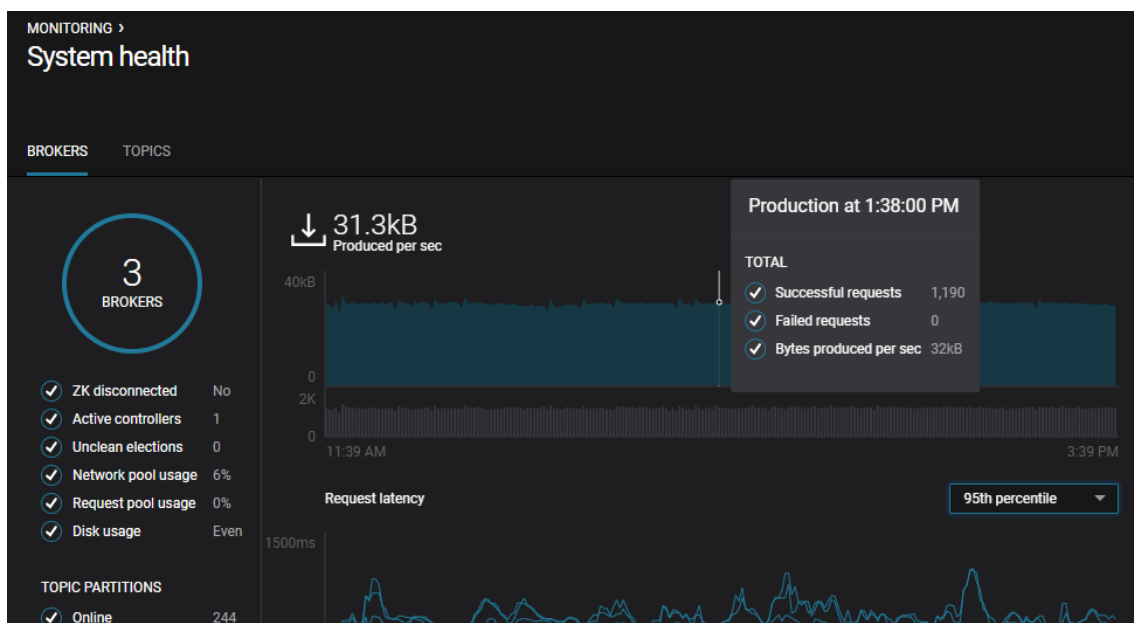
Broker tab:



- From the Request latency percentile drop-down list, select any percentile and view the change



- Mouse over the charts and you can view details like Successful requests, Failed requests and Bytes produced/fetched per sec



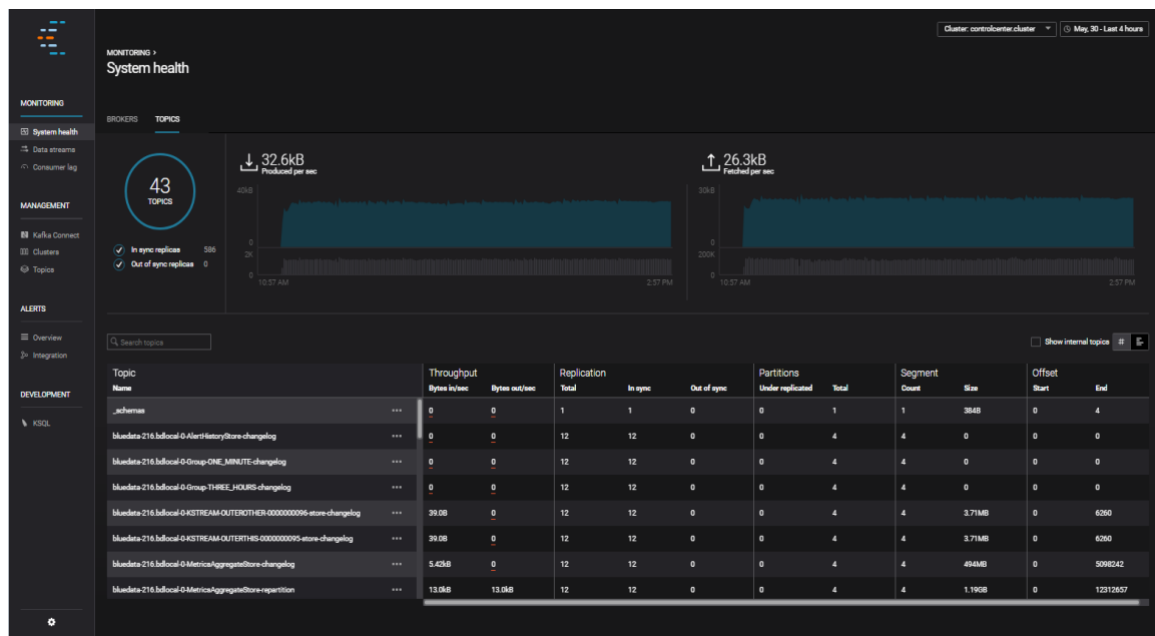
- Summary of the Produce and Fetch requests are given like below




- At the bottom, you will see the Broker Metrics table

Broker ID		Throughput		Latency (produce)				Latency (fetch)				Replication Total Count	Segment Size	Rack ID
		Bytes in/sec	Bytes out/sec	99.9th %ile	99th %ile	95th %ile	Median	99.9th %ile	99th %ile	95th %ile	Median			
4	---	9.24kB	7.87kB	2163ms	1819ms	535ms	13ms	1893ms	1039ms	597ms	501ms	194	2.30GB	
5	---	12.4kB	11.4kB	1246ms	754ms	307ms	12ms	2321ms	2150ms	754ms	501ms	197	2.34GB	
6	---	10.2kB	7.31kB	2012ms	1728ms	586ms	16ms	2528ms	1046ms	616ms	501ms	195	2.36GB	

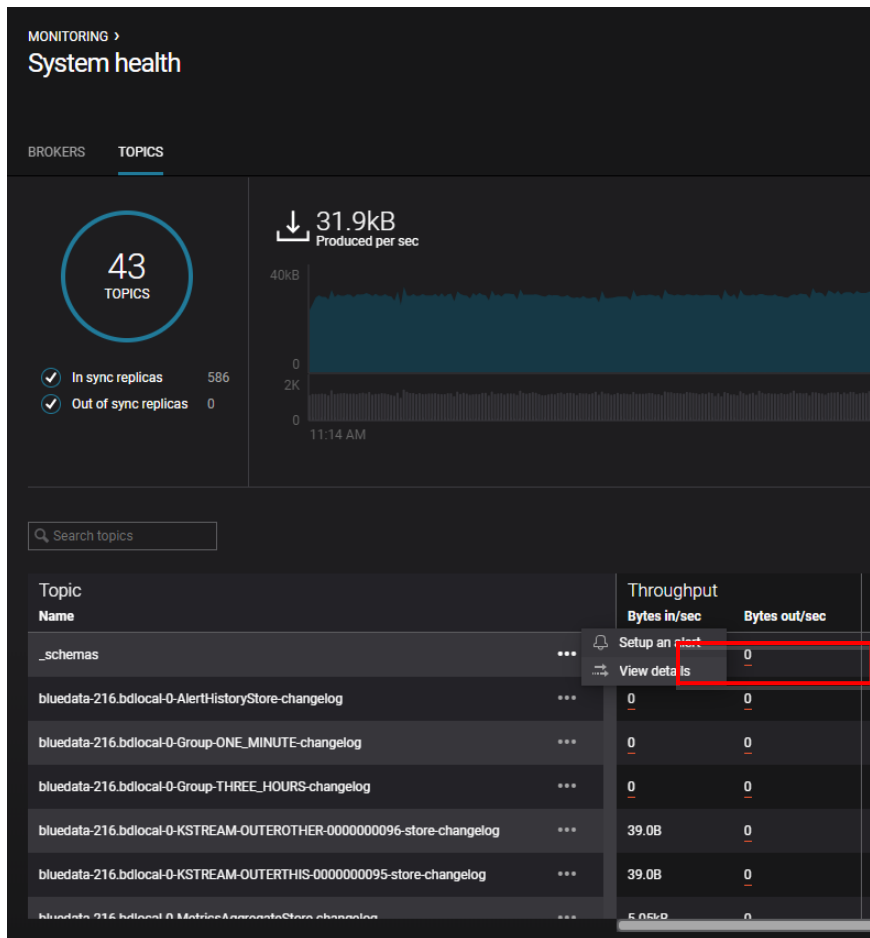
Topic tab:



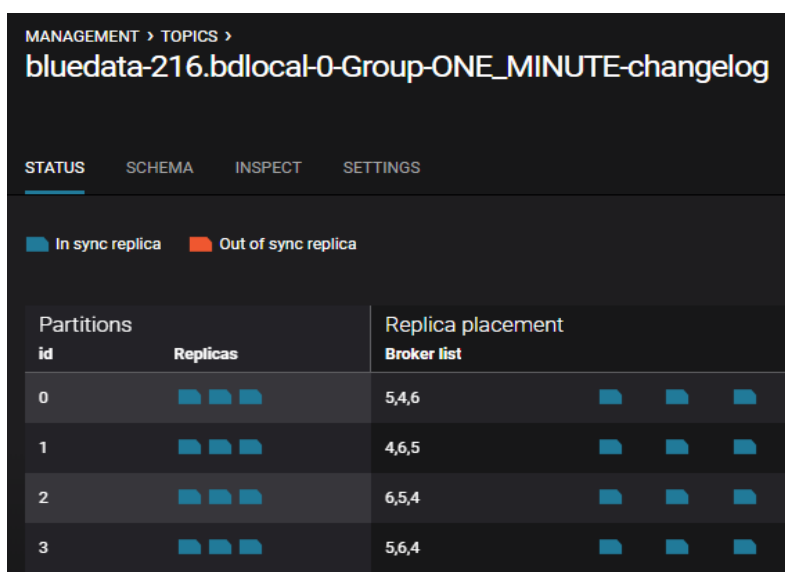
- Topic Aggregate Metrics shows the total number of topics, In sync replicas and out of sync replicas



- Click on the ellipsis of any topic, under Topics section and click on View details



- You will get screen like below, showing the In sync replica and out of sync replica (if any)





- Click on SETTINGS, to edit settings or Show full config button, as per requirement

MANAGEMENT > TOPICS > **bluedata-216.bdlocal-0-Group-ONE_MINUTE-changelog**

STATUS SCHEMA INSPECT **SETTINGS**

min.insync.replicas	2
cleanup.policy	compact,delete
retention.ms	432000000
retention.bytes	-1

 Edit settings  Show full config

- At the bottom, you will see Topic Metric table

Topic Name	Throughput		Replication			Partitions		Segment		Offset	
	Bytes in/sec	Bytes out/sec	Total	In sync	Out of sync	Under replicated	Total	Count	Size	Start	End
bluedata-216.bdlocal-0-AlertHistoryStore-changelog	0	0	12	12	0	0	4	4	0	0	0
bluedata-216.bdlocal-0-Group-ONE_MINUTE-changelog	0	0	12	12	0	0	4	4	0	0	0
bluedata-216.bdlocal-0-Group-THREE_HOURS-changelog	0	0	12	12	0	0	4	4	0	0	0
bluedata-216.bdlocal-0-KSTREAM-OUTEROTHER-0000000095-store-changelog	39.0B	0	12	12	0	0	4	4	3.83MB	0	6451
bluedata-216.bdlocal-0-KSTREAM-OUTERTHS-0000000095-store-changelog	39.0B	0	12	12	0	0	4	4	3.83MB	0	6451
bluedata-216.bdlocal-0-MetricsAggregateStore-changelog	4.31kB	0	12	12	0	0	4	4	509MB	0	5253713

- Likewise explore other options to get more information


10.1.2 Kafka Connect

- From the left-hand menu, under MANAGEMENT section, click on Kafka Connect
- To create a source, provide the value of following parameters and click on Submit

MANAGEMENT > KAFKA CONNECT > SOURCES > **New**

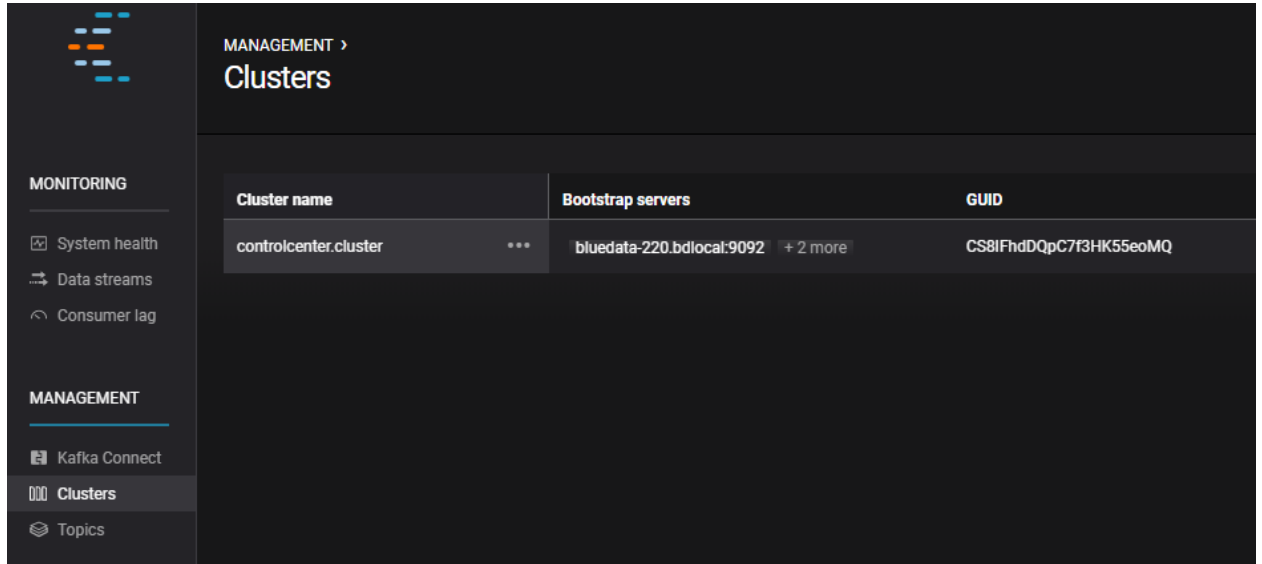
SOURCES SINKS

How should we connect to your data?

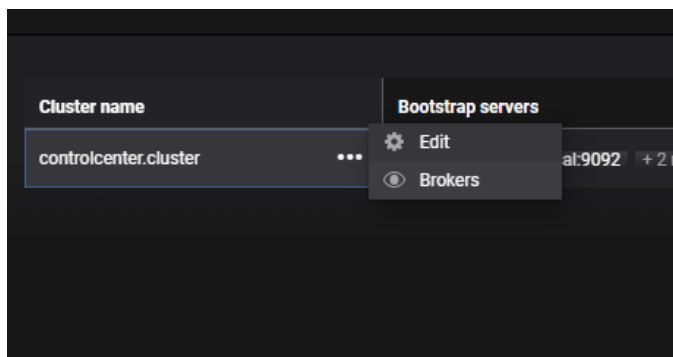
Connector Class*  Name*

10.1.3 Clusters

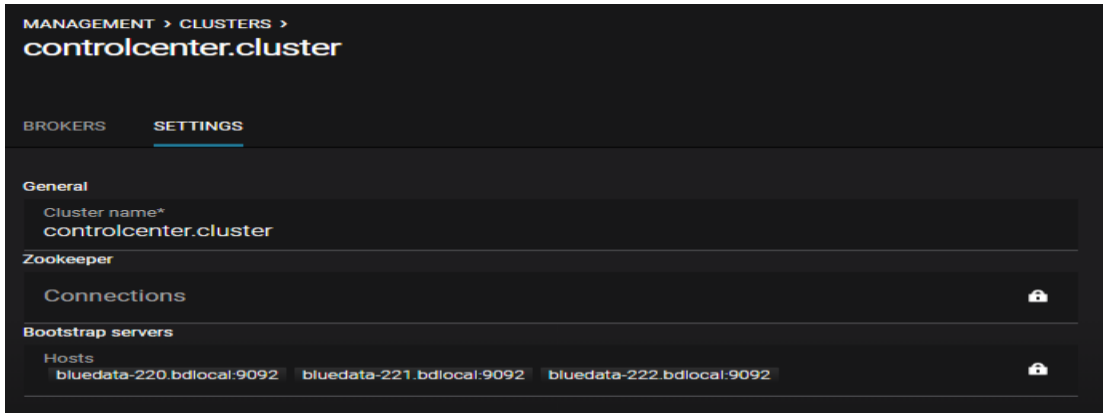
- From the left-hand menu, under MANAGEMENT section, click on Clusters



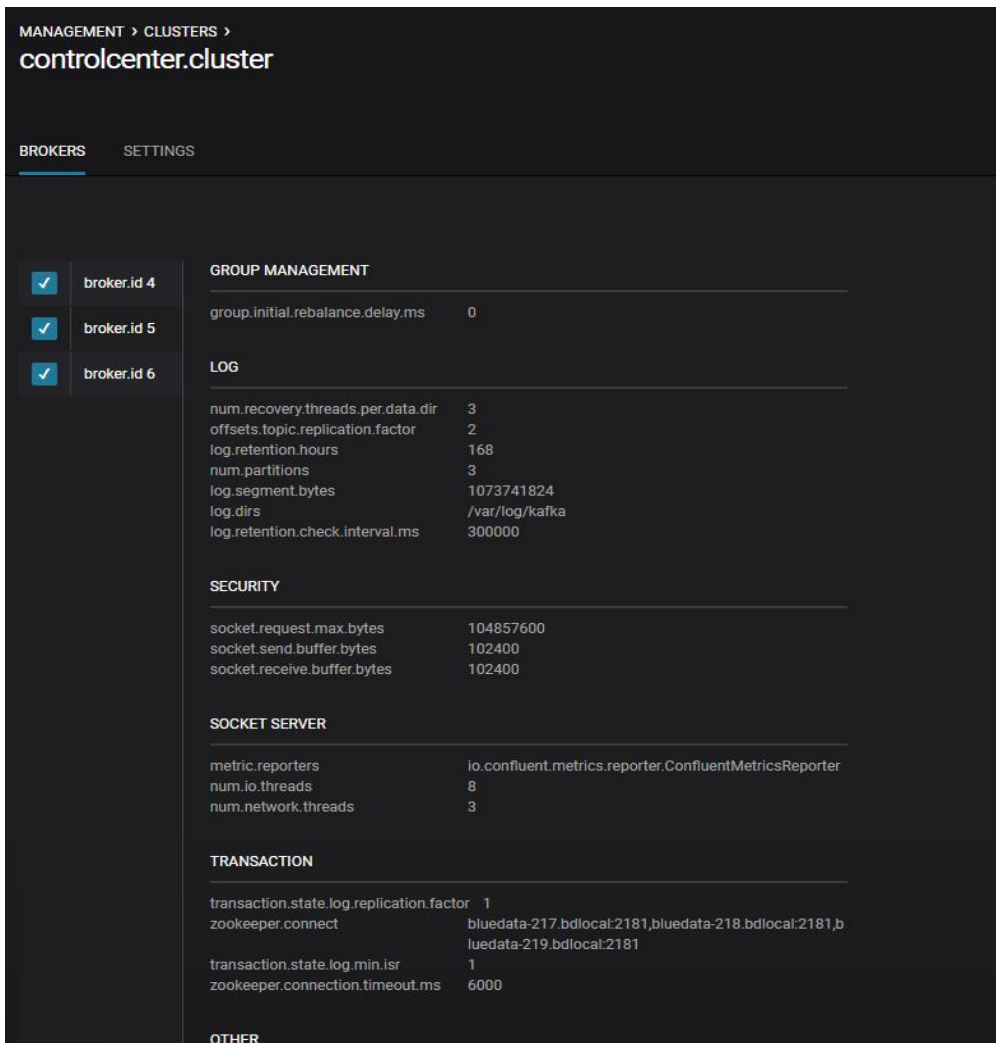
- Click on ellipsis of the cluster name to edit or Brokers button as per requirement



- The Edit window looks like this below



- The Broker information is listed as below, it can also be downloaded



10.1.4 Topics

- From the left-hand menu, under MANAGEMENT section, click on Topics

MANAGEMENT > Topics

Search topics ☐ Show internal topics

Topics Name	Partitions Total	Replication Factor	% In sync	Data flow Consumer Groups
__schemas	1	x1	100%	0
bluedata-216.bdlocal-0-AlertHistoryStore-changelog	4	x3	100%	0
bluedata-216.bdlocal-0-Group-ONE_MINUTE-changelog	4	x3	100%	0
bluedata-216.bdlocal-0-Group-THREE_HOURS-changelog	4	x3	100%	0
bluedata-216.bdlocal-0-KSTREAM-OUTEROTHER-0000000096-store-changelog	4	x3	100%	0
bluedata-216.bdlocal-0-KSTREAM-OUTERTHIS-0000000095-store-changelog	4	x3	100%	0
bluedata-216.bdlocal-0-MetricsAggregateStore-changelog	4	x3	100%	0
bluedata-216.bdlocal-0-MetricsAggregateStore-repartition	4	x3	100%	0
bluedata-216.bdlocal-0-MonitoringMessageAggregatorWindows-ONE_MINUTE-chang...	4	x3	100%	0
bluedata-216.bdlocal-0-MonitoringMessageAggregatorWindows-THREE_HOURS-cha...	4	x3	100%	0
bluedata-216.bdlocal-0-MonitoringStream-ONE_MINUTE-changelog	4	x3	100%	0
bluedata-216.bdlocal-0-MonitoringStream-THREE_HOURS-changelog	4	x3	100%	0
bluedata-216.bdlocal-0-MonitoringTriggerStore-changelog	4	x3	100%	0
bluedata-216.bdlocal-0-MonitoringVerifierStore-changelog	4	x3	100%	0
bluedata-216.bdlocal-0-TriggerActionsStore-changelog	4	x3	100%	0

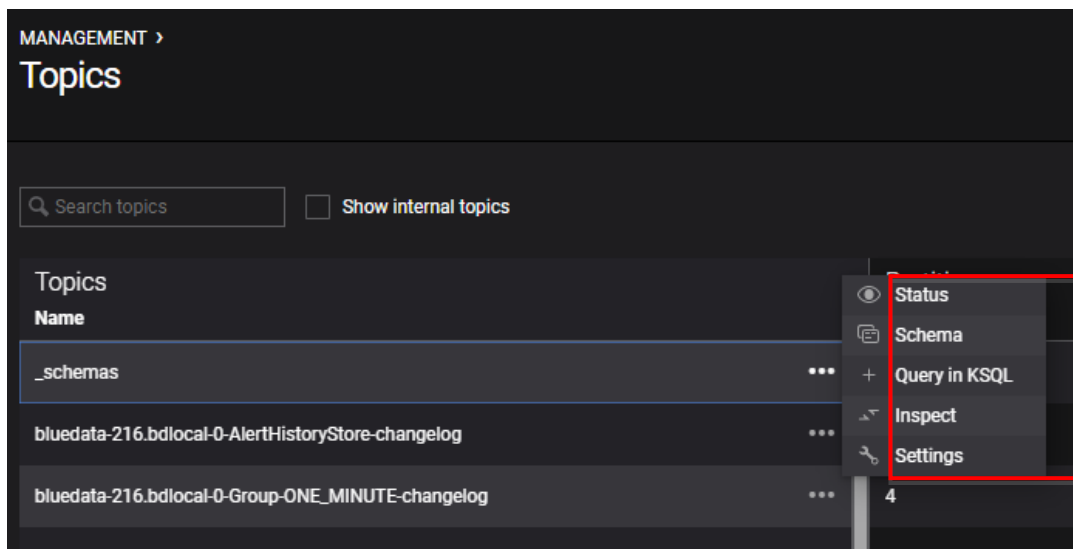
- Check Show internal topics, to view all the internal topics

MANAGEMENT > Topics


Search topics ☒ Show internal topics

Topics Name	Partitions Total
__confluent.support.metrics	1
__confluent-command	1

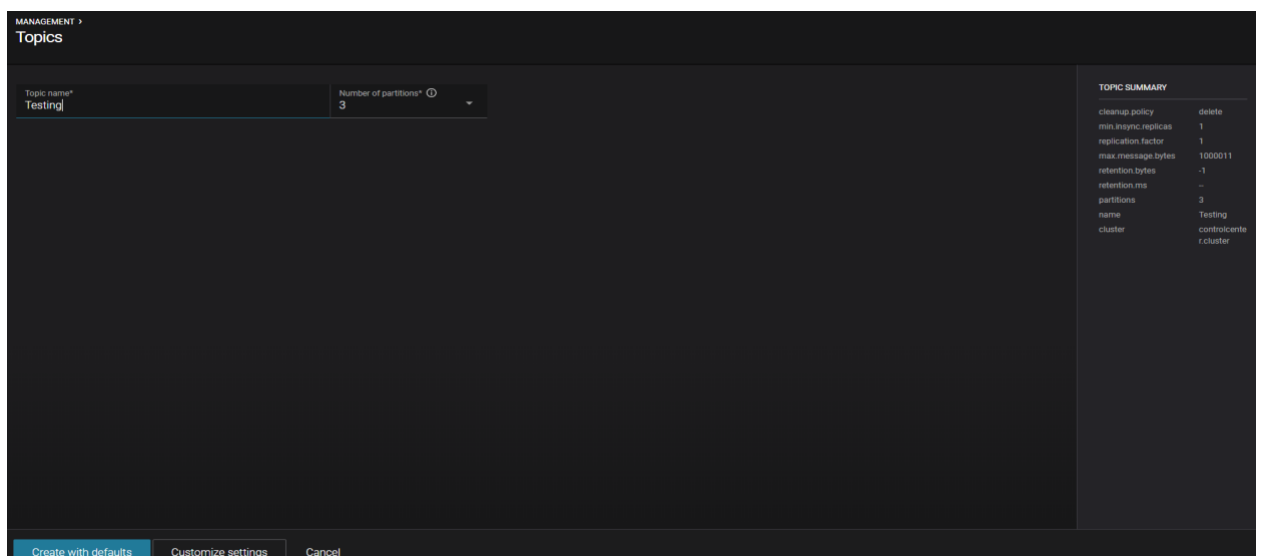
- Click on the Topics ellipsis to view the multiple options related to that topic



- To create a new topic, click on Create topic, at right-hand side

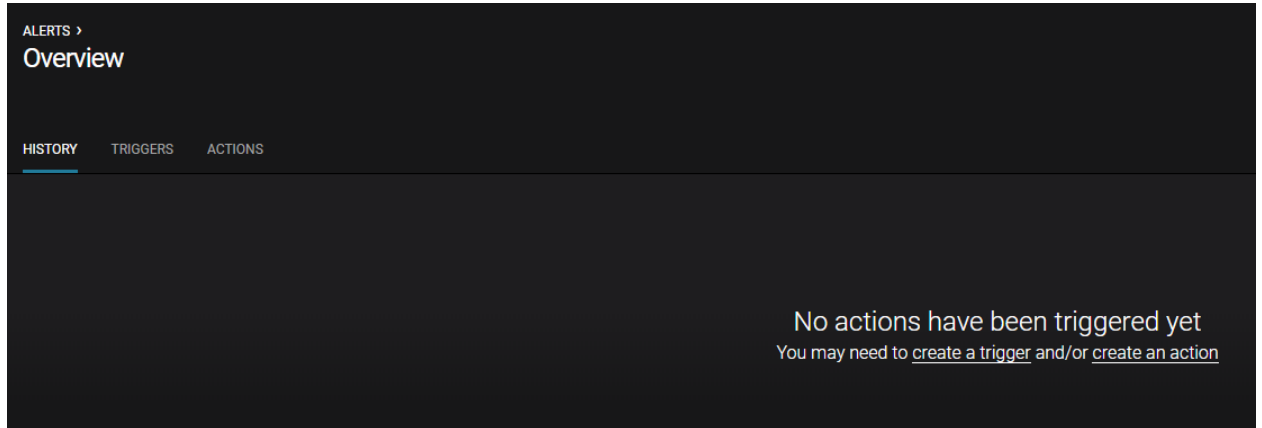


- Provide topic name and number of partitions for topic creation and click on Create with defaults or Customize settings, according to requirement

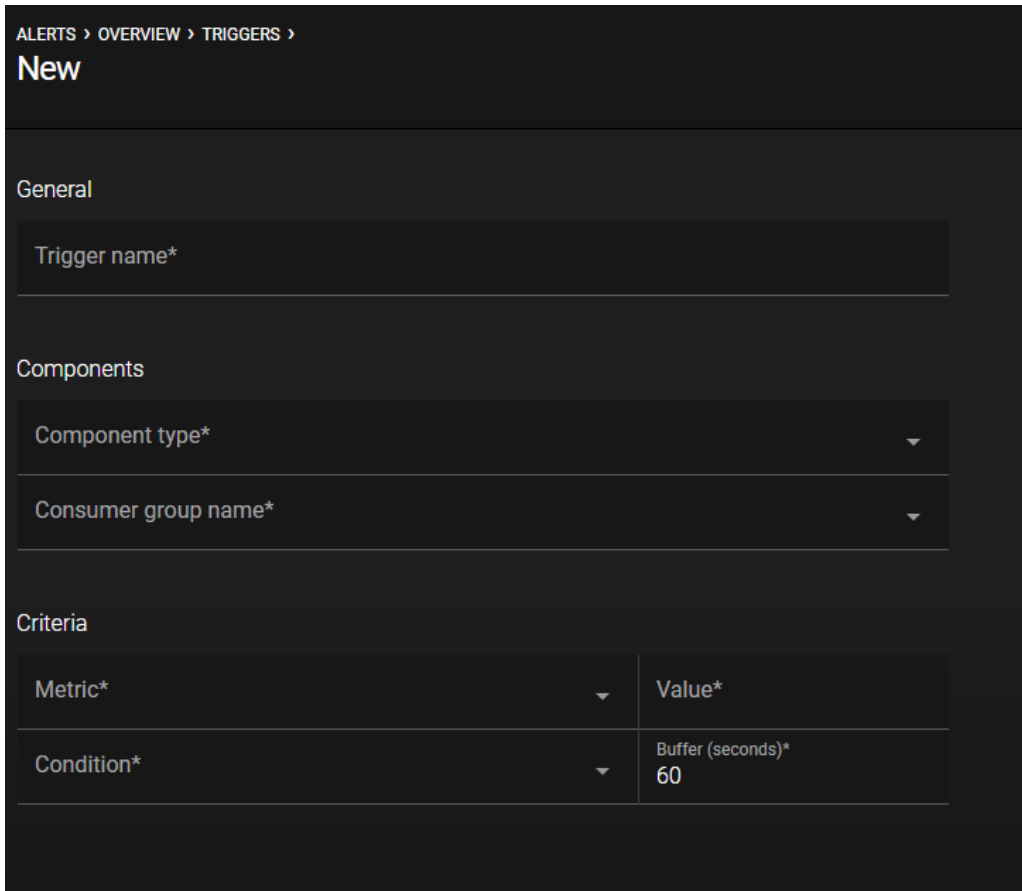


10.1.5 Alerts

- From the left-hand menu, under ALERTS section, click on Overview



- It allows you to create a Trigger and/or an Action
- To create a Trigger, provide the values of the following parameters and click on Submit



Criteria	
Metric*	Value*
Condition*	Buffer (seconds)* 60

- To create an Action, provide the value for the following parameters and click on Submit

ALERTS > OVERVIEW > ACTIONS >

New

General

Action Name* enabled ☐

Triggers

Triggers* ▼

Actions

Action* 🔒
Send email

Recipient email address*

Subject*

Max send rate* Per hour ▼

- Click on Integration, under the ALERTS section, it provides REST API endpoints

ALERTS >

Integration

REST API

[GET](#) /2.0/alerts/history

Sample Response http://10.0.1.163:10000/2.0/alerts/history

```
[ ]
```

Optional URL parameters

limit	Integer representing the maximum number of records to return.	Default	50
ts	Timestamp defining the most recent alert to return (in milliseconds since epoch).	Default	null

11 TESTING KAFKA CONNECT

- Get a list of active connectors (Initially empty)

```
curl -X GET http://<Kafka_Connect_IP_Address>:8083/connectors
```

Output:

```
[]
```

- Create a new connector

- HDFS Sink Connector

```
curl -H "Content-Type: application/json" \
--request POST \
--data '{
  "name": "hdfs-sink-connector",
  "config": {
    "connector.class": "io.confluent.connect.hdfs.HdfsSinkConnector",
    "tasks.max": "10",
    "topics": "test-topic",
    "hdfs.url": "hdfs://<IP_Address>:9000",
    "hadoop.conf.dir": "/opt/hadoop/conf",
    "hadoop.home": "/opt/hadoop",
    "flush.size": "100",
    "rotate.interval.ms": "1000"
  }
}' \
http://<Kafka_Connect_IP_Address>:8083/connectors
```

Output:

```
{"name":"hdfs-sink-connector","config":{"connector.class":"io.confluent.connect.hdfs.HdfsSinkConnector",
"tasks.max":"10","topics":"test-topic","hdfs.url":"hdfs://10.39.250.64:9000","hadoop.conf.dir":"/opt/hadoop/conf","hadoop.home":"/opt/hadoop","flush.size":"100","rotate.interval.ms":"1000","name":"hdfs-sink-connector"},"tasks":[{"connector":"hdfs-sink-connector","task":0}, {"connector":"hdfs-sink-connector","task":1}, {"connector":"hdfs-sink-connector","task":2}, {"connector":"hdfs-sink-connector","task":3}, {"connector":"hdfs-sink-connector","task":4}, {"connector":"hdfs-sink-connector","task":5}, {"connector":"hdfs-sink-connector","task":6}, {"connector":"hdfs-sink-connector","task":7}, {"connector":"hdfs-sink-connector","task":8}, {"connector":"hdfs-sink-connector","task":9}], "type":"sink"}
```

- Local File Sink

```
curl -H "Content-Type: application/json" \
--request POST \
--data '{
```

```
"name": "local-file-sink",
"config": {
  "connector.class": "FileStreamSinkConnector",
  "tasks.max": "2",
  "topics": "connect-test",
  "hdfs.url": "hdfs://<IP_Address>:9000",
  "hadoop.conf.dir": "/opt/hadoop/conf",
  "hadoop.home": "/opt/hadoop",
  "flush.size": "100",
  "rotate.interval.ms": "1000"
}
}' \
http://<Kafka_Connect_IP_Address>:8083/connectors
```

Output:

```
{"name":"local-file-sink","config":{"connector.class":"FileStreamSinkConnector","tasks.max":"2","topics":"connect-test","hdfs.url":"hdfs://10.39.250.16:9000","hadoop.conf.dir":"/opt/hadoop/conf","hadoop.home":"/opt/hadoop","flush.size":"100","rotate.interval.ms":"1000","name":"local-file-sink"},"tasks":[],"type":"sink"}
```

- Verify the connector is added by executing

```
curl -X GET http://<Kafka_Connect_IP_Address>:8083/connectors
```

Output:

```
["local-file-sink","hdfs-sink-connector"]
```

Note: hdfs-sink-connector and local-file-sink are the connector name here, which will be used in later commands. Replace it with any other string name (connector name) according to your environment.

- Get connector status

```
curl -X GET http://<Kafka_Connect_IP_Address>:8083/connectors/local-file-sink/status
```

Output:

```
{"name":"local-file-sink","connector":{"state":"RUNNING","worker_id":"bluedata-4424.bdlocal:8083"},"tasks":[{"id":0,"state":"RUNNING","worker_id":"bluedata-4424.bdlocal:8083"},{"id":1,"state":"RUNNING","worker_id":"bluedata-4424.bdlocal:8083"}],"type":"sink"}
```

- Get worker's version

```
curl -X GET http://<Kafka_Connect_IP_Address>:8083/
```

Output:

```
{"version":"2.0.1-cp4","commit":"49da0fef3e389dc2","kafka_cluster_id":"CS8lFhdDQpC7f3HK55eoMQ"}
```

- Get the list of connector plugins available on the worker
`curl -X GET http://<Kafka_Connect_IP_Address>:8083/connector-plugins`

Output:

```
[{"class": "io.confluent.connect.activemq.ActiveMQSourceConnector", "type": "source", "version": "5.0.3"}, {"class": "io.confluent.connect.elasticsearch.ElasticsearchSinkConnector", "type": "sink", "version": "5.0.3"}, {"class": "io.confluent.connect.hdfs.HdfsSinkConnector", "type": "sink", "version": "5.0.3"}, {"class": "io.confluent.connect.hdfs.tools.SchemaSourceConnector", "type": "source", "version": "2.0.1-cp4"}, {"class": "io.confluent.connect.ibm.mq.IbmMQSourceConnector", "type": "source", "version": "5.0.3"}, {"class": "io.confluent.connect.jdbc.JdbcSinkConnector", "type": "sink", "version": "5.0.3"}, {"class": "io.confluent.connect.jdbc.JdbcSourceConnector", "type": "source", "version": "5.0.3"}, {"class": "io.confluent.connect.jms.JmsSourceConnector", "type": "source", "version": "5.0.3"}, {"class": "io.confluent.connect.replicator.ReplicatorSourceConnector", "type": "source", "version": "5.0.3"}, {"class": "io.confluent.connect.s3.S3SinkConnector", "type": "sink", "version": "5.0.3"}, {"class": "io.confluent.connect.storage.tools.SchemaSourceConnector", "type": "source", "version": "2.0.1-cp4"}, {"class": "org.apache.kafka.connect.file.FileStreamSinkConnector", "type": "sink", "version": "2.0.1-cp4"}, {"class": "org.apache.kafka.connect.file.FileStreamSourceConnector", "type": "source", "version": "2.0.1-cp4"}]
```

- Get tasks for a connector

```
curl -X GET http://<Kafka_Connect_IP_Address>:8083/connectors/hdfs-sink-connector/tasks
```

Output:

```
{
  "id": "connector": "hdfs-sink-connector",
  "task": 0,
  "config": {
    "connector.class": "io.confluent.connect.hdfs.HdfsSinkConnector",
    "task.class": "io.confluent.connect.hdfs.HdfsSinkTask",
    "hadoop.conf.dir": "/opt/hadoop/conf",
    "flush.size": "100",
    "tasks.max": "10",
    "topics": "test-topic",
    "hdfs.url": "hdfs://10.39.250.64:9000",
    "name": "hdfs-sink-connector",
    "rotate.interval.ms": "1000",
    "hadoop.home": "/opt/hadoop"
  },
  "id": "connector": "hdfs-sink-connector",
  "task": 1,
  "config": {
    "connector.class": "io.confluent.connect.hdfs.HdfsSinkConnector",
    "task.class": "io.confluent.connect.hdfs.HdfsSinkTask",
    "hadoop.conf.dir": "/opt/hadoop/conf",
    "flush.size": "100",
    "tasks.max": "10",
    "topics": "test-topic",
    "hdfs.url": "hdfs://10.39.250.64:9000",
    "name": "hdfs-sink-connector",
    "rotate.interval.ms": "1000",
    "hadoop.home": "/opt/hadoop"
  },
  "id": "connector": "hdfs-sink-connector",
  "task": 2,
  "config": {
    "connector.class": "io.confluent.connect.hdfs.HdfsSinkConnector",
    "task.class": "io.confluent.connect.hdfs.HdfsSinkTask",
    "hadoop.conf.dir": "/opt/hadoop/conf",
    "flush.size": "100",
    "tasks.max": "10",
    "topics": "test-topic",
    "hdfs.url": "hdfs://10.39.250.64:9000",
    "name": "hdfs-sink-connector",
    "rotate.interval.ms": "1000",
    "hadoop.home": "/opt/hadoop"
  },
  "id": "connector"
}
```

```

": "hdfs-sink-
connector", "task": 3}, "config": {"connector.class": "io.confluent.connect.hdfs.HdfsSinkCo
nnect", "task.class": "io.confluent.connect.hdfs.HdfsSinkTask", "hadoop.conf.dir": "/opt/
hadoop/conf", "flush.size": "100", "tasks.max": "10", "topics": "test-
topic", "hdfs.url": "hdfs://10.39.250.64:9000", "name": "hdfs-sink-
connector", "rotate.interval.ms": "1000", "hadoop.home": "/opt/hadoop"}}, {"id": "connector
": "hdfs-sink-
connector", "task": 4}, "config": {"connector.class": "io.confluent.connect.hdfs.HdfsSinkCo
nnect", "task.class": "io.confluent.connect.hdfs.HdfsSinkTask", "hadoop.conf.dir": "/opt/
hadoop/conf", "flush.size": "100", "tasks.max": "10", "topics": "test-
topic", "hdfs.url": "hdfs://10.39.250.64:9000", "name": "hdfs-sink-
connector", "rotate.interval.ms": "1000", "hadoop.home": "/opt/hadoop"}}, {"id": "connector
": "hdfs-sink-
connector", "task": 5}, "config": {"connector.class": "io.confluent.connect.hdfs.HdfsSinkCo
nnect", "task.class": "io.confluent.connect.hdfs.HdfsSinkTask", "hadoop.conf.dir": "/opt/
hadoop/conf", "flush.size": "100", "tasks.max": "10", "topics": "test-
topic", "hdfs.url": "hdfs://10.39.250.64:9000", "name": "hdfs-sink-
connector", "rotate.interval.ms": "1000", "hadoop.home": "/opt/hadoop"}}, {"id": "connector
": "hdfs-sink-
connector", "task": 6}, "config": {"connector.class": "io.confluent.connect.hdfs.HdfsSinkCo
nnect", "task.class": "io.confluent.connect.hdfs.HdfsSinkTask", "hadoop.conf.dir": "/opt/
hadoop/conf", "flush.size": "100", "tasks.max": "10", "topics": "test-
topic", "hdfs.url": "hdfs://10.39.250.64:9000", "name": "hdfs-sink-
connector", "rotate.interval.ms": "1000", "hadoop.home": "/opt/hadoop"}}, {"id": "connector
": "hdfs-sink-
connector", "task": 7}, "config": {"connector.class": "io.confluent.connect.hdfs.HdfsSinkCo
nnect", "task.class": "io.confluent.connect.hdfs.HdfsSinkTask", "hadoop.conf.dir": "/opt/
hadoop/conf", "flush.size": "100", "tasks.max": "10", "topics": "test-
topic", "hdfs.url": "hdfs://10.39.250.64:9000", "name": "hdfs-sink-
connector", "rotate.interval.ms": "1000", "hadoop.home": "/opt/hadoop"}}, {"id": "connector
": "hdfs-sink-
connector", "task": 8}, "config": {"connector.class": "io.confluent.connect.hdfs.HdfsSinkCo
nnect", "task.class": "io.confluent.connect.hdfs.HdfsSinkTask", "hadoop.conf.dir": "/opt/
hadoop/conf", "flush.size": "100", "tasks.max": "10", "topics": "test-
topic", "hdfs.url": "hdfs://10.39.250.64:9000", "name": "hdfs-sink-
connector", "rotate.interval.ms": "1000", "hadoop.home": "/opt/hadoop"}}, {"id": "connector
": "hdfs-sink-
connector", "task": 9}, "config": {"connector.class": "io.confluent.connect.hdfs.HdfsSinkCo
nnect", "task.class": "io.confluent.connect.hdfs.HdfsSinkTask", "hadoop.conf.dir": "/opt/
hadoop/conf", "flush.size": "100", "tasks.max": "10", "topics": "test-
topic", "hdfs.url": "hdfs://10.39.250.64:9000", "name": "hdfs-sink-
connector", "rotate.interval.ms": "1000", "hadoop.home": "/opt/hadoop"}}]

```

- Restart connector and its tasks

```
curl -X POST http://<Kafka_Connect_IP_Address>:8083/connectors/hdfs-sink-
connector/restart
```

Output:

<No response when successful>

- Pause the connector and its tasks

```
curl -X PUT http://<Kafka_Connect_IP_Address>:8083/connectors/hdfs-sink-connector/pause
```

Output:

<No response when successful>

- Resume the connector and its tasks

```
curl -X PUT http://<Kafka_Connect_IP_Address>:8083/connectors/hdfs-sink-connector/resume
```

Output:

<No response when successful>

- Restart an individual task

```
curl -X POST http://<Kafka_Connect_IP_Address>:8083/connectors/hdfs-sink-connector/tasks/0/restart
```

Output:

<No response when successful>

- Delete a connector

```
curl -X DELETE http://<Kafka_Connect_IP_Address>:8083/connectors/hdfs-sink-connector
```

Output:

<No response when successful>

12 CONFIGURING CONFLUENT KAFKA WITH HDFS AND DTAP USING CONNECTORS

Here in this section, we will see how to receive message from a Kafka topic and store it in HDFS and DTAP, using connector.

- SSH into Kafka connect node

Note: It is assumed that Cloudera Manager Repo is installed in the /etc/yum.repos.d/ directory and Hadoop-client is installed on Kafka connect node.

- Create a Kafka topic

```
/bin/kafka-topics --create --zookeeper  
<Zookeeper_1_IP_Address>,<Zookeeper_2_IP_Address>,<Zookeeper_3_IP_Addres  
>:2181 --replication-factor 1 --partitions 1 --topic blue
```

Output:

Created topic "blue".

- Produce some data in the created topic

```
/bin/kafka-console-producer --broker-list  
<Broker_1_IP_Address>:9092,<Broker_2_IP_Address>:9092,<Broker_3_IP_Address  
>:9092 --topic blue
```

Note: When prompted for data, enter some data for the topic and return to command prompt, by pressing Ctrl+Z.

- Consume the generated data

```
/bin/kafka-console-consumer --bootstrap-server  
<Broker_1_IP_Address>:9092,<Broker_2_IP_Address>:9092,<Broker_3_IP_Address  
>:9092 --topic blue --from-beginning
```

Output:

```
kafka  
message  
here  
^CProcessed a total of 3 messages.
```

- Open /etc/schema-registry/connect-avro-standalone.properties file (with sudo privilege), edit the following, save and exit the file

```
bootstrap.servers=<Broker_1_IP_Address>:9092,<Broker_2_IP_Address>:9092,<Bro  
ker_3_IP_Address>:9092
```

```
# The converters specify the format of data in Kafka and how to translate it into
Connect data.
# Every Connect user will need to configure these based on the format they want their
data in
# when loaded from or stored into Kafka
key.converter=org.apache.kafka.connect.storage.StringConverter
key.converter.schema.registry.url=http://<Schema_Registry_IP_Address>:8081
value.converter=org.apache.kafka.connect.storage.StringConverter
value.converter.schema.registry.url=http://<Schema_Registry_IP_Address>:8081
```

Note: Update the highlighted parameters with the values according to your Confluent Kafka cluster.

- Navigate to `/etc/kafka-connect-hdfs/quickstart-hdfs.properties` (with sudo privilege), update the HDFS URL and topic details, save and exit the file

```
topics=blue
hdfs.url=hdfs://<CDH_Controller_IP_Address>:8020
```

Note: Update the highlighted parameters.

- Run the HDFS Connector


```
sudo -u hdfs /bin/connect-standalone /etc/schema-registry/connect-avro-standalone.properties /etc/kafka-connect-hdfs/quickstart-hdfs.properties
```

Note: Kill any connector service, if running. Use `ps aux | grep "Connect"` to get the list of running connector service. Use `sudo kill -9 process_id` to kill the service.

Output:

```
[2019-06-06 04:59:09,286] INFO Committed hdfs://10.39.250.11:8020/topics/blue/partition=0/blue+0+0000000000+0000000002.avro for blue-0 (io.confluent.connect.hdfs.TopicPartitionWriter:782)
```

Note: The snapshot shows the path where the data is stored in HDFS.

- SSH into the HDFS cluster Controller node and look for the path


```
hdfs dfs -ls /topics/blue/partition=0
```

Output:

```
Found 1 items
-rw-r--r--  3 hdfs supergroup      108 2019-06-06 05:00
/topics/blue/partition=0/blue+0+0000000000+0000000002.avro
```

Note: Here `blue+0+0000000000+0000000002.avro` file contains the data sent from Kafka Connect node.

- Copy the file to `/tmp` directory

```
hdfs dfs -copyToLocal /topics/blue/partition=0/blue+0+0000000000+0000000002.avro
/tmp
```

- Convert the generated avro file to .json file

```
java -jar avro-tools-1.8.2.jar tojson /tmp/blue+0+0000000000+0000000002.avro >
kdata.json
```

Output:

```
log4j:WARN No appenders could be found for logger
(org.apache.hadoop.metrics2.lib.MutableMetricsFactory).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
```

Note: Ignore the Warning messages. Assuming, Avro tools is installed in the HDFS node.

- View the content of the generated .json file
cat kdata.json

Output:

```
{"string":"kafka"}
{"string":"message"}
{"string":"here"}
```

- From the Kafka Connect node, open the /etc/hadoop/conf.empty/core-site.xml file (with sudo privilege), add given content to it, save and exit the file

```
<configuration>
<property>
  <name>fs.s3a.access.key</name>
  <value></value>
</property>
<property>
  <name>fs.s3a.secret.key</name>
  <value></value>
</property>
<property>
  <name>fs.s3a.impl</name>
  <value>org.apache.hadoop.fs.s3a.S3AFileSystem</value>
</property>
<property>
  <name>fs.dtap.impl</name>
  <value>com.bluedata.hadoop.bdfs.Bdfs</value>
  <description>The FileSystem for BlueData dtap: URIs.</description>
</property>
</configuration>
```

- Open /etc/kafka-connect-hdfs/quickstart-hdfs.properties file (with sudo privilege), update the following details for DTAP connection, save and exit the file

```
topics=blue
hdfs.url=dtap://TenantStorage/kafka
hadoop.conf.dir=/etc/hadoop/conf.empty/
```


Note: The hdfs.url should be the path from EPIC DTAP.

- Run the executor

```
sudo -u hdfs /bin/connect-standalone /etc/schema-registry/connect-avro-standalone.properties /etc/kafka-connect-hdfs/quickstart-hdfs.properties
```

Output:

```
2019-06-06 06:14:02,415] INFO Opening record writer for: dtap://TenantStorage/message/topics//+tmp/blue/partition=0/00a398ce-fcce-4915-96e8-093038ac8069_tmp.avro (io.confluent.connect.hdfs.avro.AvroRecordWriterProvider:65)
```

- Verify in the EPIC UI

