

SMOKE TEST

Confluent Kafka 5

Date Prepared: Oct 2019

Document Information

Project Name	EPIC Accelerator Deployment & Integration Services		
Project Owner		Document Version No	0.1
Quality Review Method			
Prepared By	Priyanka	Preparation Date	Oct 2019
Reviewed By		Review Date	

Table of Contents

1 PREREQUISITES4

2 LAUNCH THE KSQL CLI5

3 INSPECT KAFKA TOPICS BY USING SHOW AND PRINT STATEMENTS6

4 CREATE A STREAM AND TABLE8

5 WRITE QUERIES10

6 KSQL REST API TESTING16

7 TESTING CONFLUENT KAFKA BROKER AND KAFKA ZOOKEEPRE SERVICES17

8 TESTING SCHEMA REGISTRY20

9 MONITOR AND MANAGE – CONTROL CENTER23

10 TESTING KAFKA CONNECT35

11 CONFIGURING CONFLUENT KAFKA WITH HDFS AND DTAP USING CONNECTORS
43

Table of Tables

NO TABLE OF FIGURES ENTRIES FOUND.

1 PREREQUISITES

In this section, we will see pre-requisites:

- This installation includes a Kafka broker, KSQL, Control Center, Zookeeper, Schema Registry, REST Proxy, and Kafka Connect.
- If you installed Confluent Platform via TAR or ZIP, navigate into the installation directory. The paths and commands used throughout this tutorial assume that you are in this installation directory.
- Java: Minimum version 1.8. Install Oracle Java JRE or JDK ≥ 1.8 on your local machine.

Create and produce data to the Kafka topics pageviews and users. These steps use the KSQL datagen that is included Confluent Platform.

- Create the pageviews topic and produce data using the data generator. The following example continuously generates data with a value in DELIMITED format.

```
<path-to-confluent>/bin/ksql-datagen quickstart=pageviews  
format=delimited topic=pageviews maxInterval=500
```

- Produce Kafka data to the user's topic using the data generator. The following example continuously generates data with a value in JSON format.

```
<path-to-confluent>/bin/ksql-datagen quickstart=users format=json  
topic=users maxInterval=100
```

2 LAUNCH THE KSQL CLI

After KSQL is started, your terminal should resemble this:

```
=====
=                                     =
=  [KSQL]                           =
=  <S>                               =
=  \S\                               =
=                                     =
=  Streaming SQL Engine for Apache Kafka® =
=====
```

Copyright 2018 Confluent Inc.

CLI v5.1.1, Server v5.1.1 located at <http://localhost:8088>

Having trouble? Type 'help' (case-insensitive) for a rundown of how things work!

ksql>

3 INSPECT KAFKA TOPICS BY USING SHOW AND PRINT STATEMENTS

KSQL enables inspecting Kafka topics and messages in real time.

- Use the SHOW TOPICS statement to list the available topics in the Kafka cluster.
- Use the PRINT statement to see a topic's messages as they arrive.

In the KSQL CLI, run the following statement:

```
SHOW TOPICS;
```

Output:

Kafka Topic	Registered	Partitions	Partition Replicas	Consumers ConsumerGroups
-----	-----	-----	-----	-----
_confluent-metrics	false	12	1	0 0
_schemas	false	1	1	0 0
pageviews	false	1	1	0 0
users	false	1	1	0 0
-----	-----	-----	-----	-----

1. Inspect the users topic by using the PRINT statement:

```
PRINT 'users';
```

Output:

Format:JSON

```
{ "ROWTIME":1540254230041,"ROWKEY":"User_1","registertime":1516754966866,"user
id":"User_1","regionid":"Region_9","gender":"MALE"}
{"ROWTIME":1540254230081,"ROWKEY":"User_3","registertime":1491558386780,"user
id":"User_3","regionid":"Region_2","gender":"MALE"}
{"ROWTIME":1540254230091,"ROWKEY":"User_7","registertime":1514374073235,"user
id":"User_7","regionid":"Region_2","gender":"OTHER"}
^C{"ROWTIME":1540254232442,"ROWKEY":"User_4","registertime":1510034151376,"u
serid":"User_4","regionid":"Region_8","gender":"FEMALE"}
Topic printing ceased
```

2. Inspect the pageviews topic by using the PRINT statement:

```
PRINT 'pageviews';
```

Output:

Format:STRING

10/23/18 12:24:03 AM UTC , 9461 , 1540254243183,User_9,Page_20

10/23/18 12:24:03 AM UTC , 9471 , 1540254243617,User_7,Page_47

10/23/18 12:24:03 AM UTC , 9481 , 1540254243888,User_4,Page_27


^C10/23/18 12:24:05 AM UTC , 9521 , 1540254245161,User_9,Page_62

Topic printing ceased

ksql>

4 CREATE A STREAM AND TABLE

These examples query messages from Kafka topics called pageviews and users using the following schemas:



1. Create a stream pageviews_original from the Kafka topic pageviews, specifying the value_format of DELIMITED.
2. Describe the new STREAM. Notice that KSQL created additional columns called ROWTIME, which corresponds to the Kafka message timestamp, and ROWKEY, which corresponds to the Kafka message key.

```
CREATE STREAM pageviews_original (viewtime bigint, userid  
varchar, pageid varchar) WITH \  
(kafka_topic='pageviews', value_format='DELIMITED');
```

Output:

Message

Stream created

3. Create a table users_original from the Kafka topic users, specifying the value_format of JSON.

```
CREATE TABLE users_original (registertime BIGINT, gender  
VARCHAR, regionid VARCHAR, userid VARCHAR) WITH \  
(kafka_topic='users', value_format='JSON', key = 'userid');
```


Output:

Message

Table created

4. Show all streams and tables.

```
ksql> SHOW STREAMS;
```

<i>Stream Name</i>	<i> Kafka Topic</i>	<i> Format</i>

<i>PAGEVIEWS_ORIGINAL</i>	<i> pageviews</i>	<i> DELIMITED</i>

```
ksql> SHOW TABLES;
```

<i>Table Name</i>	<i> Kafka Topic</i>	<i> Format</i>	<i> Windowed</i>

<i>USERS_ORIGINAL</i>	<i> users</i>	<i> JSON</i>	<i> false</i>

5 WRITE QUERIES

These examples write queries using KSQL.

Note: By default KSQL reads the topics for streams and tables from the latest offset.

1. Use SELECT to create a query that returns data from a STREAM. This query includes the LIMIT keyword to limit the number of rows returned in the query result. Note that exact data output may vary because of the randomness of the data generation.

```
SELECT pageid FROM pageviews_original LIMIT 3;
```

Output:

```
Page_24
Page_73
Page_78
LIMIT reached
Query terminated
```

2. Create a *persistent query* by using the CREATE STREAM keywords to precede the SELECT statement. The continual results from this query are written to the PAGEVIEWS_ENRICHED Kafka topic. The following query enriches the pageviews STREAM by doing a LEFT JOIN with the users_original TABLE on the user ID.

```
CREATE STREAM pageviews_enriched AS \
SELECT users_original.userid AS userid, pageid, regionid,
gender \
FROM pageviews_original \
LEFT JOIN users_original \
ON pageviews_original.userid = users_original.userid;
```

Output:

Message

```
-----
Stream created and running
-----
```

- Use SELECT to view query results as they come in. To stop viewing the query results, press <ctrl-c>. This stops printing to the console but it does not terminate the actual query. The query continues to run in the underlying KSQL application.

```
SELECT * FROM pageviews_enriched;
```

Output:

```
1519746861328 | User_4 | User_4 | Page_58 | Region_5 | OTHER
1519746861794 | User_9 | User_9 | Page_94 | Region_9 | MALE
1519746862164 | User_1 | User_1 | Page_90 | Region_7 | FEMALE
^CQuery terminated
```

- Create a new persistent query where a condition limits the streams content, using WHERE. Results from this query are written to a Kafka topic called PAGEVIEWS_FEMALE.

```
CREATE STREAM pageviews_female AS \
SELECT * FROM pageviews_enriched \
WHERE gender = 'FEMALE';
```

Output:

Message

```
-----
Stream created and running
-----
```

- Create a new persistent query where another condition is met, using LIKE. Results from this query are written to the **pageviews_enriched_r8_r9** Kafka topic.

```
CREATE STREAM pageviews_female_like_89 \
WITH (kafka_topic='pageviews_enriched_r8_r9') AS \
SELECT * FROM pageviews_female \
WHERE regionid LIKE '%_8' OR regionid LIKE '%_9';
```

Output:

Message

Stream created and running

6. Create a new persistent query that counts the pageviews for each region and gender combination in a tumbling window of 30 seconds when the count is greater than one. Results from this query are written to the PAGEVIEWS_REGIONS Kafka topic in the Avro format. KSQL will register the Avro schema with the configured Schema Registry when it writes the first message to the PAGEVIEWS_REGIONS topic.

```
CREATE TABLE pageviews_regions \  
  WITH (VALUE_FORMAT='avro') AS \  
SELECT gender, regionid , COUNT(*) AS numusers \  
FROM pageviews_enriched \  
  WINDOW TUMBLING (size 30 second) \  
GROUP BY gender, regionid \  
HAVING COUNT(*) > 1;
```

Output:

Message

Table created and running

7. **Optional:** View results from the above queries using **SELECT**.

```
SELECT gender, regionid, numusers FROM pageviews_regions LIMIT  
5;
```

Output:

FEMALE / Region_6 / 3

FEMALE / Region_1 / 4

FEMALE / Region_9 / 6

MALE / Region_8 / 2

OTHER / Region_5 / 4

LIMIT reached

Query terminated

ksql>

8. **Optional:** Show all persistent queries.

```
SHOW QUERIES;
```

Output:

Query ID	Kafka Topic	Query String
CSAS_PAGEVIEWS_FEMALE_1	PAGEVIEWS_FEMALE	CREATE STREAM pageviews_female AS SELECT * FROM pageviews_enriched WHERE gender = 'FEMALE';
CTAS_PAGEVIEWS_REGIONS_3	PAGEVIEWS_REGIONS	CREATE TABLE pageviews_regions WITH (VALUE_FORMAT='avro') AS SELECT gender, regionid , COUNT(*) AS numusers FROM pageviews_enriched WINDOW TUMBLING (size 30 second) GROUP BY gender, regionid HAVING COUNT(*) > 1;
CSAS_PAGEVIEWS_FEMALE_LIKE_89_2	PAGEVIEWS_FEMALE_LIKE_89	CREATE STREAM pageviews_female_like_89 WITH (kafka_topic='pageviews_enriched_r8_r9') AS SELECT * FROM pageviews_female WHERE regionid LIKE '%_8' OR regionid LIKE '%_9';
CSAS_PAGEVIEWS_ENRICHED_0	PAGEVIEWS_ENRICHED	CREATE STREAM pageviews_enriched AS SELECT users_original.userid AS userid, pageid, regionid, gender FROM pageviews_original LEFT JOIN users_original ON pageviews_original.userid = users_original.userid;

For detailed information on a Query run: EXPLAIN <Query ID>;

9. **Optional:** Examine query run-time metrics and details. Observe that information including the target Kafka topic is available, as well as throughput figures for the messages being processed.

```
DESCRIBE EXTENDED PAGEVIEWS_REGIONS;
```

Output:

Name : *PAGEVIEWS_REGIONS*
Type : *TABLE*
Key field : *KSQL_INTERNAL_COL_0/+|KSQL_INTERNAL_COL_1*
Key format : *STRING*
Timestamp field : *Not set - using <ROWTIME>*
Value format : *AVRO*
Kafka topic : *PAGEVIEWS_REGIONS (partitions: 4, replication: 1)*

Field | Type

```

-----
ROWTIME | BIGINT      (system)
ROWKEY  | VARCHAR(STRING) (system)
GENDER  | VARCHAR(STRING)
REGIONID | VARCHAR(STRING)
NUMUSERS | BIGINT
-----

```

Queries that write into this TABLE

```

-----
CTAS_PAGEVIEWS_REGIONS_3: CREATE TABLE pageviews_regions WITH
(value_format='avro') AS SELECT gender, regionid, COUNT(*) AS numusers
FROM pageviews_enriched WINDOW TUMBLING (size 30 second) GROUP
BY gender, regionid HAVING COUNT(*) > 1;

```

For query topology and execution plan please run: EXPLAIN <QueryId>

Local runtime statistics

***messages-per-sec: 3.06 total-messages: 1827 last-message: 7/19/18
4:17:55 PM UTC***

failed-messages: 0 failed-messages-per-sec: 0 last-failed: n/a

***(Statistics of the local KSQL server interaction with the Kafka topic
PAGEVIEWS_REGIONS)***

ksql>

6 KSQL REST API TESTING

1. GET the status of KSQL Server

```
curl -sX GET "http://172.18.0.25:8088/info";
```

Output:

```
{"KsqlServerInfo":{"version":"5.0.1","kafkaClusterId":"sNmrikrMSGGxZO3elbvQig",  
"ksqlServiceId":"default_"}}
```

2. Show all streams and tables using rest API

```
curl -X "POST" "http://172.18.0.25:8088/ksql" \  
-H "Content-Type: application/vnd.ksql.v1+json; charset=utf-8" \  
-d $'{  
  "ksql": "LIST STREAMS;",  
  "streamsProperties": {}  
'
```

Output:

```
[  
  {"@type":"streams","statementText":"LIST STREAMS;","streams":[  
    {"type":"STREAM","name":"PAGEVIEWS","topic":"_confluent-  
metrics","format":"DELIMITED"},  
    {"type":"STREAM","name":"TEST_STREAM","topic":"_confluent-  
metrics","format":"DELIMITED"}  
  ]  
}
```


7 TESTING CONFLUENT KAFKA BROKER AND KAFKA ZOOKEEPRE SERVICES

1. Create a Kafka Topic

To create a topic called *Test* in running cluster, use the following command:

```
/bin/kafka-topics --create --zookeeper  
172.18.0.26:2181,172.18.0.19:2181,172.18.0.23:2181 --  
replication-factor 1 --partitions 1 --topic Test
```

Output:

Created topic "Test".

Note: kafka_zookeeper service will be used to create Kafka Topic

2. List the Kafka topic

To list the topics in running cluster, use the following command:

```
/bin/kafka-topics --list --zookeeper  
172.18.0.26:2181,172.18.0.19:2181,172.18.0.23:2181
```

Output:

TEST-STREAMING

Test

__confluent.support.metrics

__consumer_offsets

__confluent-command

__confluent-ksql-default__command_topic

__confluent-metrics

__confluent-monitoring

__schemas

connect-configs

connect-offsets

connect-status

Note: This command returns the list with the names of all of the running topics in the cluster.

3. Produce the data in the Kafka topic

To create a Kafka Producer, run the following:

```
/bin/kafka-console-producer --broker-list  
172.18.0.21:9092,172.18.0.20:9092,172.18.0.27:9092 --topic Test
```

Now we can start sending the messages to the Kafka cluster from the console. The messages will be published to the Kafka Topic, "test".

Output:

>1, First Entry

>2, Second Entry

>3, Third Entry

>4, Fourth Entry

>5, Fifth Entry

Note: The output is just an example you can send message according to your Requirement.

4. Read data from a Kafka topic

To read data from Kafka topic, run the following command:

```
/bin/kafka-console-consumer --bootstrap-server  
172.18.0.20:9092,172.18.0.21:9092,172.18.0.27:9092 --topic Test  
--from-beginning
```

Start typing messages in the producer. Consumer would get the messages via Kafka Topic

Output:

1, First Entry

2, Second Entry

3, Third Entry

4, Fourth Entry

5, Fifth Entry

Processed a total of 5 messages

8 TESTING SCHEMA REGISTRY

1. Registering a New Version of a Schema under the subject "Kafka-key"

```
curl -X POST -H "Content-Type:
application/vnd.schemaregistry.v1+json" \
  --data '{"schema": "{\"type\": \"string\"}"}' \
  http://<Schema\_registry\_IP\_address>:8081/subjects/Kafka-key/versions
```

Output:

```
{"id":1}
```

2. Registering a New Version of a Schema under the subject "Kafka-value"

```
curl -X POST -H "Content-Type:
application/vnd.schemaregistry.v1+json" \
  --data '{"schema": "{\"type\": \"string\"}"}' \
  http://<Schema\_registry\_IP\_address>:8081/subjects/Kafka-value/versions
```

Output:

```
{"id":1}
```

3. Listing all Subjects

```
curl -X GET http://<Schema\_Registry\_IP\_address>:8081/subjects
```

Output:

```
["Kafka-value","Kafka-key"]
```

4. Fetching a Schema by Globally Unique ID

```
curl -X GET http://<Schema\_Registry\_IP\_address>:8081/schemas/ids/1
```

Output:

`{"schema": "\"string\""}"`

5. Listing all Schema Versions registered under the subject “Kafka-value”

```
curl -X GET http://<Schema\_registry\_IP\_address>:8081/subjects/Kafka-value/versions
```

Output:

`[1]`

6. Fetch Version 1 of the Schema registered under subject “Kafka-value”

```
curl -X GET http://<Schema\_Registry\_IP\_address>:8081/subjects/Kafka-value/versions/1
```

Output:

`{"subject": "Kafka-value", "version": 1, "id": 1, "schema": "\"string\""}"`

7. Checking if a Schema is registered under subject “Kafka-key”

```
curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" \
--data '{"schema": "{\"type\": \"string\"}"}' \
http://<Schema\_Registry\_IP\_address>:8081/subjects/Kafka-key
```

Output:

`{"subject": "Kafka-key", "version": 1, "id": 1, "schema": "\"string\""}"`

8. Testing compatibility of a Schema with the latest schema under subject “Kafka-value”

```
curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" \
--data '{"schema": "{\"type\": \"string\"}"}' \
```

```
http://<Schema\_Registry\_IP\_address>:8081/compatibility/subjects/Kafka-value/versions/latest
```

Output:

```
{"is_compatible":true}
```

9. Getting the Top level Config

```
curl -X GET http://<Schema\_Registry\_IP\_address>:8081/config
```

Output:

```
{"compatibilityLevel":"BACKWARD"}
```

10. Updating Compatibility requirements globally

```
curl -X PUT -H "Content-Type: application/vnd.schemaregistry.v1+json" \
  --data '{"compatibility": "NONE"}' \
  http://<Schema\_Registry\_IP\_address>:8081/config
```

Output:

```
{"compatibility":"NONE"}
```

11. Deleting all Schema Versions registered under subject “Kafka-value”

```
curl -X DELETE http://<Schema\_registry\_IP\_address>:8081/subjects/Kafka-value
```

Output:

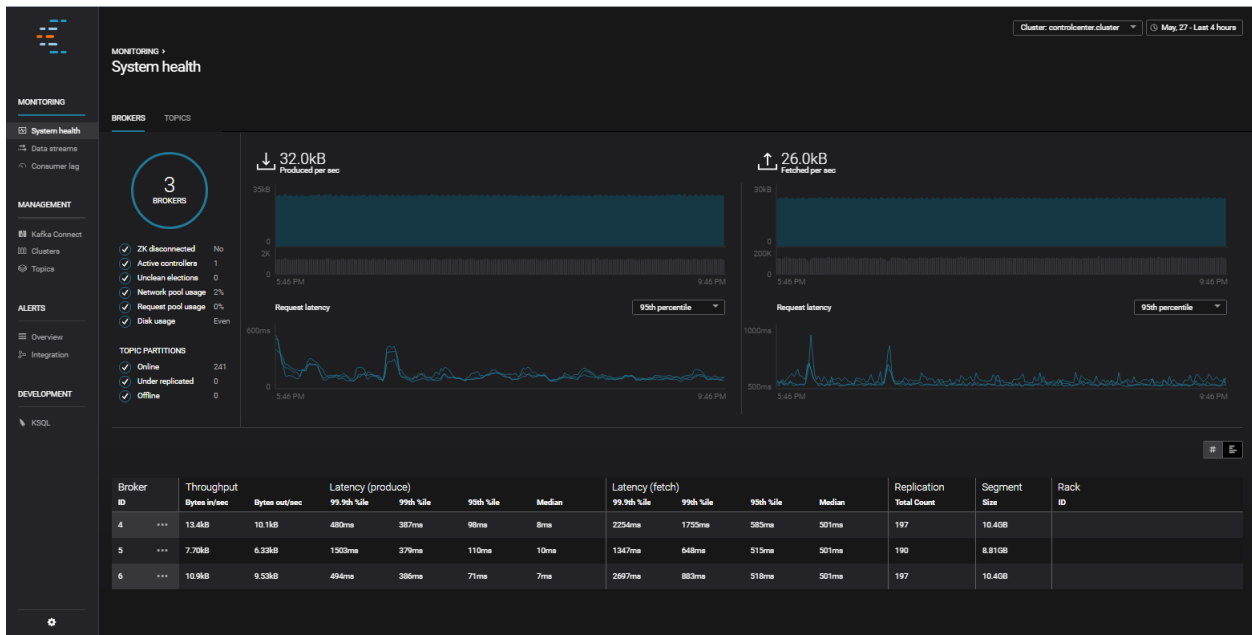
```
[1]
```

Smoke Test Document Confluent Kafka 5



9 MONITOR AND MANAGE – CONTROL CENTER

Open your web browser and navigate to <http://<IP Address>:9021> (by default). As you will enter you will see the Confluent dashboard, as shown below:



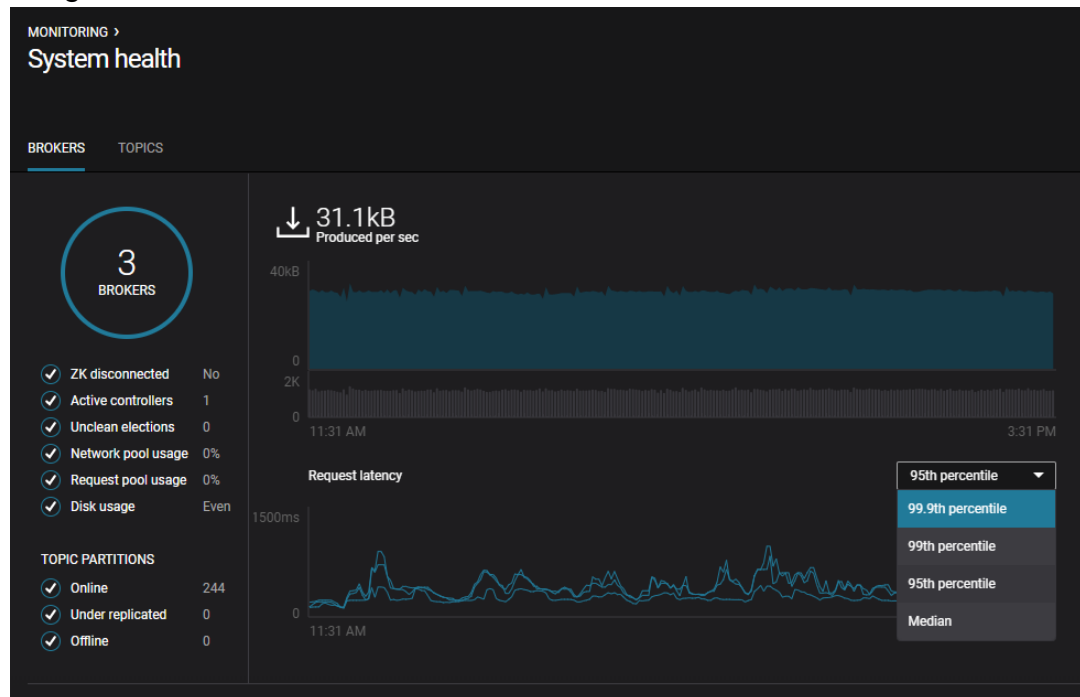
9.1.1 System Health

From the left-hand menu, click on **System Health**. System Health page will appear and it will display the health of Broker and Topic.

Broker tab:



- From the Request latency percentile drop-down list, select any percentile and view the change



- Hover your mouse over the charts and you can view details like Successful requests, Failed requests and Bytes produced/fetched per sec



Smoke Test Document Confluent Kafka 5



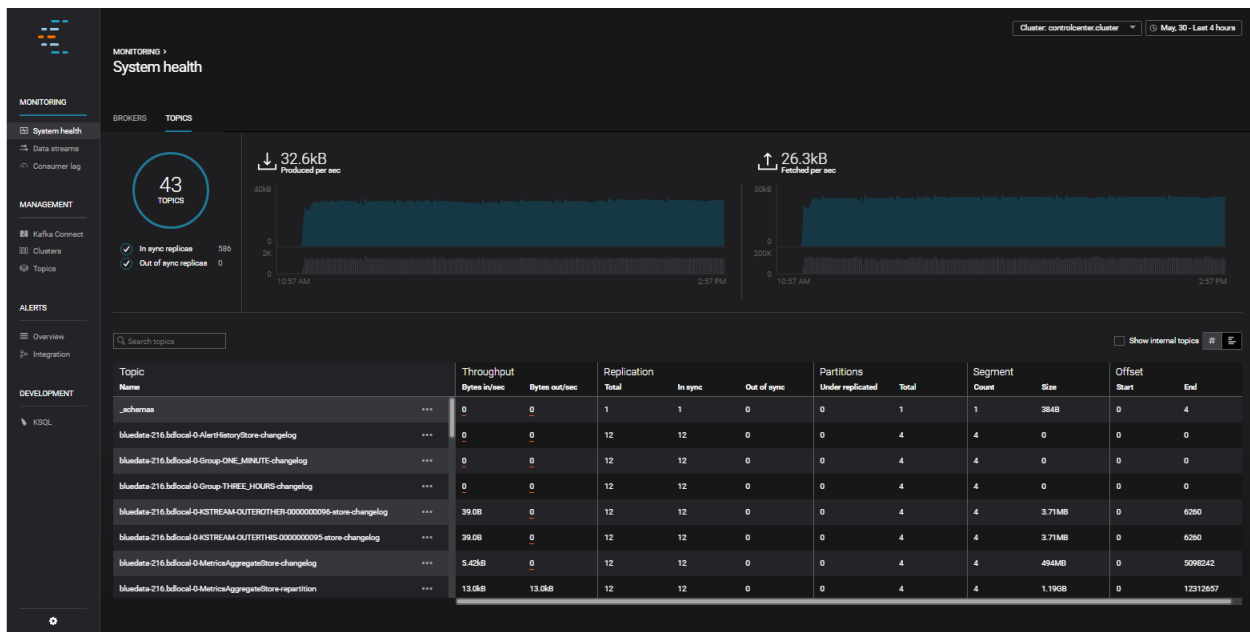
- Summary of the Produce and Fetch requests are given like below



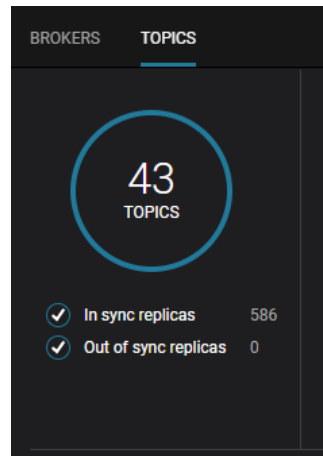
- At the bottom, you will see the Broker Metrics table

Broker ID	Throughput		Latency (produce)				Latency (fetch)				Replication Total Count	Segment Size	Rack ID
	Bytes in/sec	Bytes out/sec	99.9th %ile	99th %ile	95th %ile	Median	99.9th %ile	99th %ile	95th %ile	Median			
4	9.24kB	7.87kB	2163ms	1819ms	535ms	13ms	1893ms	1029ms	597ms	501ms	194	2.90GB	
5	12.6kB	11.4kB	1246ms	754ms	307ms	12ms	2321ms	2150ms	754ms	501ms	197	2.34GB	
6	10.2kB	7.31kB	2012ms	1728ms	586ms	16ms	2528ms	1046ms	616ms	501ms	195	2.36GB	

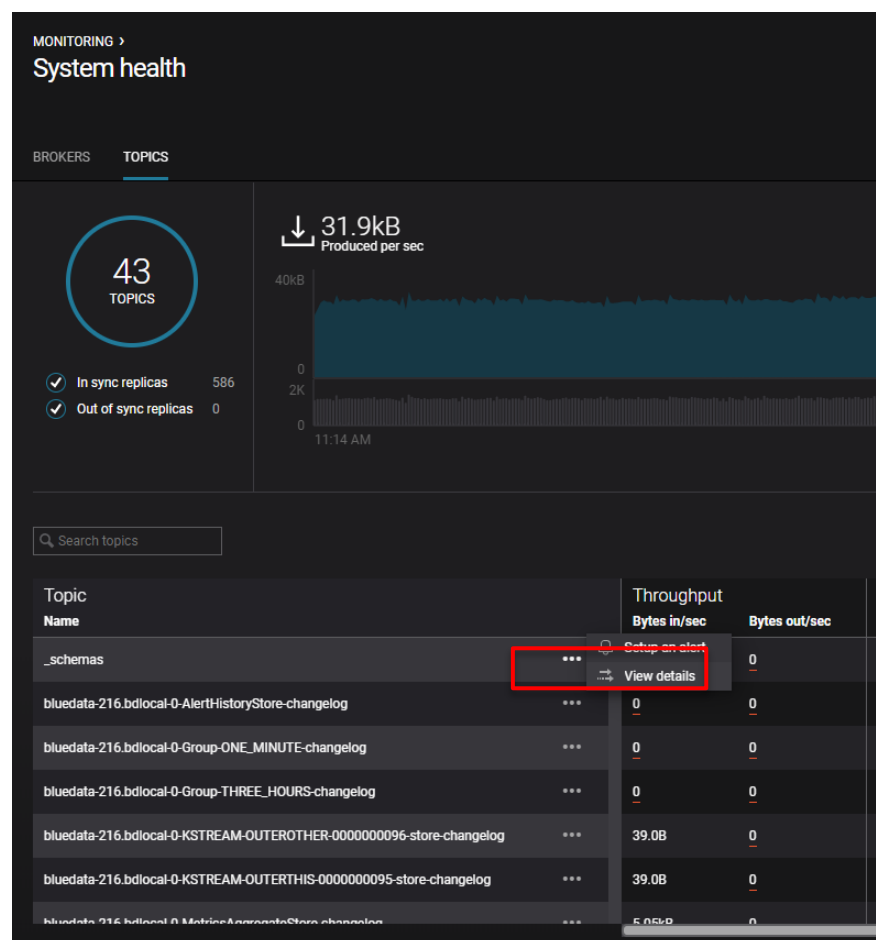
Topic tab:



- Topic Aggregate Metrics shows the total no. of topics, In sync replicas and Out of sync replicas



- Click on the ellipsis of any Topic, under Topics section and click on **View details**



- You will get screen like below, showing the In sync replica and Out of sync replica (if any)

MANAGEMENT > TOPICS > **bluedata-216.bdlocal-0-Group-ONE_MINUTE-changelog**

STATUS SCHEMA INSPECT SETTINGS

■ In sync replica ■ Out of sync replica

Partitions	Replicas	Replica placement
id		Broker list
0	■ ■ ■	5,4,6 ■ ■ ■
1	■ ■ ■	4,6,5 ■ ■ ■
2	■ ■ ■	6,5,4 ■ ■ ■
3	■ ■ ■	5,6,4 ■ ■ ■

- Click on **SETTINGS**, to **Edit settings** or **Show full config** button, as per requirement

MANAGEMENT > TOPICS > **bluedata-216.bdlocal-0-Group-ONE_MINUTE-changelog**

STATUS SCHEMA INSPECT **SETTINGS**

min.insync.replicas	2
cleanup.policy	compact,delete
retention.ms	432000000
retention.bytes	-1

Edit settings Show full config

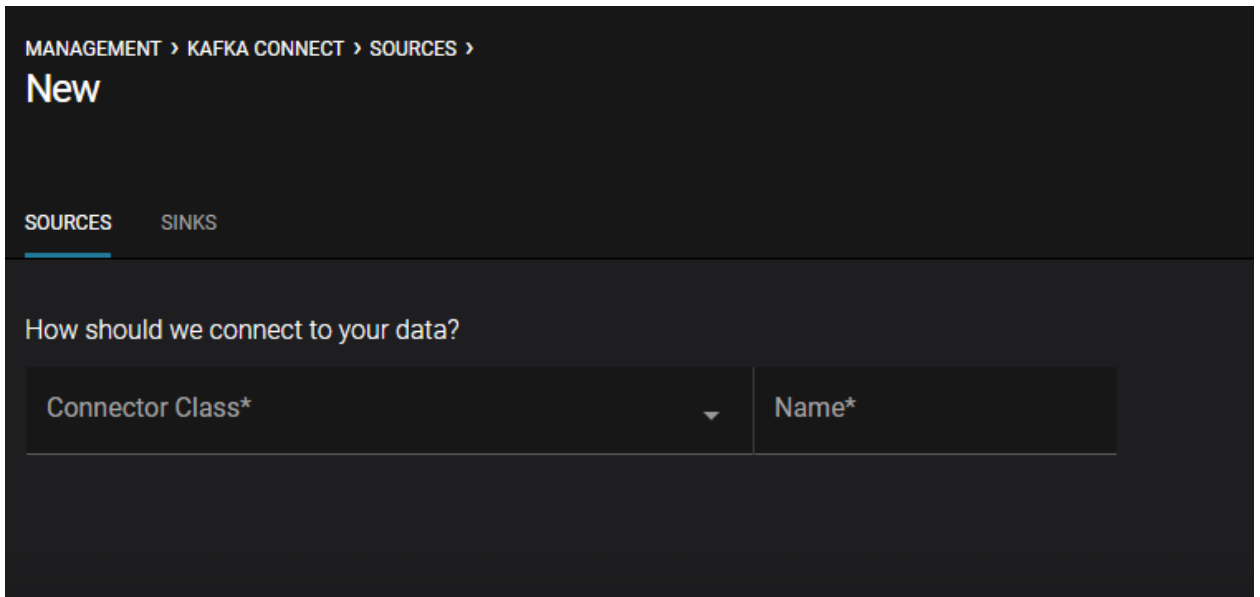
- At the bottom, you will see Topic Metric table

Topic Name	Throughput		Replication			Partitions		Segment		Offset	
	Bytes in/sec	Bytes out/sec	Total	In sync	Out of sync	Under replicated	Total	Count	Size	Start	End
bluedata-216.bdlocal-0-AlertHistoryStore-changelog	0	0	12	12	0	0	4	4	0	0	0
bluedata-216.bdlocal-0-Group-ONE_MINUTE-changelog	0	0	12	12	0	0	4	4	0	0	0
bluedata-216.bdlocal-0-Group-THREE_HOURS-changelog	0	0	12	12	0	0	4	4	0	0	0
bluedata-216.bdlocal-0-KSTREAM-OUTEROTHER-0000000096-store-changelog	39.08	0	12	12	0	0	4	4	3.83MB	0	6451
bluedata-216.bdlocal-0-KSTREAM-OUTERTHIS-0000000095-store-changelog	39.08	0	12	12	0	0	4	4	3.83MB	0	6451
bluedata-216.bdlocal-0-MetricsAggregateStore-changelog	4.31KB	0	12	12	0	0	4	4	509MB	0	5253713

- Likewise explore other options to get more information

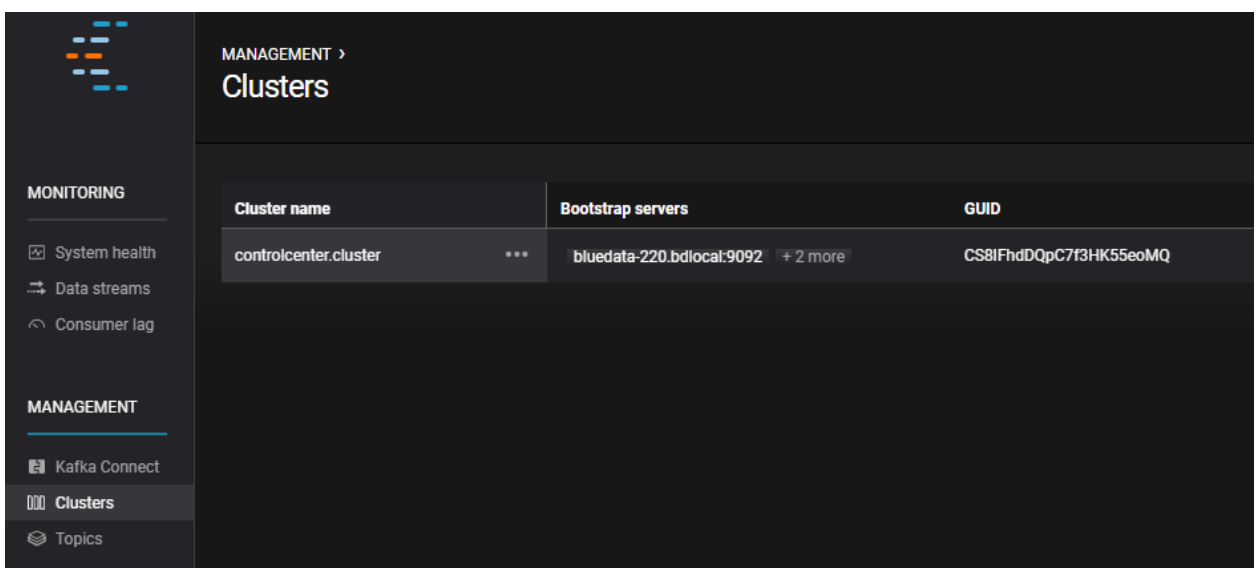
9.1.2 Kafka Connect

- From the left-hand menu, under **MANAGEMENT** section, click on **Kafka Connect**
- To create a source, provide the value of following parameters and click on **Submit**



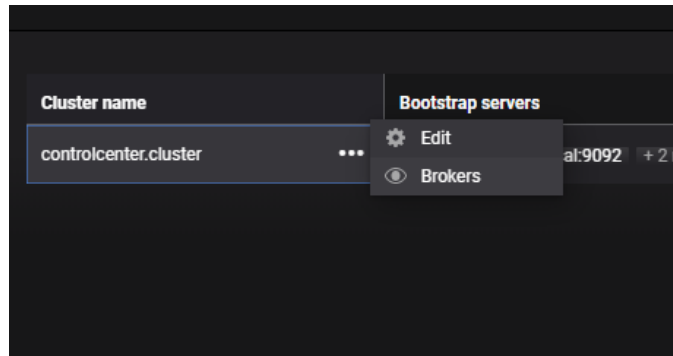
9.1.3 Clusters

- From the left-hand menu, under **MANAGEMENT** section, click on **Clusters**

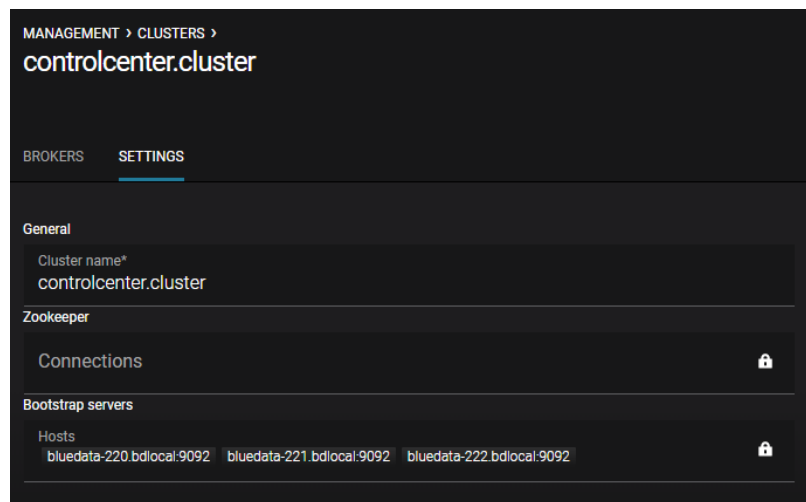


Cluster name	Bootstrap servers	GUID
controlcenter.cluster	bluedata-220.bdlocal:9092 + 2 more	CS8lFhdDQpC7f3HK55eoMQ

- Click on ellipsis of the cluster name to **Edit** or **Brokers** button as per requirement



- The Edit window looks like this below



- The Broker information is listed as below, it can also be downloaded

MANAGEMENT > CLUSTERS >
controlcenter.cluster

BROKERS SETTINGS

☒ broker.id 4
☒ broker.id 5
☒ broker.id 6

GROUP MANAGEMENT

group.initial.rebalance.delay.ms	0
----------------------------------	---

LOG

num.recovery.threads.per.data.dir	3
offsets.topic.replication.factor	2
log.retention.hours	168
num.partitions	3
log.segment.bytes	1073741824
log.dirs	/var/log/kafka
log.retention.check.interval.ms	300000

SECURITY

socket.request.max.bytes	104857600
socket.send.buffer.bytes	102400
socket.receive.buffer.bytes	102400

SOCKET SERVER

metric.reporters	io.confluent.metrics.reporter.ConfluentMetricsReporter
num.io.threads	8
num.network.threads	3

TRANSACTION

transaction.state.log.replication.factor	1
zookeeper.connect	bluedata-217.bdlocal:2181,bluedata-218.bdlocal:2181,bluedata-219.bdlocal:2181
transaction.state.log.min.isr	1
zookeeper.connection.timeout.ms	6000

OTHER

9.1.4 Topics

- From the left-hand menu, under **MANAGEMENT** section, click on **Topics**

MANAGEMENT >

Topics

☐ Show internal topics

Topics		Partitions	Replication		Data flow
Name		Total	Factor	% In sync	Consumer Groups
_schemas	...	1	x1	100%	0
bluedata-216.bdlocal-0-AlertHistoryStore-changelog	...	4	x3	100%	0
bluedata-216.bdlocal-0-Group-ONE_MINUTE-changelog	...	4	x3	100%	0
bluedata-216.bdlocal-0-Group-THREE_HOURS-changelog	...	4	x3	100%	0
bluedata-216.bdlocal-0-KSTREAM-OUTEROTHER-0000000096-store-changelog	...	4	x3	100%	0
bluedata-216.bdlocal-0-KSTREAM-OUTERTHIS-0000000095-store-changelog	...	4	x3	100%	0
bluedata-216.bdlocal-0-MetricsAggregateStore-changelog	...	4	x3	100%	0
bluedata-216.bdlocal-0-MetricsAggregateStore-repartition	...	4	x3	100%	0
bluedata-216.bdlocal-0-MonitoringMessageAggregatorWindows-ONE_MINUTE-chang...	...	4	x3	100%	0
bluedata-216.bdlocal-0-MonitoringMessageAggregatorWindows-THREE_HOURS-cha...	...	4	x3	100%	0
bluedata-216.bdlocal-0-MonitoringStream-ONE_MINUTE-changelog	...	4	x3	100%	0
bluedata-216.bdlocal-0-MonitoringStream-THREE_HOURS-changelog	...	4	x3	100%	0
bluedata-216.bdlocal-0-MonitoringTriggerStore-changelog	...	4	x3	100%	0
bluedata-216.bdlocal-0-MonitoringVerifierStore-changelog	...	4	x3	100%	0
bluedata-216.bdlocal-0-TriggerActionsStore-changelog	...	4	x3	100%	0

- Check **Show internal topics**, to view all the internal topics

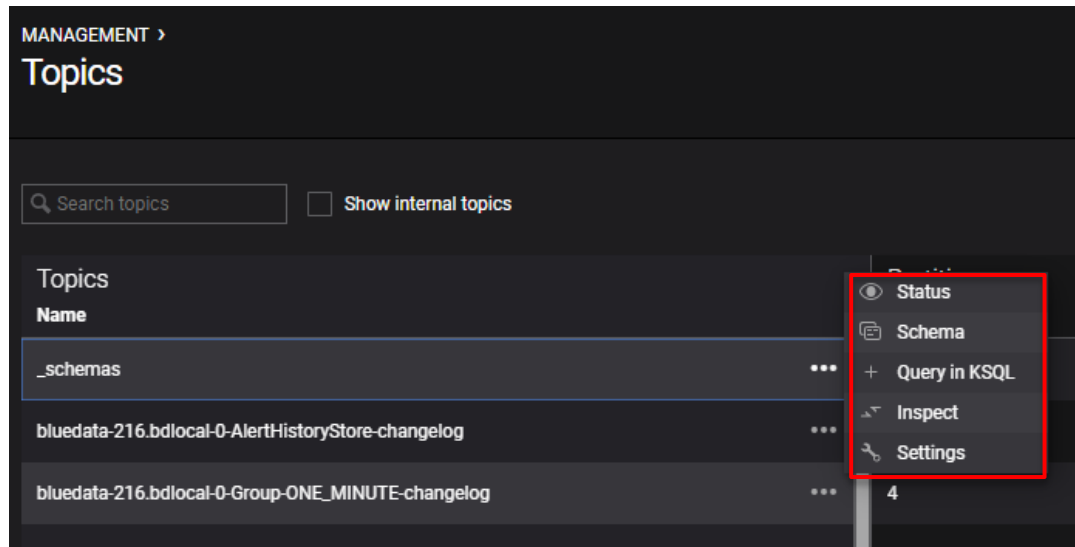
MANAGEMENT >

Topics


☒ Show internal topics

Topics		Partitions
Name		Total
__confluent.support.metrics	...	1
__confluent-command	...	1

- Click on the topics ellipsis to view the multiple options related to that topic



- To create a new Topic, click on **Create topic**, at right-hand side



- Provide **Topic name** and **Number of partitions** for Topic creation and click on **Create with defaults** or **Customize settings**, according to requirement

MANAGEMENT >
Topics

Topic name*
Testing

Number of partitions* 3

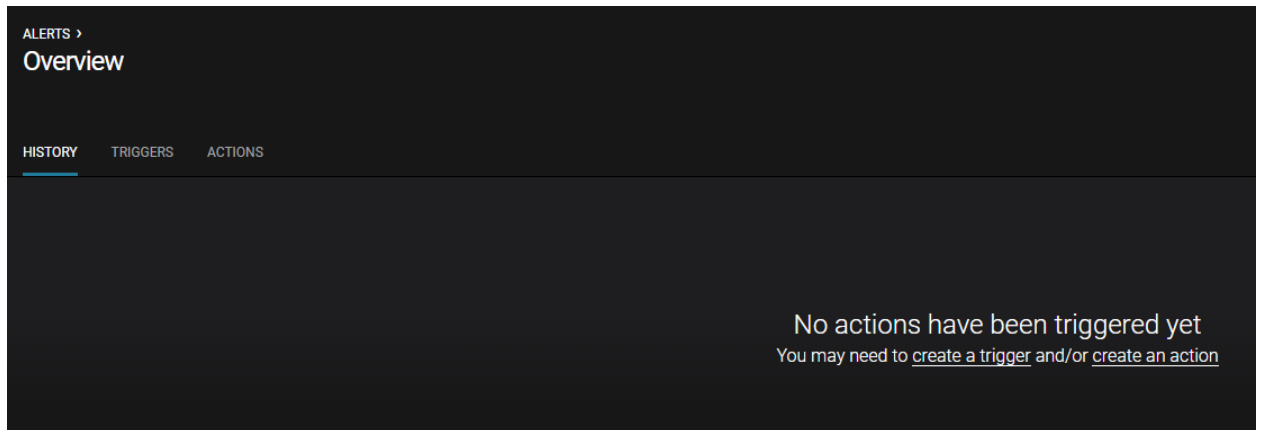
TOPIC SUMMARY

cleanup.policy	delete
min.insync.replicas	1
replication.factor	1
max.message.bytes	1000011
retention.bytes	-1
retention.ms	-
partitions	3
name	Testing
cluster	confluent-enterprise

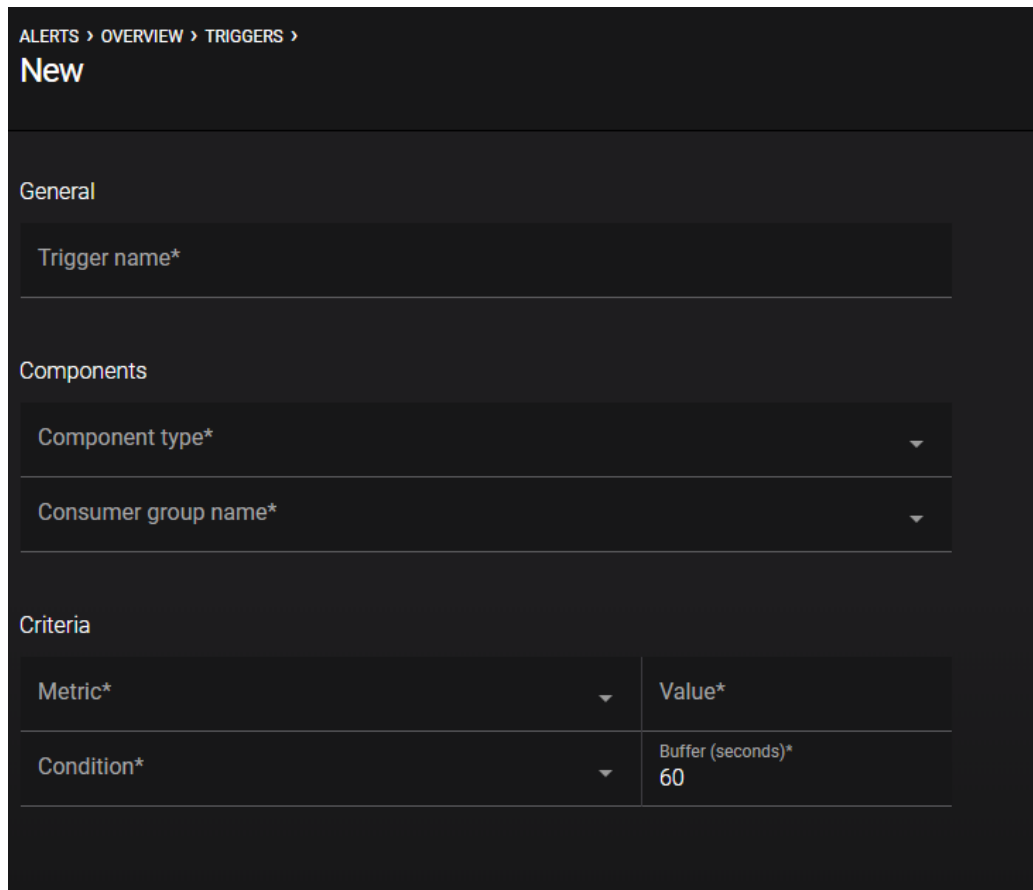
Create with defaults Customize settings Cancel

9.1.5 Alerts

- From the left-hand menu, under **ALERTS** section, click on **Overview**

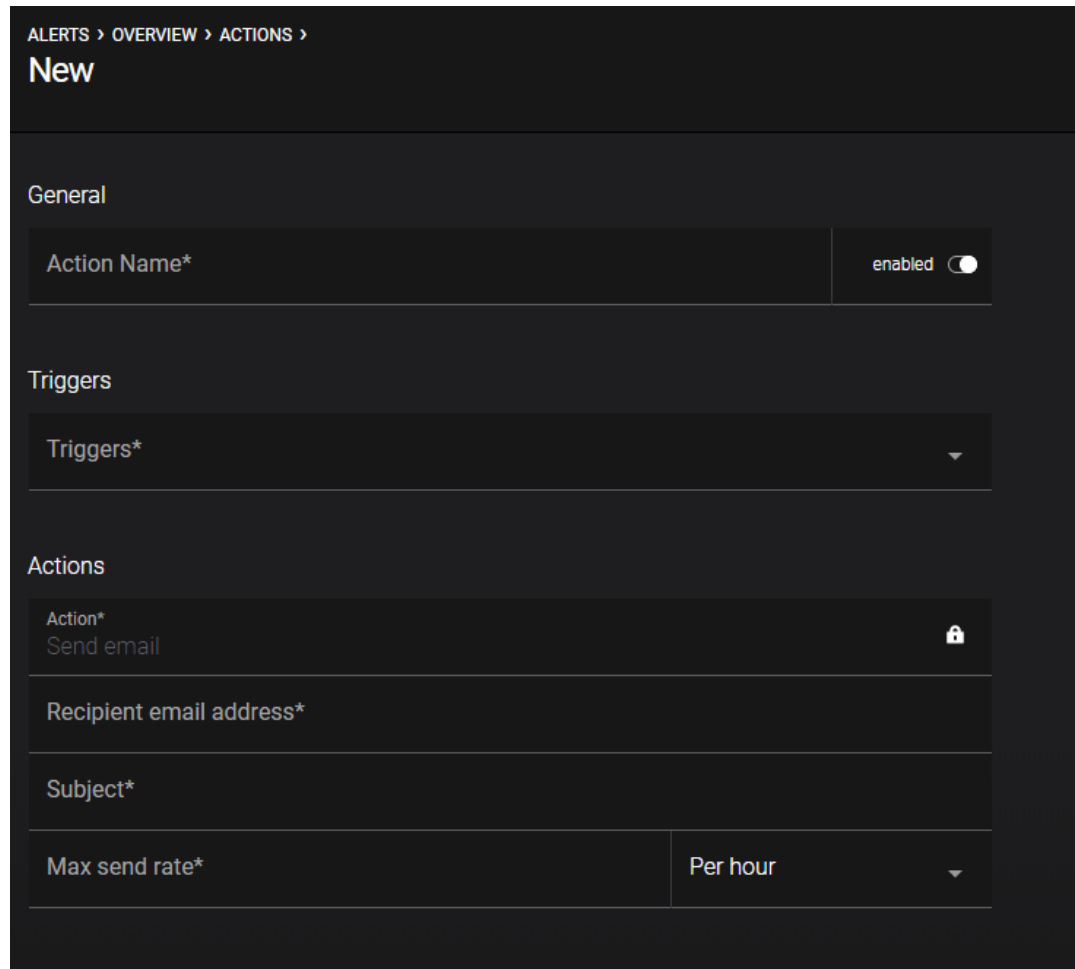


- It allows you to create a Trigger and/or an Action
- To create a trigger, provide the values of the following parameters and click on **Submit**



Criteria	
Metric*	Value*
Condition*	Buffer (seconds)* 60

- To create an Action, provide the value for the following parameters and click on **Submit**



ALERTS > OVERVIEW > ACTIONS > **New**

General

Action Name* enabled ☐

Triggers

Triggers* ▼

Actions

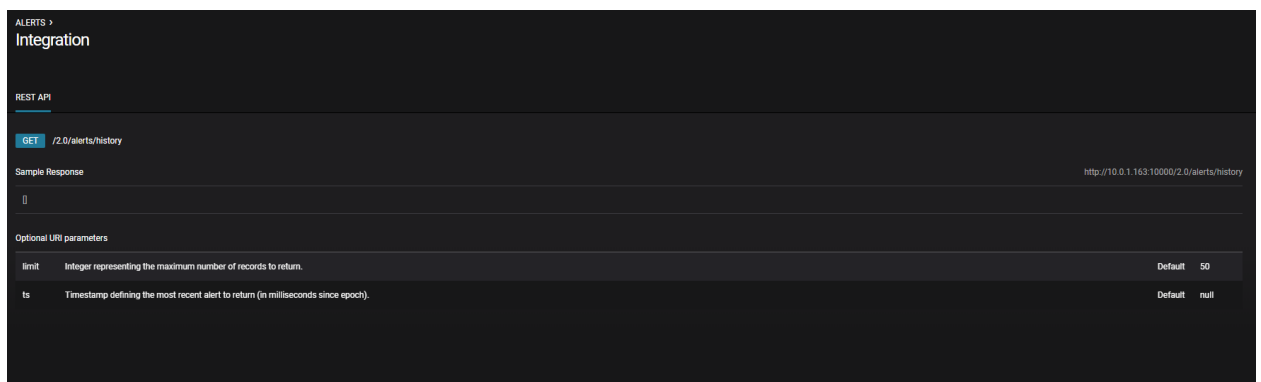
Action* 🔒
Send email

Recipient email address*

Subject*

Max send rate* Per hour ▼

- Click on **Integration**, under the ALERTS section, it provides REST API endpoints



ALERTS > **Integration**

REST API

GET /2.0/alerts/history

Sample Response http://10.0.1.163:10000/2.0/alerts/history

```
[ ]
```

Optional URI parameters

Parameter	Description	Default
limit	Integer representing the maximum number of records to return.	50
ts	Timestamp defining the most recent alert to return (in milliseconds since epoch).	null

10 TESTING KAFKA CONNECT

1. Get a list of active connectors (Initially empty)

```
curl -X GET http://<Kafka\_Connect\_IP\_Address>:8083/connectors
```

Output:

```
[]
```

2. Create a new connector
 - a. HDFS Sink Connector

```
curl -H "Content-Type: application/json" \
--request POST \
--data '{
  "name": "hdfs-sink-connector",
  "config": {
    "connector.class":
"io.confluent.connect.hdfs.HdfsSinkConnector",
    "tasks.max": "10",
    "topics": "test-topic",
    "hdfs.url": "hdfs://<IP_Address>:9000",
    "hadoop.conf.dir": "/opt/hadoop/conf",
    "hadoop.home": "/opt/hadoop",
    "flush.size": "100",
    "rotate.interval.ms": "1000"
  }
}' \
http://<Kafka\_Connect\_IP\_Address>:8083/connectors
```

Output:

```
{"name":"hdfs-sink-connector","config":{"connector.class":"io.confluent.connect.hdfs.HdfsS
```

```
inkConnector","tasks.max":"10","topics":"test-  
topic","hdfs.url":"hdfs://10.39.250.64:9000","hadoop.conf.dir":"/opt/hado  
op/conf","hadoop.home":"/opt/hadoop","flush.size":"100","rotate.interval  
.ms":"1000","name":"hdfs-sink-connector"},"tasks":[{"connector":"hdfs-  
sink-connector","task":0},{"connector":"hdfs-sink-  
connector","task":1},{"connector":"hdfs-sink-  
connector","task":2},{"connector":"hdfs-sink-  
connector","task":3},{"connector":"hdfs-sink-  
connector","task":4},{"connector":"hdfs-sink-  
connector","task":5},{"connector":"hdfs-sink-  
connector","task":6},{"connector":"hdfs-sink-  
connector","task":7},{"connector":"hdfs-sink-  
connector","task":8},{"connector":"hdfs-sink-  
connector","task":9}], "type":"sink"}
```

b. Local File Sink

```
curl -H "Content-Type: application/json" \  
--request POST \  
--data '{  
  "name": "local-file-sink",  
  "config": {  
    "connector.class": "FileStreamSinkConnector",  
    "tasks.max": "2",  
    "topics": "connect-test",  
    "hdfs.url": "hdfs://<IP_Address>:9000",  
    "hadoop.conf.dir": "/opt/hadoop/conf",  
    "hadoop.home": "/opt/hadoop",  
    "flush.size": "100",  
    "rotate.interval.ms": "1000"  
  }  
}' \  
http://<Kafka Connect IP Address>:8083/connectors
```

Output:

```
{"name":"local-file-sink","config":{"connector.class":"FileStreamSinkConnector","tasks.max":"2","topics":"connect-test","hdfs.url":"hdfs://10.39.250.16:9000","hadoop.conf.dir":"/opt/hadoop/conf","hadoop.home":"/opt/hadoop","flush.size":"100","rotate.interval.ms":"1000","name":"local-file-sink"},"tasks":[],"type":"sink"}
```

3. Verify the connector is added by executing

```
curl -X GET http://<Kafka\_Connect\_IP\_Address>:8083/connectors
```

Output:

```
["local-file-sink","hdfs-sink-connector"]
```

Note: **hdfs-sink-connector** and **local-file-sink** are the connector name here, which will be used in later commands. Replace it with any other string name (connector name) according to your environment.

4. Get connector status

```
curl -X GET http://<Kafka\_Connect\_IP\_Address>:8083/connectors/local-file-sink/status
```

Output:

```
{"name":"local-file-sink","connector":{"state":"RUNNING","worker_id":"bluedata-4424.bdlocal:8083"},"tasks":[{"id":0,"state":"RUNNING","worker_id":"bluedata-4424.bdlocal:8083"},{"id":1,"state":"RUNNING","worker_id":"bluedata-4424.bdlocal:8083"}],"type":"sink"}
```

5. Get worker's version

```
curl -X GET http://<Kafka_Connect_IP_Address>:8083/
```

Output:

```
{"version":"2.0.1-  
cp4","commit":"49da0fef3e389dc2","kafka_cluster_id":"CS8IFhdDQpC7f3HK55  
eoMQ"}
```

6. Get the list of connector plugins available on the worker

```
curl -X GET http://<Kafka_Connect_IP_Address>:8083/connector-plugins
```

Output:

```
[{"class":"io.confluent.connect.activemq.ActiveMQSourceConnector","type":"s  
ource","version":"5.0.3"}, {"class":"io.confluent.connect.elasticsearch.Elasticse  
archSinkConnector","type":"sink","version":"5.0.3"}, {"class":"io.confluent.con  
nect.hdfs.HdfsSinkConnector","type":"sink","version":"5.0.3"}, {"class":"io.conf  
luent.connect.hdfs.tools.SchemaSourceConnector","type":"source","version":"  
2.0.1-  
cp4"}, {"class":"io.confluent.connect.ibm.mq.IbmMQSourceConnector","type":"  
source","version":"5.0.3"}, {"class":"io.confluent.connect.jdbc.JdbcSinkConnec  
tor","type":"sink","version":"5.0.3"}, {"class":"io.confluent.connect.jdbc.JdbcSo  
urceConnector","type":"source","version":"5.0.3"}, {"class":"io.confluent.conne  
ct.jms.JmsSourceConnector","type":"source","version":"5.0.3"}, {"class":"io.co  
nfluent.connect.replicator.ReplicatorSourceConnector","type":"source","versio  
n":"5.0.3"}, {"class":"io.confluent.connect.s3.S3SinkConnector","type":"sink","v  
ersion":"5.0.3"}, {"class":"io.confluent.connect.storage.tools.SchemaSourceCo  
nnector","type":"source","version":"2.0.1-  
cp4"}, {"class":"org.apache.kafka.connect.file.FileStreamSinkConnector","type"  
:"sink","version":"2.0.1-
```

```
cp4"},"{"class":"org.apache.kafka.connect.file.FileStreamSourceConnector","type":"source","version":"2.0.1-cp4"}]]
```

7. Get tasks for a connector

```
curl -X GET http://<Kafka_Connect_IP_Address>:8083/connectors/hdfs-sink-connector/tasks
```

Output:

```
[{"id":{"connector":"hdfs-sink-connector","task":0},"config":{"connector.class":"io.confluent.connect.hdfs.HdfsSinkConnector","task.class":"io.confluent.connect.hdfs.HdfsSinkTask","hadoop.conf.dir":"/opt/hadoop/conf","flush.size":"100","tasks.max":"10","topics":"test-topic","hdfs.url":"hdfs://10.39.250.64:9000","name":"hdfs-sink-connector","rotate.interval.ms":"1000","hadoop.home":"/opt/hadoop"}}, {"id":{"connector":"hdfs-sink-connector","task":1},"config":{"connector.class":"io.confluent.connect.hdfs.HdfsSinkConnector","task.class":"io.confluent.connect.hdfs.HdfsSinkTask","hadoop.conf.dir":"/opt/hadoop/conf","flush.size":"100","tasks.max":"10","topics":"test-topic","hdfs.url":"hdfs://10.39.250.64:9000","name":"hdfs-sink-connector","rotate.interval.ms":"1000","hadoop.home":"/opt/hadoop"}}, {"id":{"connector":"hdfs-sink-connector","task":2},"config":{"connector.class":"io.confluent.connect.hdfs.HdfsSinkConnector","task.class":"io.confluent.connect.hdfs.HdfsSinkTask","hadoop.conf.dir":"/opt/hadoop/conf","flush.size":"100","tasks.max":"10","topics":"test-topic","hdfs.url":"hdfs://10.39.250.64:9000","name":"hdfs-sink-connector","rotate.interval.ms":"1000","hadoop.home":"/opt/hadoop"}}, {"id":{"connector":"hdfs-sink-connector","task":3},"config":{"connector.class":"io.confluent.connect.hdfs.HdfsSinkConnector","task.class":"io.confluent.connect.hdfs.HdfsSinkTask","hadoop.conf.dir":"/opt/hadoop/conf","flush.size":"100","tasks.max":"10","topics":"test-topic","hdfs.url":"hdfs://10.39.250.64:9000","name":"hdfs-sink-
```

```
connector","rotate.interval.ms":"1000","hadoop.home":"/opt/hadoop"}},{  
connector":"hdfs-sink-  
connector","task":4,"config":{"connector.class":"io.confluent.connect.hdfs.Hd  
fsSinkConnector","task.class":"io.confluent.connect.hdfs.HdfsSinkTask","hado  
op.conf.dir":"/opt/hadoop/conf","flush.size":"100","tasks.max":"10","topics":"t  
est-topic","hdfs.url":"hdfs://10.39.250.64:9000","name":"hdfs-sink-  
connector","rotate.interval.ms":"1000","hadoop.home":"/opt/hadoop"}},{  
connector":"hdfs-sink-  
connector","task":5,"config":{"connector.class":"io.confluent.connect.hdfs.Hd  
fsSinkConnector","task.class":"io.confluent.connect.hdfs.HdfsSinkTask","hado  
op.conf.dir":"/opt/hadoop/conf","flush.size":"100","tasks.max":"10","topics":"t  
est-topic","hdfs.url":"hdfs://10.39.250.64:9000","name":"hdfs-sink-  
connector","rotate.interval.ms":"1000","hadoop.home":"/opt/hadoop"}},{  
connector":"hdfs-sink-  
connector","task":6,"config":{"connector.class":"io.confluent.connect.hdfs.Hd  
fsSinkConnector","task.class":"io.confluent.connect.hdfs.HdfsSinkTask","hado  
op.conf.dir":"/opt/hadoop/conf","flush.size":"100","tasks.max":"10","topics":"t  
est-topic","hdfs.url":"hdfs://10.39.250.64:9000","name":"hdfs-sink-  
connector","rotate.interval.ms":"1000","hadoop.home":"/opt/hadoop"}},{  
connector":"hdfs-sink-  
connector","task":7,"config":{"connector.class":"io.confluent.connect.hdfs.Hd  
fsSinkConnector","task.class":"io.confluent.connect.hdfs.HdfsSinkTask","hado  
op.conf.dir":"/opt/hadoop/conf","flush.size":"100","tasks.max":"10","topics":"t  
est-topic","hdfs.url":"hdfs://10.39.250.64:9000","name":"hdfs-sink-  
connector","rotate.interval.ms":"1000","hadoop.home":"/opt/hadoop"}},{  
connector":"hdfs-sink-  
connector","task":8,"config":{"connector.class":"io.confluent.connect.hdfs.Hd  
fsSinkConnector","task.class":"io.confluent.connect.hdfs.HdfsSinkTask","hado  
op.conf.dir":"/opt/hadoop/conf","flush.size":"100","tasks.max":"10","topics":"t  
est-topic","hdfs.url":"hdfs://10.39.250.64:9000","name":"hdfs-sink-  
connector","rotate.interval.ms":"1000","hadoop.home":"/opt/hadoop"}},{  
connector":"hdfs-sink-
```



```
connector","task":9},"config":{"connector.class":"io.confluent.connect.hdfs.HdfsSinkConnector","task.class":"io.confluent.connect.hdfs.HdfsSinkTask","hadoop.conf.dir":"/opt/hadoop/conf","flush.size":"100","tasks.max":"10","topics":["test-topic"],"hdfs.url":"hdfs://10.39.250.64:9000","name":"hdfs-sink-connector","rotate.interval.ms":"1000","hadoop.home":"/opt/hadoop"}]]
```

8. Restart connector and its tasks

```
curl -X POST http://<Kafka\_Connect\_IP\_Address>:8083/connectors/hdfs-sink-connector/restart
```

Output:

<No response when successful>

9. Pause the connector and its tasks

```
curl -X PUT http://<Kafka\_Connect\_IP\_Address>:8083/connectors/hdfs-sink-connector/pause
```

Output:

<No response when successful>

10. Resume the connector and its tasks

```
curl -X PUT http://<Kafka\_Connect\_IP\_Address>:8083/connectors/hdfs-sink-connector/resume
```

Output:

<No response when successful>

11. Restart an individual task

```
curl -X POST http://<Kafka\_Connect\_IP\_Address>:8083/connectors/hdfs-sink-connector/tasks/0/restart
```

Output:

<No response when successful>

12. Delete a connector

```
curl -X DELETE http://<Kafka Connect IP Address>:8083/connectors/hdfs-sink-connector
```

Output:

<No response when successful>

11 CONFIGURING CONFLUENT KAFKA WITH HDFS AND DTAP USING CONNECTORS

Here in this section, we will see how to receive message from a Kafka Topic and store it into HDFS and DTAP, using connector.

1. SSH into Kafka connect node

Note: It is assumed that Cloudera Manager Repo is installed in the `/etc/yum.repos.d/` directory and Hadoop-client is installed on Kafka connect node.

2. Create a Kafka Topic

```
/bin/kafka-topics --create --zookeeper  
<Zookeeper_1_IP_Address>,<Zookeeper_2_IP_Address>,<Zookeeper_3_  
IP_Address>:2181 --replication-factor 1 --partitions 1 --topic  
blue
```

Output:

Created topic "blue".

3. Produce some data in the created Topic

```
/bin/kafka-console-producer --broker-list  
<Broker_1_IP_Address>:9092,<Broker_2_IP_Address>:9092,<Broker_3_  
_IP_Address>:9092 --topic blue
```

Note: When prompted for data, enter some data for the Topic and return to command prompt, by pressing **Ctrl+Z**.

4. Consume the generated data

```
/bin/kafka-console-consumer --bootstrap-server  
<Broker_1_IP_Address>:9092,<Broker_2_IP_Address>:9092,<Broker_3_  
_IP_Address>:9092 --topic blue --from-beginning
```

Output:

kafka

message

here

^CProcessed a total of 3 messages.

5. Open **/etc/schema-registry/connect-avro-standalone.properties** file (with sudo privilege), edit the following, save and exit the file

```
bootstrap.servers=<Broker_1_IP_Address>:9092,<Broker_2_IP_Addre
ss>:9092,<Broker_3_IP_Address>:9092
# The converters specify the format of data in Kafka and how to
translate it into Connect data.
# Every Connect user will need to configure these based on the
format they want their data in
# when loaded from or stored into Kafka
key.converter=org.apache.kafka.connect.storage.StringConverter
key.converter.schema.registry.url=http://<Schema_Registry_IP_Ad
dress>:8081
value.converter=org.apache.kafka.connect.storage.StringConverte
r
value.converter.schema.registry.url=http://<Schema_Registry_IP_
Address>:8081
```

Note: Update the highlighted parameters with the values according to your Confluent Kafka cluster.

6. Navigate to **/etc/kafka-connect-hdfs/quickstart-hdfs.properties** (with sudo privilege), update the HDFS URL and Topic details, save and exit the file

```
topics=blue
hdfs.url=hdfs://<CDH_Controller_IP_Address>:8020
```

Note: Update the highlighted parameters.

7. Run the HDFS Connector

```
sudo -u hdfs /bin/connect-standalone /etc/schema-registry/connect-avro-standalone.properties /etc/kafka-connect-hdfs/quickstart-hdfs.properties
```

Note: Kill any connector service, if running. Use **ps aux | grep "Connect"** to get the list of running connector service. Use **sudo kill -9 process_id** to kill the service.

Output:

```
[2019-06-06 04:59:09,286] INFO Committed hdfs://10.39.250.11:8020/topics/blue/partition=0/blue+0+0000000000+0000000002.avro for blue-0 (io.confluent.connect.hdfs.TopicPartitionWriter:782)
```

Note: The snapshot shows the path where the data is stored in HDFS.

- SSH into the HDFS cluster Controller node and look for the path

```
hdfs dfs -ls /topics/blue/partition=0
```

Output:

Found 1 items

```
-rw-r--r--      3 hdfs supergroup          108 2019-06-06 05:00
/topics/blue/partition=0/blue+0+0000000000+0000000002.avro
```

Note: Here **blue+0+0000000000+0000000002.avro** file contains the data sent from Kafka Connect node.

- Copy the file to **/tmp** directory

```
hdfs dfs -copyToLocal
/topics/blue/partition=0/blue+0+0000000000+0000000002.avro /tmp
```

- Convert the generated .avro file to .json file

```
java -jar avro-tools-1.8.2.jar tojson
/tmp/blue+0+0000000000+0000000002.avro > kdata.json
```

Output:

log4j:WARN No appenders could be found for logger (org.apache.hadoop.metrics2.lib.MutableMetricsFactory).

log4j:WARN Please initialize the log4j system properly.

log4j:WARN See <http://logging.apache.org/log4j/1.2/faq.html#noconfig> for more info.

Note: Ignore the Warning messages. Assuming, Avro tools is installed in the HDFS node.

11. View the content of the generated .json file

```
cat kdata.json
```

Output:

{"string":"kafka"}

{"string":"message"}

{"string":"here"}

12. From the Kafka Connect node, open the **/etc/hadoop/conf.empty/core-site.xml** file (with sudo privilege), add given content to it, save and exit the file

```
<configuration>
<property>
  <name>fs.s3a.access.key</name>
  <value></value>
</property>
<property>
  <name>fs.s3a.secret.key</name>
  <value></value>
</property>
<property>
  <name>fs.s3a.impl</name>
  <value>org.apache.hadoop.fs.s3a.S3AFileSystem</value>
</property>
```

```
</property>
<property>
  <name>fs.dtap.impl</name>
  <value>com.bluedata.hadoop.bdfs.Bdfs</value>
  <description>The FileSystem for BlueData dtap:
URIs.</description>
</property>
</configuration>
```

13. Open **/etc/kafka-connect-hdfs/quickstart-hdfs.properties** file (with sudo privilege), update the following details for DTAP connection, save and exit the file

```
topics=blue
hdfs.url=dtap://TenantStorage/kafka
hadoop.conf.dir=/etc/hadoop/conf.empty/
```

Note: The hdfs.url should be the path from EPIC DTAP.

14. Run the executor


```
sudo -u hdfs /bin/connect-standalone /etc/schema-
registry/connect-avro-standalone.properties /etc/kafka-connect-
hdfs/quickstart-hdfs.properties
```


Output:


```


[2019-06-06 06:14:02,415] INFO Opening record writer for: dtap://TenantStorage/message/topics//+tmp/blue/partitio
n=0/00a398ce-fcce-4915-96e8-093038ac8069_tmp.avro (io.confluent.connect.hdfs.avro.AvroRecordWriterProvider:65)
```


15. Verify in the EPIC UI


 BlueData


 Dashboard


 Clusters 4


 DataTaps 1

 FS Mounts 0






 Nodes 28

 Users 4

 Flavors



TenantStorage Browser



dtap://TenantStorage/kafka

/

kafka

tmp

logs

topics

+tmp

blue

partition=0

blue 0 0000000000 0000000002.avro