

SMOKE TEST

HDF 3.3.1

Date Prepared: July 2019

Document Information

Project Name	HDF Smoke Test Document		
Project Owner		Document Version No	1.0
Quality Review Method	By email/HP SharePoint		
Prepared By		Preparation Date	July 2019
Reviewed By	Refer to version history	Review Date	

Table of Contents

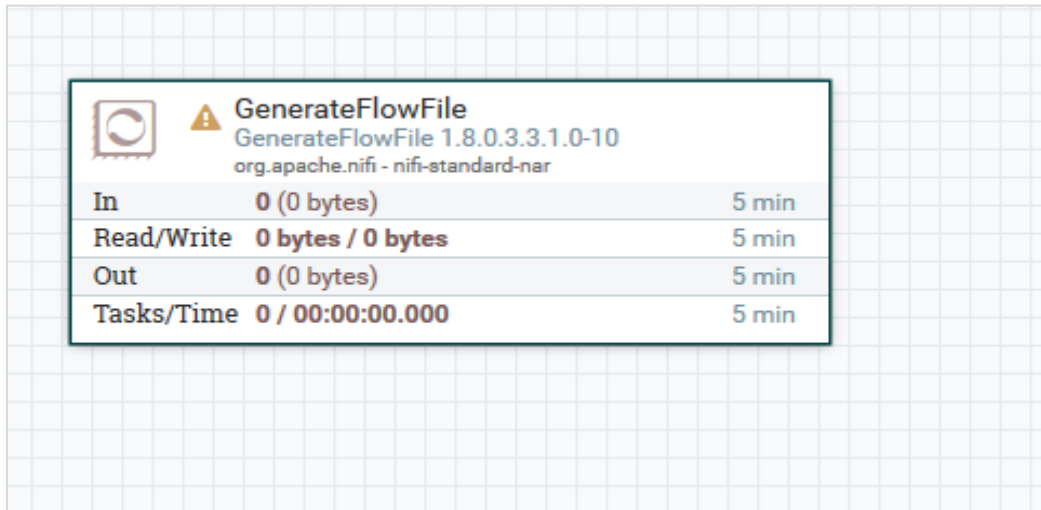
1	TESTING NIFI	5
1.1	SELECT A GENERATEFLOWFILE PROCESSOR IN NIFI UI	5
1.2	CONFIGURE GENERATEFLOWFILE PROCESSOR TO GENERATE THE DATA.....	5
1.3	SELECT A PUTFILE PROCESSOR IN NIFI UI	6
1.4	CONFIGURE PUTFILE PROCESSOR TO STORE GENERATED DATA FROM GENERATEFLOWFILE PROCESSOR.....	6
1.5	CONNECT AND RUN THE PROCESSORS	7
1.6	VERIFY THE DATA	7
2	TESTING NIFI REGISTRY.....	8
2.1	CREATE HELLOWORLD BUCKET IN NIFI REGISTRY UI	8
2.2	CREATE A PROCESS GROUP	8
2.3	CONFIGURE A PROCESS GROUP	9
2.4	VERIFY VERSION CONTROL	9
3	TESTING GRAFANA	11
3.1	SELECTION OF COMPONENT	11
4	TESTING SOLR.....	12
4.1	GO TO THE SOLR ADMIN UI IN HDF CLUSTER	12
4.2	GO TO THE SOLR CONTAINER AND CREATE A NEW COLLECTION	12
4.3	VERIFY HELLOWORLD COLLECTION CREATION IN SOLR ADMIN UI.....	13
4.4	PLACE SOME EXAMPLE JSON DATA IN HELLOWORLD COLLECTION	14
4.5	FETCH JSON DATA FROM HELLOWORLD COLLECTION.....	14
5	TESTING KAFKA	16
5.1	CREATE A NEW KAFKA TOPIC.....	16
5.2	CREATE A NEW TOPIC	16
6	TESTING LOG SEARCH	17
6.1	LOG SEARCH DASHBOARD FOR SERVICE LOGS.....	17
6.2	LOG SEARCH DASHBOARD FOR AUDIT LOGS.....	17
7	TESTING STREAMING ANALYTICS MANAGER.....	18
7.1	GO TO THE STREAMLINE UI	18
7.2	CREATE A SERVICE POOL	18
7.3	CREATE A NEW ENVIRONMENT FOR HDF CLUSTER FROM SERVICE POOL	19
7.4	CREATE A NEW APPLICATION FOR YOUR ENVIRONMENT	19
8	TESTING RANGER	21

9	TESTING STORM.....	22
9.1	DOWNLOAD STORM WORDCOUNT EXAMPLE	22
9.2	CHANGE DIRECTORY TO THE STORM-EXAMPLE	22
9.3	BUILDING JAR	22
9.4	CREATE TOPOLOGY FOR STORM WORDCOUNT EXAMPLE	23
9.5	VERIFY WORDCOUNT TOPOLOGY FROM STORM UI	24

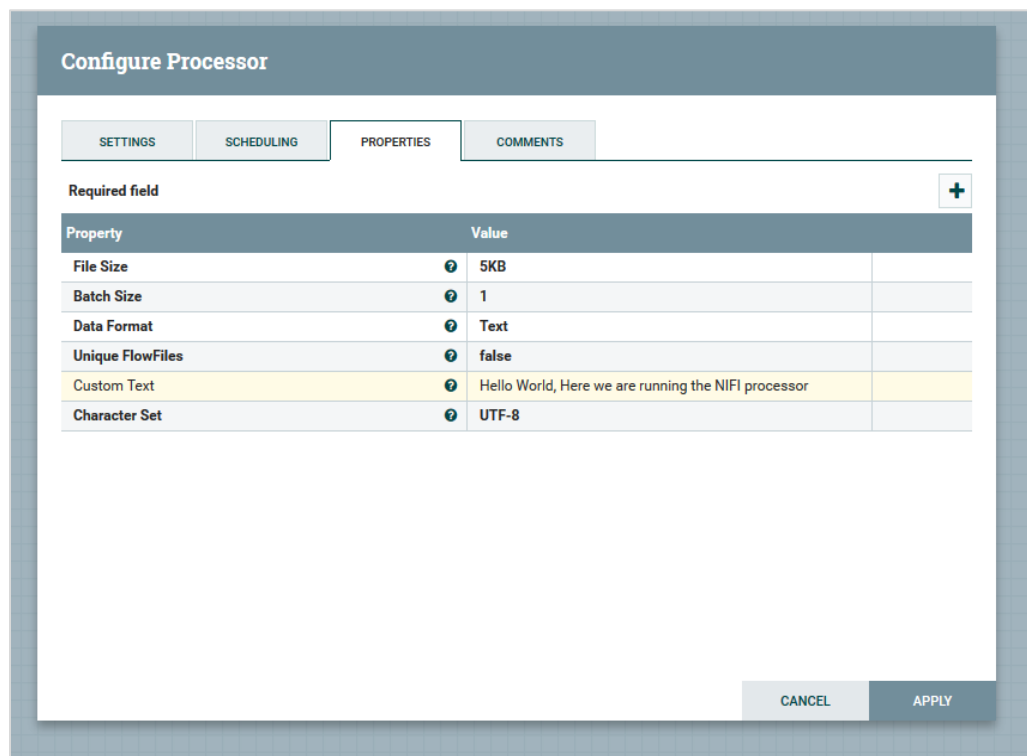
1 TESTING NIFI

We will use NIFI UI to create a HelloWorld dataflow.

1.1 Select a GenerateFlowFile processor in NIFI UI



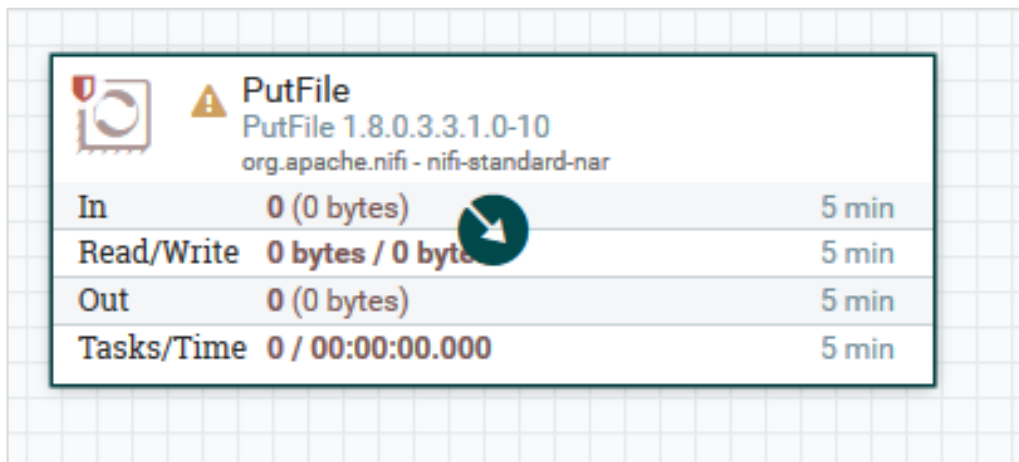
1.2 Configure GenerateFlowFile processor to generate the data



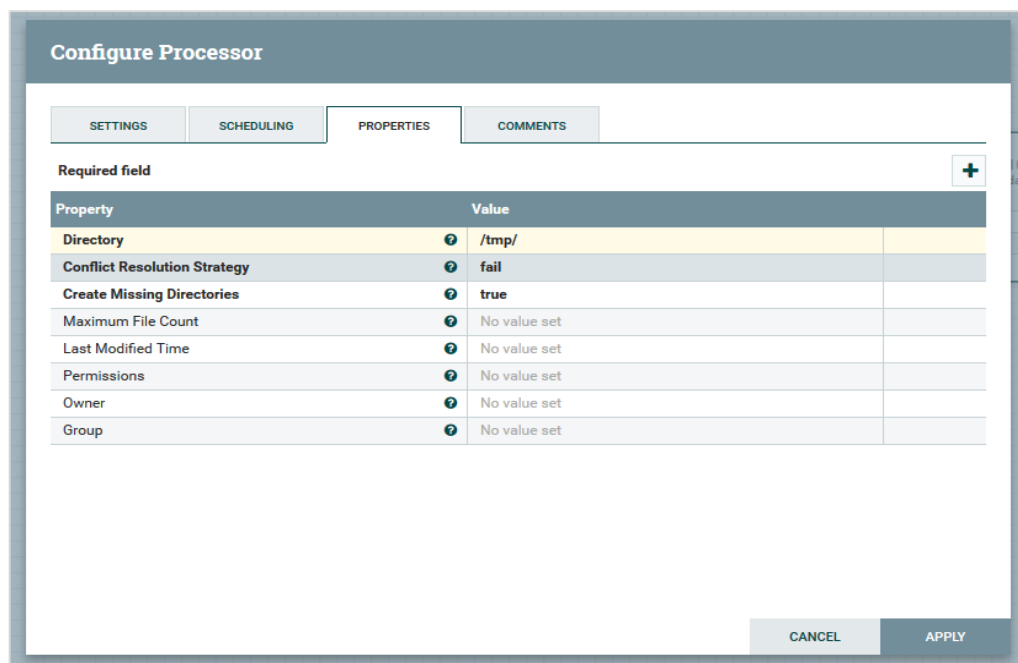
Following are the configuration you need to perform for GenerateFlowFile processor:

- Set the file size
- Enter custom text in Custom Text section

1.3 Select a PutFile processor in NIFI UI



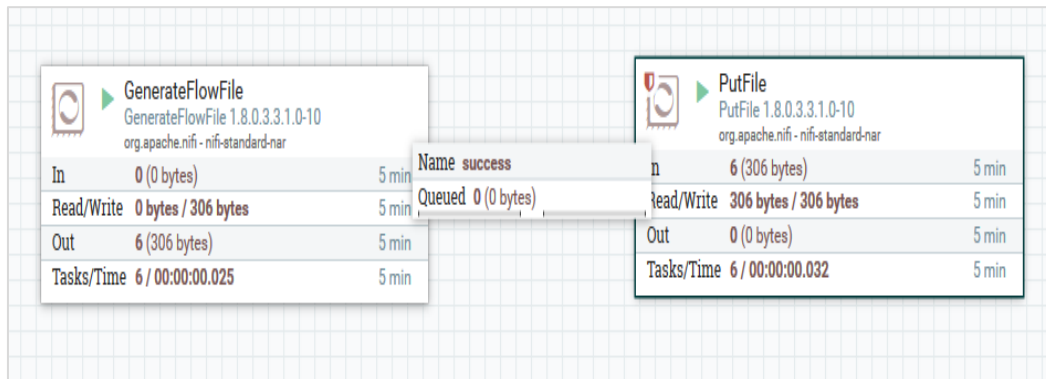
1.4 Configure PutFile processor to store generated data from GenerateFlowFile processor




You need to perform following configuration for PutFile processor:

- Enter a location where you want to store generated data from GenerateFlowFile processor.
(Ex. /tmp/)

1.5 Connect and run the processors

	GenerateFlowFile GenerateFlowFile 1.8.0.3.3.1.0-10 org.apache.nifi - nifi-standard-nar		
In	0 (0 bytes)	5 min	Name success
Read/Write	0 bytes / 306 bytes	5 min	Queued 0 (0 bytes)
Out	6 (306 bytes)	5 min	
Tasks/Time	6 / 00:00:00.025	5 min	

	PutFile PutFile 1.8.0.3.3.1.0-10 org.apache.nifi - nifi-standard-nar		
In	6 (306 bytes)	5 min	
Read/Write	306 bytes / 306 bytes	5 min	
Out	0 (0 bytes)	5 min	
Tasks/Time	6 / 00:00:00.032	5 min	

1.6 Verify the data

Go to NIFI container and check into **/tmp** directory, if the data is stored or not.

```
ls /tmp/
```

```
[bluedata@bluedata-6670 ~]$ ls /tmp/
0f415ed5-b405-422c-89a6-69bd68f17bcf  6ae13641-8aef-4eab-9e0a-bc53e70a3f1d  bds-20190616225559.log          eella6ce-974d-4276-bd1f-5b4b54495e66
1c96755c-209b-4690-b8b5-71e9186d50fc  6f5eef59-f810-4a31-a838-ced3e4f86650  bd_vagent_bundle.517.log       efc2534f-e25f-412d-a0f5-309bc920b95a
2019466d-1b12-44f9-83f2-033252867b9f  7852fa6c-d820-41f1-bb38-34158882a2d0  cfa81c94-579d-48fb-917b-abb48bfd6e04  fe448eb1-b30b-493a-af79-4cb05730d89c
23994186-4bef-4f26-9638-074eb4b9e20b  787f475d-c6a4-4f2a-ab98-e786f0cd5212  d6c6e001-57c2-41f5-aabc-ebdc894122f0  hspcrfdata_mynifiuser
3540a303-9e91-42ad-a66c-ada17c3f1673  87111caa-0238-4323-867b-9602378bbc5c  dbbab28a-5666-4f93-b992-904a7265fe63  hspcrfdata_registry
358ef550-c2cc-42fb-837e-e71b75422b2e  8c0c012e-ed1c-4eab-9027-903d5b49f2b7  de538e75-172d-4631-9684-cda97eb5d315  hspcrfdata_root
4f44eda8-5801-43ae-965d-fded68baa078  8d970764-d334-4f9d-b964-3c2f6c1e428d  e3a9c207-d4dc-44df-88a5-90f30e6ad659  jna-551163445
647c88ff-ea0f-4536-9612-2ce45f79992c  a2c32c88-76bb-4210-a9ad-f1f43d3c44b6  e9e1a5f6-915e-4438-aa4d-c25ea7985a37  snappy-1.0.5-libsnapyyjava.so
65b8d7ee-9635-4de4-9d2e-8622f8504fb1  a477f930-f8b8-4873-8cbd-a2e73a10e849  ec394c58-beda-469a-a17f-fd83a262ef70
```


Use **cat** command to check the stored data.

```
[bluedata@bluedata-6670 ~]$ cat /tmp/0f415ed5-b405-422c-89a6-69bd68f17bcf
Hello World, Here we are running the NIFI processor[bluedata@bluedata-6670 ~]$
[bluedata@bluedata-6670 ~]$
```

2 TESTING NIFI REGISTRY

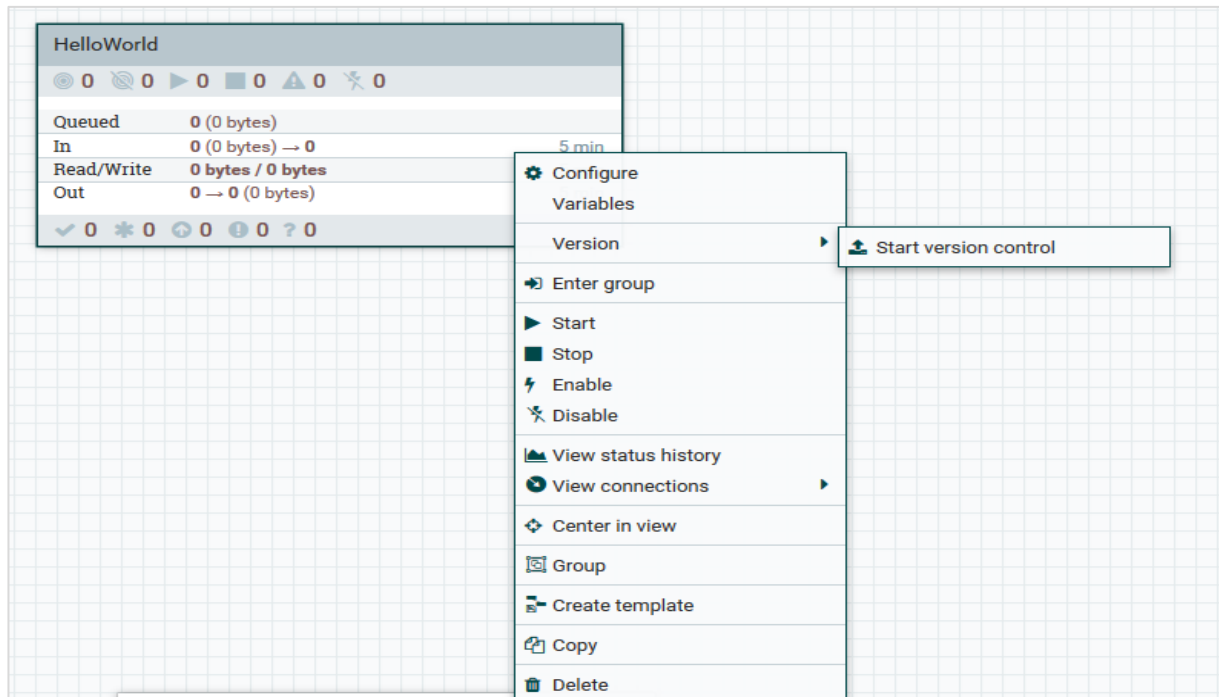
Here we will test the NIFI Registry service. Go to the NIFI Registry UI and create a bucket called HelloWorld.

2.1 Create HelloWorld bucket in NIFI Registry UI



2.2 Create a process group

Create a HelloWorld process group in NIFI UI and start version control.



2.3 Configure a process group

Configure HelloWorld process group for storing it into bucket HelloWorld in NIFI Registry.

Save Flow Version

Registry

bluedata-6700.bdlocal

Bucket

HelloWorld

Flow Name

HelloWorld

1

Flow Description

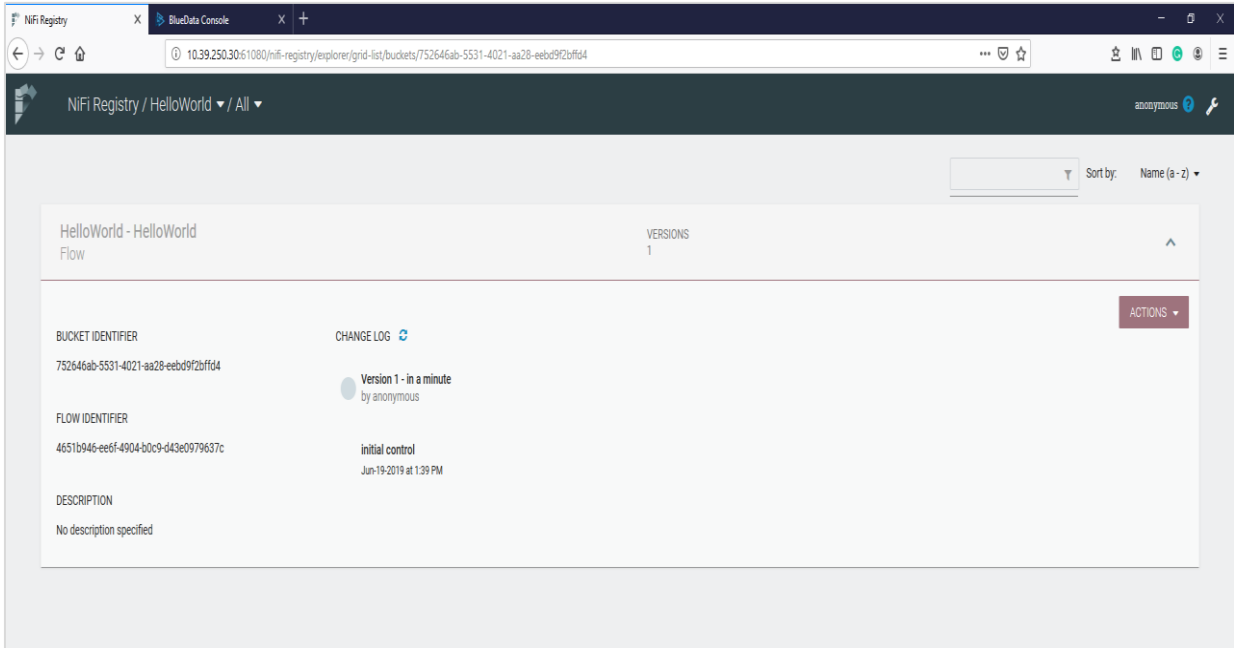
Version Comments

CANCEL


SAVE

2.4 Verify version control

Go to the NIFI Registry UI. Select HelloWorld bucket and check if version control started for process group HelloWorld.



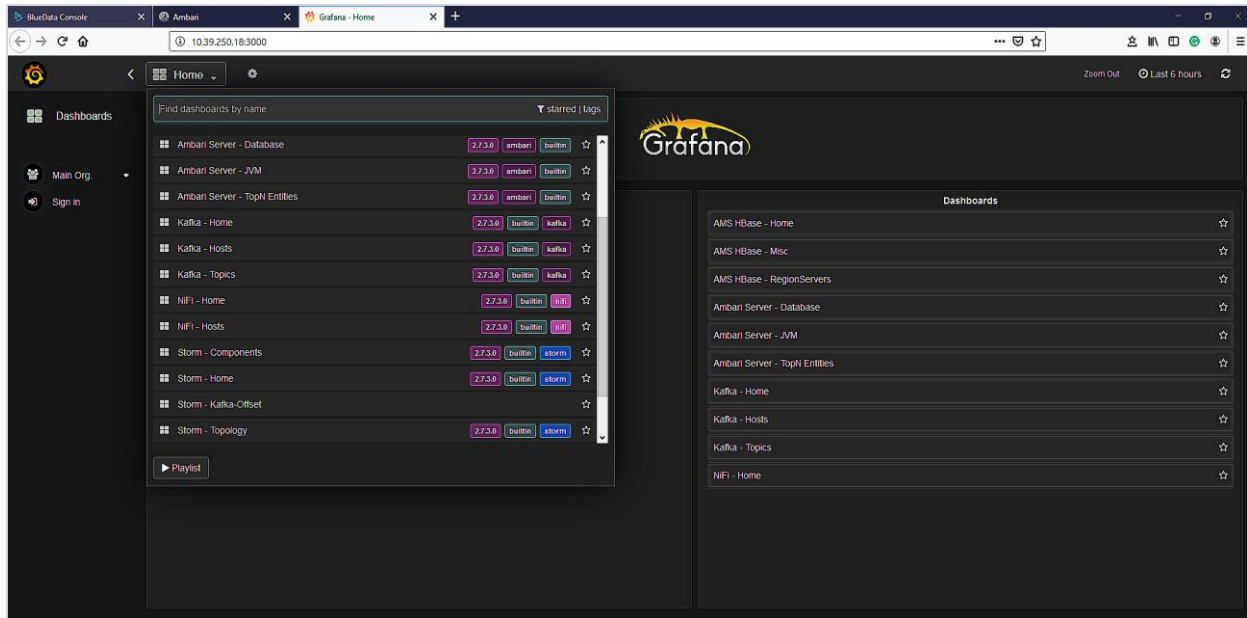
The screenshot displays the NiFi Registry web application. The browser's address bar shows the URL: `10.39.250.30:1080/nifi-registry/explorer/grid-list/buckets/752646ab-5531-4021-aa28-eebd9f2bffd4`. The page title is "NiFi Registry / HelloWorld / All". The main content area shows details for the "HelloWorld - HelloWorld Flow".

HelloWorld - HelloWorld Flow		VERSIONS 1
BUCKET IDENTIFIER	CHANGE LOG	 Version 1 - in a minute by anonymous
752646ab-5531-4021-aa28-eebd9f2bffd4		
FLOW IDENTIFIER		
4651b946-ee6f-4904-b0c9-d43e0979637c	initial control	
	Jun-19-2019 at 1:39 PM	
DESCRIPTION		
No description specified		

An "ACTIONS" button is visible in the top right corner of the details panel.

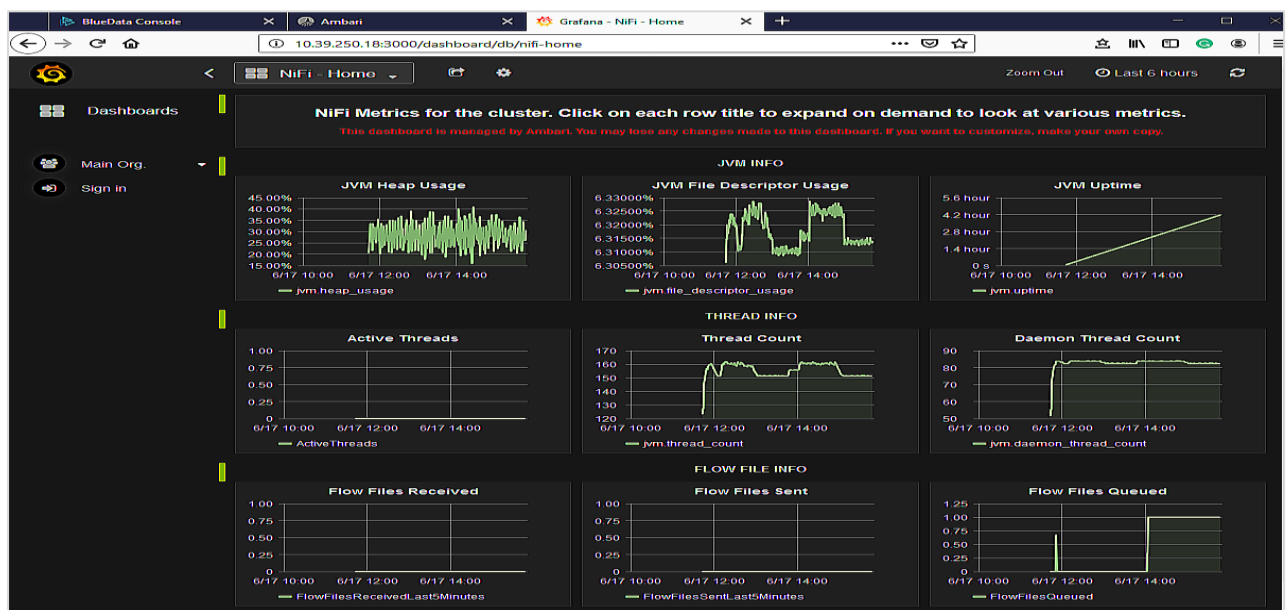
3 TESTING GRAFANA

Go to the Grafana UI and click on the search icon. Here you can see the components which is deployed into the HDF cluster.



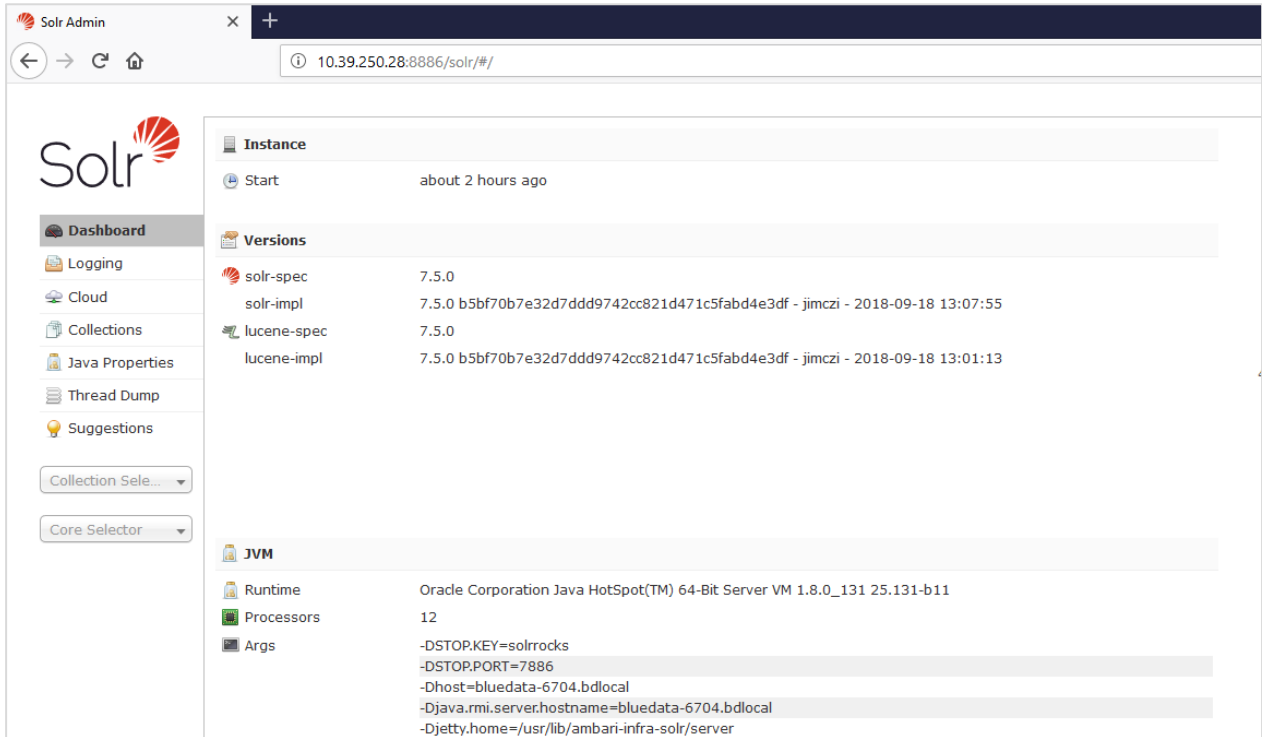
3.1 Selection of component

NIFI – Home component is selected in Grafana UI, and you can see the resources usages by the NIFI service.



4 TESTING SOLR

4.1 Go to the Solr Admin UI in HDF cluster



The screenshot shows the Solr Admin web interface. The left sidebar contains navigation links: Dashboard, Logging, Cloud, Collections, Java Properties, Thread Dump, and Suggestions. Below these are 'Collection Selector' and 'Core Selector' dropdowns. The main content area displays the following information:

- Instance:** Start button, status 'about 2 hours ago'.
- Versions:** A table listing components and their versions.

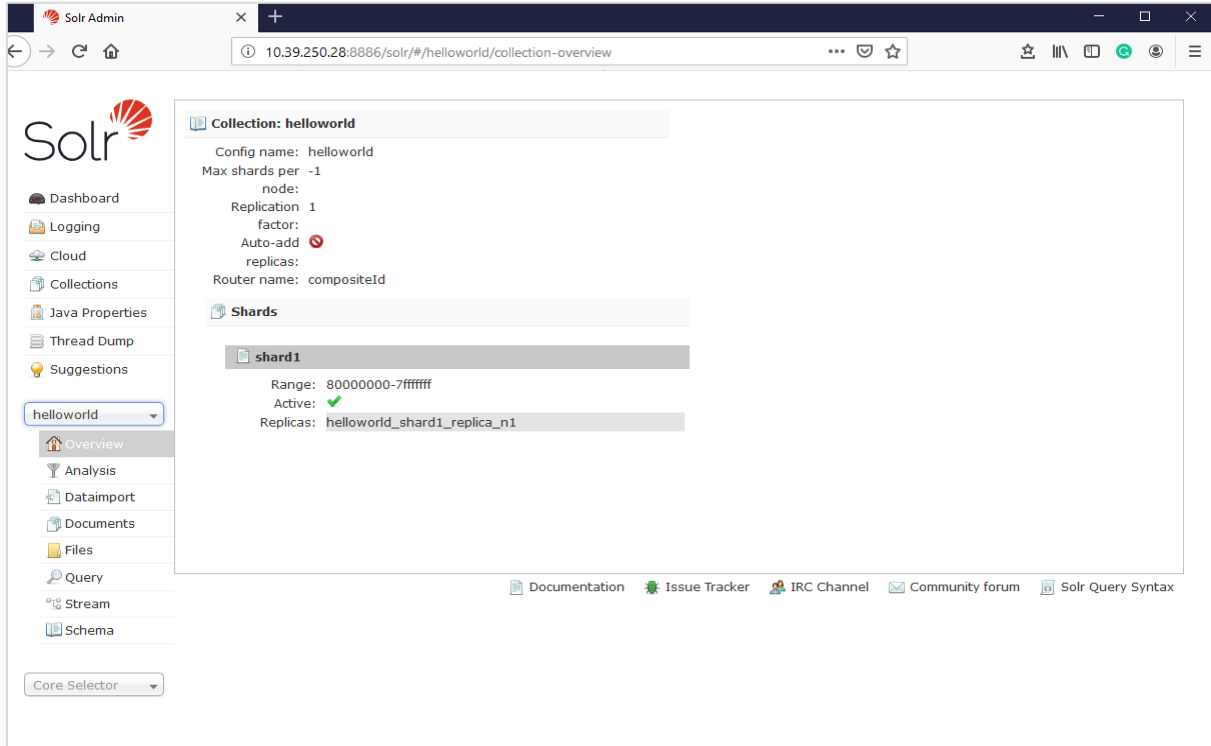
Component	Version
solr-spec	7.5.0
solr-impl	7.5.0 b5bf70b7e32d7ddd9742cc821d471c5fabd4e3df - jimcz - 2018-09-18 13:07:55
lucene-spec	7.5.0
lucene-impl	7.5.0 b5bf70b7e32d7ddd9742cc821d471c5fabd4e3df - jimcz - 2018-09-18 13:01:13
- JVM:** Runtime (Oracle Corporation Java HotSpot(TM) 64-Bit Server VM 1.8.0_131 25.131-b11), Processors (12), and Args (a list of JVM arguments including -DSTOP.KEY, -DSTOP.PORT, -Dhost, -Djava.rmi.server.hostname, and -Djetty.home).

4.2 Go to the Solr container and create a new collection

```
./bin/solr create -c helloworld
```

```
[bluedata@bluedata-6704 ambari-infra-solr]$  
[bluedata@bluedata-6704 ambari-infra-solr]$ ./bin/solr create -c helloworld  
WARNING: Using default configset with data driven schema functionality. NOT RECOMMENDED for production use.  
To turn off: bin/solr config -c helloworld -p 8886 -action set-user-property -property update.autoCreateFields -value false  
INFO - 2019-06-17 23:20:42.271; org.apache.solr.util.configuration.SSLCredentialProviderFactory; Processing SSL Credential Provider chain: env;  
sysprop  
Created collection 'helloworld' with 1 shard(s), 1 replica(s) with config-set 'helloworld'  
[bluedata@bluedata-6704 ambari-infra-solr]$
```

4.3 Verify HelloWorld collection creation in Solr Admin UI



The screenshot displays the Solr Admin web interface in a browser window. The address bar shows the URL `10.39.250.28:8886/solr/#/helloworld/collection-overview`. The left sidebar contains a navigation menu with options: Dashboard, Logging, Cloud, Collections, Java Properties, Thread Dump, Suggestions, and a dropdown menu currently showing 'helloworld'. Below this menu are links for Overview, Analysis, Dataimport, Documents, Files, Query, Stream, and Schema. At the bottom of the sidebar is a 'Core Selector' dropdown. The main content area is titled 'Collection: helloworld' and displays the following configuration details:

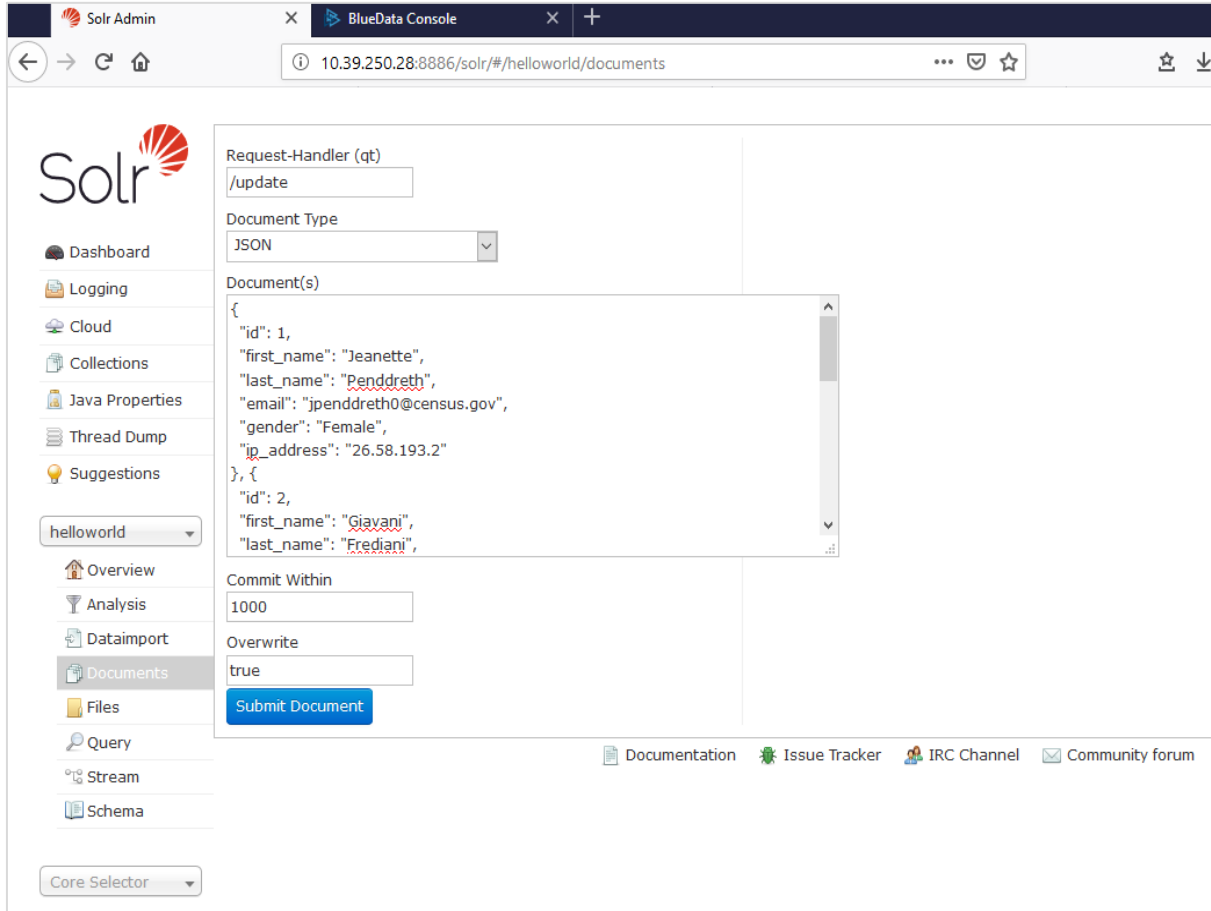
- Config name: helloworld
- Max shards per: -1
- node:
- Replication: 1
- factor:
- Auto-add: (disabled with a red 'x' icon)
- replicas:
- Router name: compositeId

Below the configuration, there is a section for 'Shards' containing a single entry, 'shard1', with the following details:

- Range: 80000000-7fffffff
- Active: (indicated by a green checkmark)
- Replicas: helloworld_shard1_replica_n1

At the bottom of the main content area, there are links to Documentation, Issue Tracker, IRC Channel, Community forum, and Solr Query Syntax.

4.4 Place some example JSON data in HelloWorld collection



Solr Admin

BlueData Console

10.39.250.28:8886/solr/#/helloworld/documents

Solr

Dashboard

Logging

Cloud

Collections

Java Properties

Thread Dump

Suggestions

helloworld

Overview

Analysis

Dataimport

Documents

Files

Query

Stream

Schema

Core Selector

Request-Handler (qt)

/update

Document Type

JSON

Document(s)

```
{
  "id": 1,
  "first_name": "Jeanette",
  "last_name": "Penddreth",
  "email": "jpenddreth0@census.gov",
  "gender": "Female",
  "ip_address": "26.58.193.2"
}, {
  "id": 2,
  "first_name": "Giovani",
  "last_name": "Frediani",
}
```

Commit Within

1000

Overwrite

true

Submit Document

Documentation

Issue Tracker


IRC Channel

Community forum

Perform below steps to store JSON data into HelloWorld collection:

- Select document type as JSON
- Place some example JSON data in document(s) section in HelloWorld collection and click on **Submit Document**

4.5 Fetch JSON data from HelloWorld collection



The screenshot shows the Solr Admin interface with the following details:

- Request-Handler (qt):** /select
- q:** *
- wt:** json
- Raw Query Parameters:** key1=val1&key2=val2
- Execute Query** button
- JSON Response:**

```
{
  "responseHeader": {
    "zkConnected": true,
    "status": 0,
    "QTime": 0,
    "params": {
      "q": "*",
      "wt": "json",
      "_: "1560840019269"
    }
  },
  "response": {
    "numFound": 5,
    "start": 0,
    "docs": [
      {
        "id": "9fc18751-416a-417d-b3e5-f1c05740bb35",
        "_version_": 1636659280990961664,
        {
          "id": "1",
          "first_name": ["Jeanette"],
          "last_name": ["Penddreth"],
          "email": ["jpenddreth0@census.gov"],
          "gender": ["Female"],
          "ip_address": ["26.58.193.2"],
          "_version_": 1636659458567307264,
          {
            "id": "2",
            "first_name": ["Giavani"],
            "last_name": ["Frediani"],
            "email": ["gfrediani1@senate.gov"],
            "gender": ["Male"],
            "ip_address": ["229.179.4.212"],
            "_version_": 1636659458581987328,
            {
              "id": "3",
              "first_name": ["Noell"],
              "last_name": ["Bea"],
              "email": ["nbea2@imageshack.us"],
              "gender": ["Female"],
              "ip_address": ["180.66.162.255"],
              "_version_": 1636659458583035904,
              {
                "id": "4",
                "first_name": ["Willard"],
                "last_name": ["Valek"],
                "email": ["wvalek3@vk.com"],
                "gender": ["Male"],
```

Perform below step to fetch JSON data from HelloWorld collection:

- Select wt as JSON and click on execute query

5 TESTING KAFKA

Go inside the Kafka container and follow the below steps.

5.1 Create a new Kafka topic

Create new Kafka topic inside the Kafka container.

5.2 Create a new topic

Execute the below command to create a new topic "SDS"

```
bin/Kafka-topics --create --zookeeper <ip address of zookeeper>:2181 --replication-factor 1 --  
partitions 1 --topic SDS
```

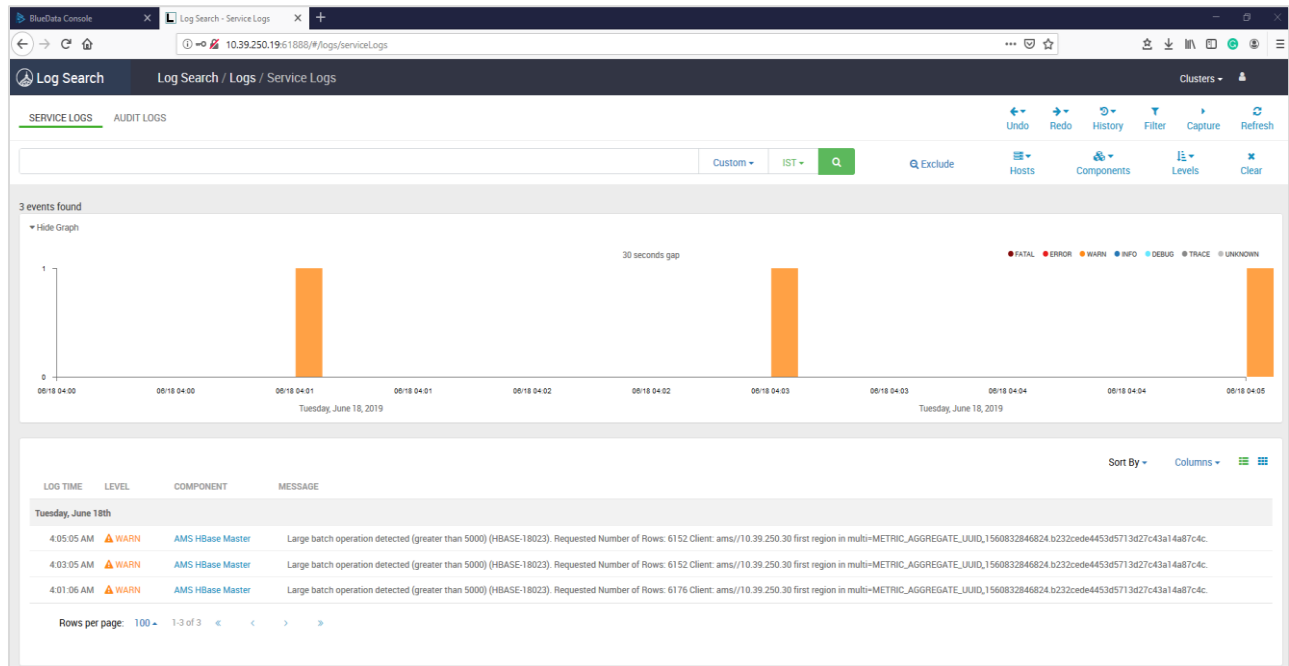
```
[bluedata@bluedata-6711 kafka]$ ls bin/  
connect-distributed.sh      kafka-configs.sh           kafka-delegation-tokens.sh  kafka-preferred-replica-elec  
connect-standalone.sh      kafka-console-consumer.sh  kafka-delete-records.sh     kafka-producer-perf-test.sh  
kafka                      kafka-console-producer.sh  kafka-dump-log.sh           kafka-reassign-partitions.sh  
kafka-acls.sh              kafka-consumer-groups.sh   kafka-log-dirs.sh           kafka-replica-verification.s  
kafka-broker-api-versions.sh kafka-consumer-perf-test.sh kafka-mirror-maker.sh       kafka-run-class.sh  
[bluedata@bluedata-6711 kafka]$  
[bluedata@bluedata-6711 kafka]$ ./bin/kafka-topics.sh --create --zookeeper 10.39.250.17,10.39.250.30,10.39.250.28:2  
181 --replication-factor 1 --partitions 1 --topic SDS  
Created topic "SDS".  
[bluedata@bluedata-6711 kafka]$
```


6 TESTING LOG SEARCH

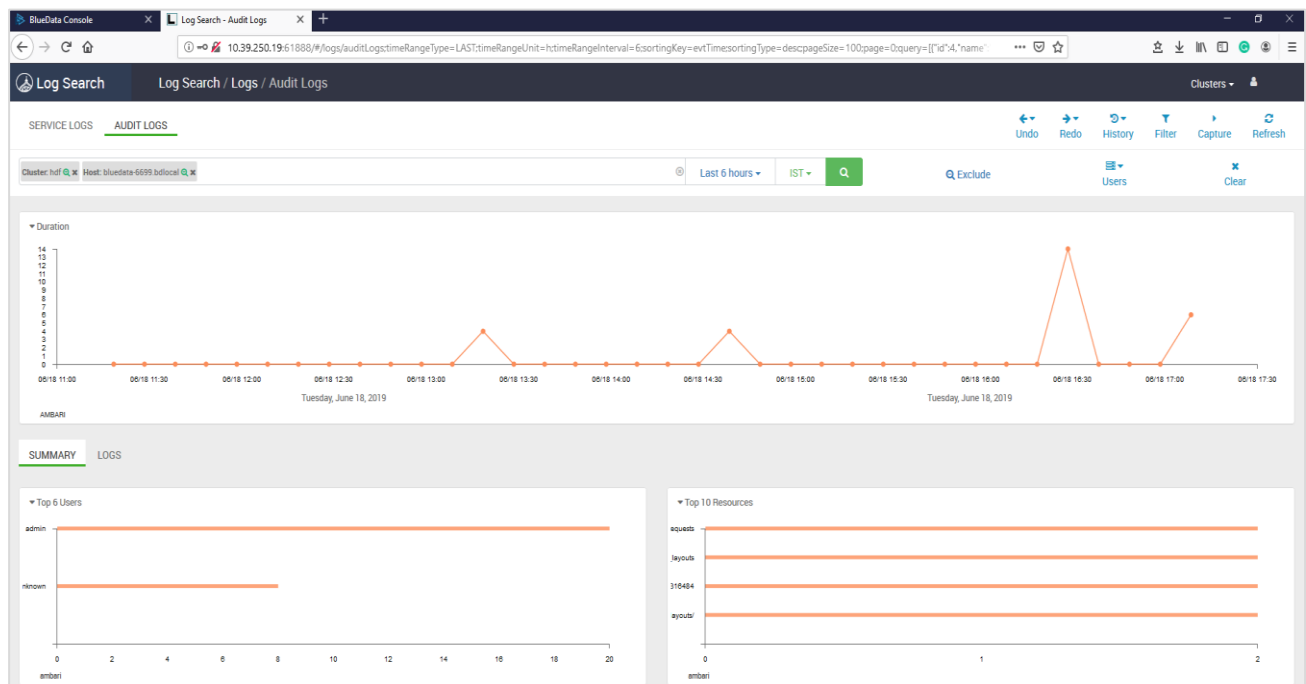
Go to the Log search UI in HDF cluster, enter username/password – admin/admin.

In Log search you can track service and audit logs for HDF component.

6.1 Log search dashboard for service logs

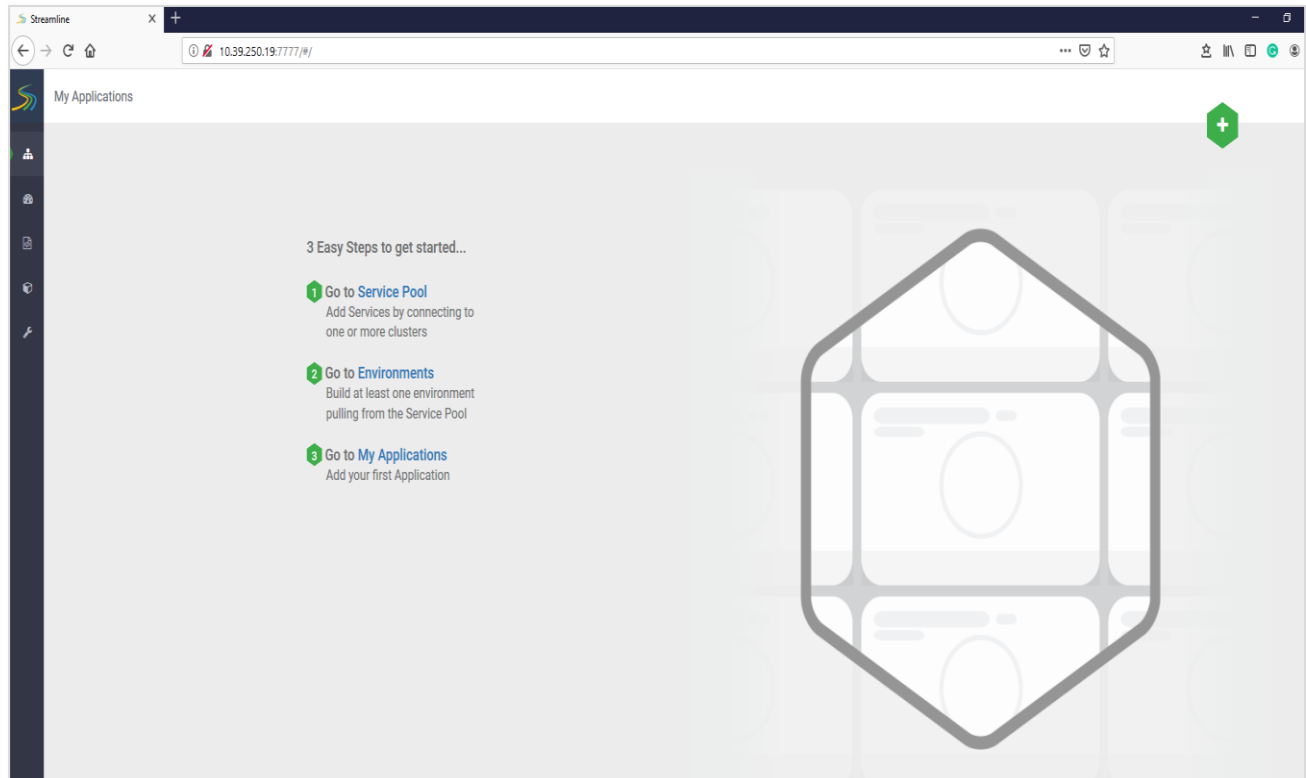


6.2 Log search dashboard for audit logs

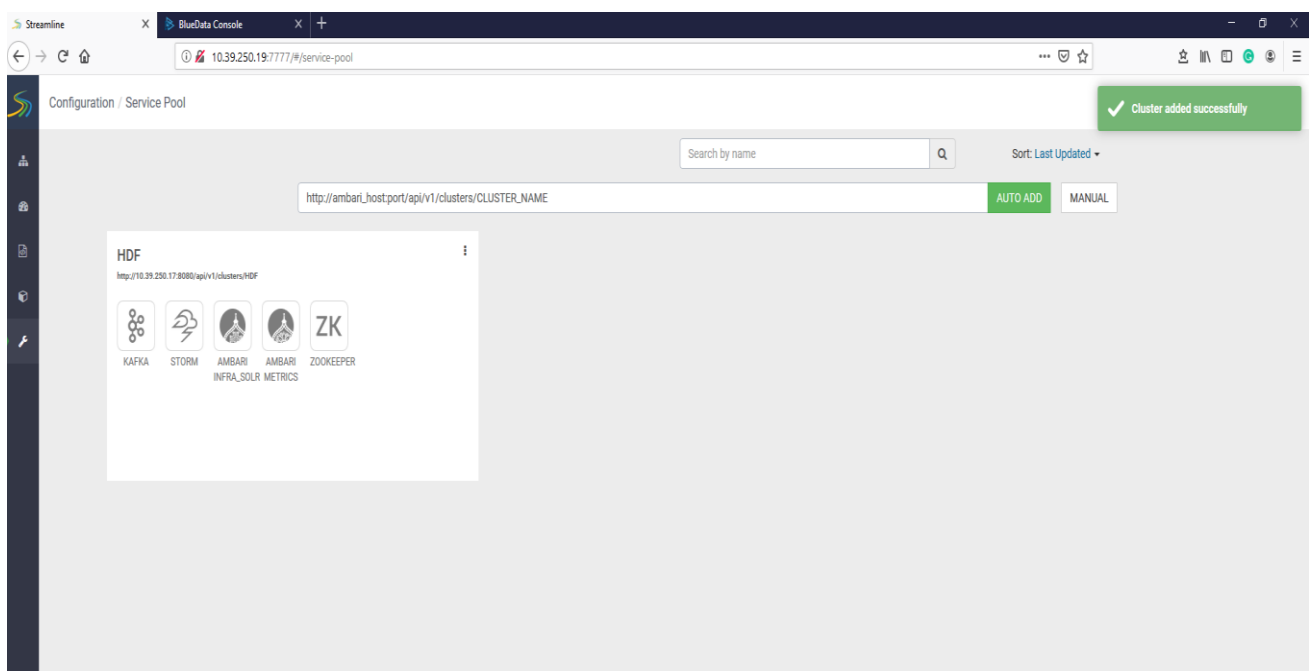


7 TESTING STREAMING ANALYTICS MANAGER

7.1 Go to the Streamline UI



7.2 Create a service pool




Perform below steps to create a service pool for HDF Cluster:

Enter url for HDF cluster into url bar and click on AUTO ADD button.

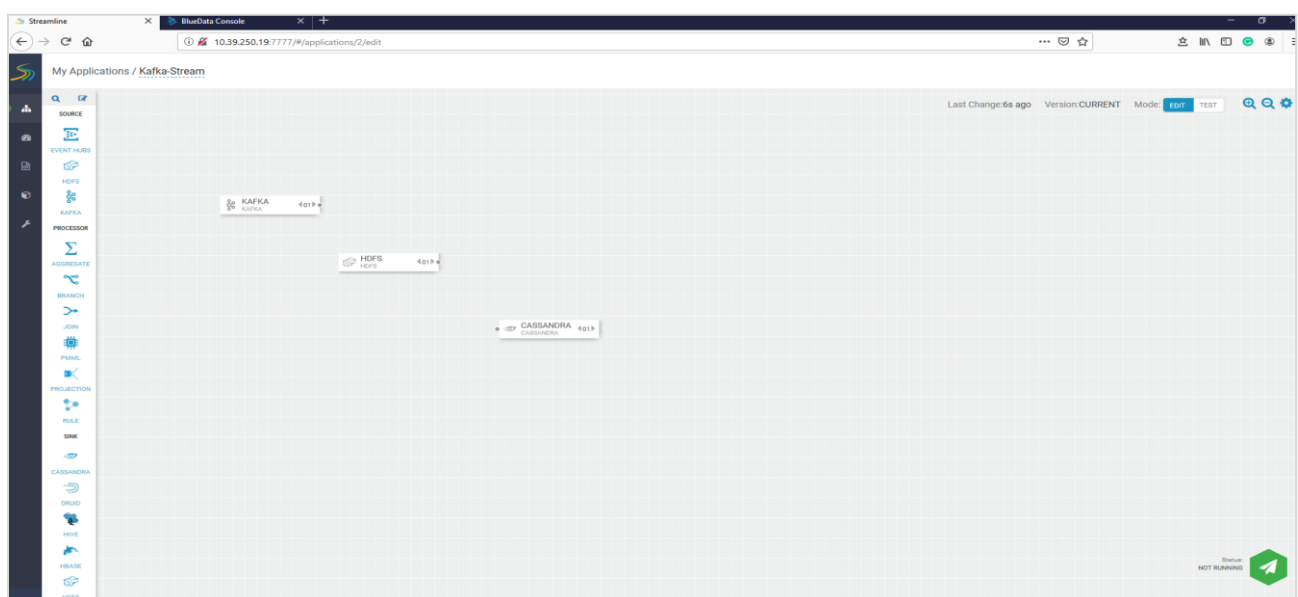
(Ex.: <http://<ip address of ambari server>:8080/api/v1/clusters/HDF>)

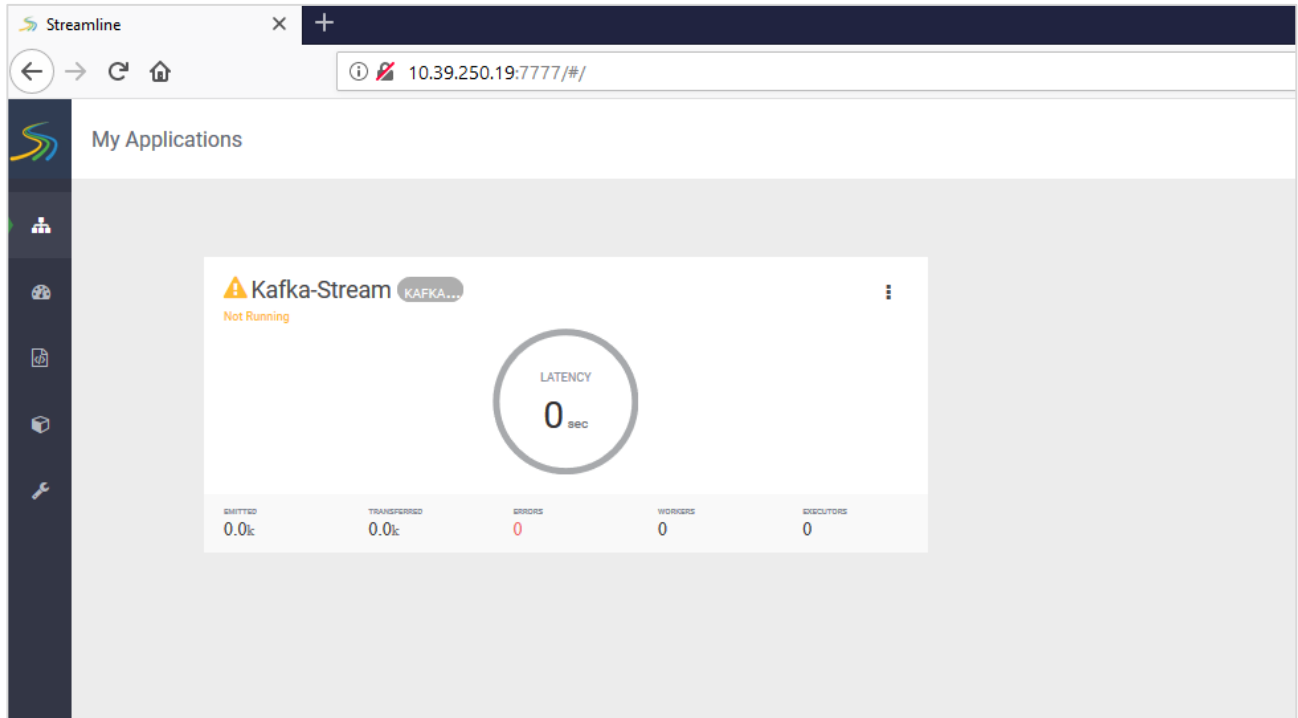
7.3 Create a new environment for HDF cluster from service pool



7.4 Create a new application for your environment

Here you can create and design your application. After running your application you can verify your application under application section.

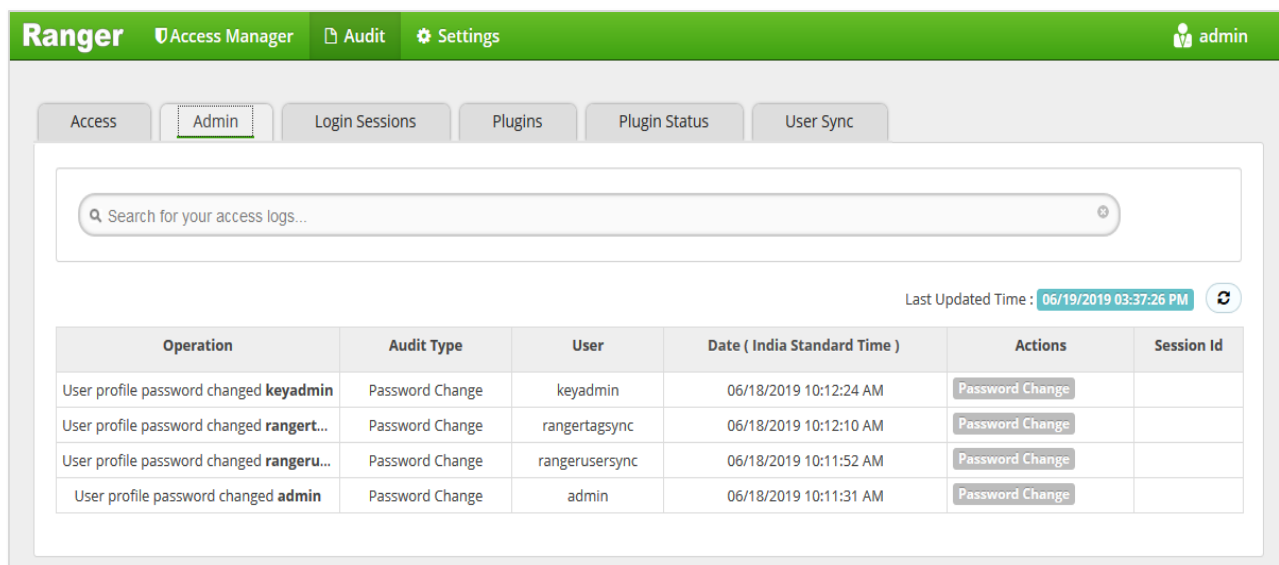
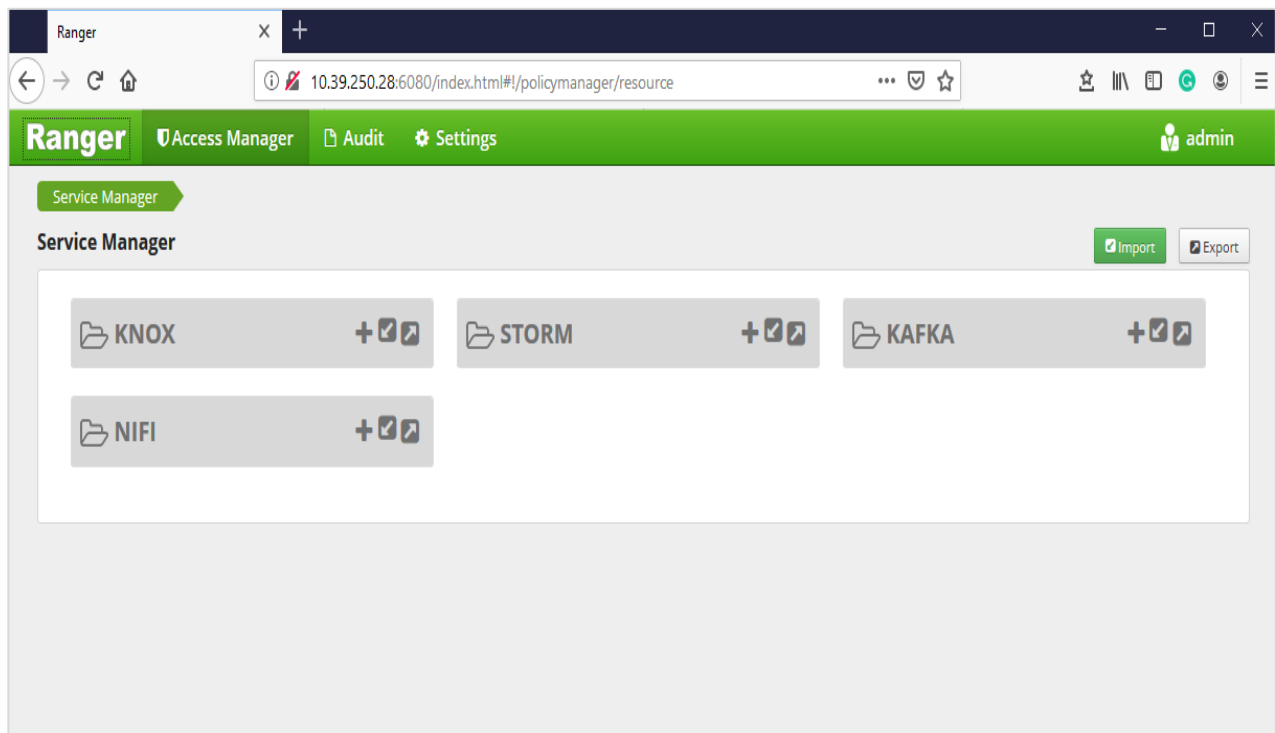




8 TESTING RANGER

Go to the Ranger UI in HDF cluster, enter username and password in login screen.

Service manager is available for Knox, Storm, Kafka and NIFI services. With the help of Ranger, you can create resource and tag based policies for different-different services.



9 TESTING STORM

Here we will create a WordCount topology for Storm. We will download a WordCount example and create a WordCount topology.

9.1 Download Storm WordCount example

Execute the below command to download Storm WordCount example

```
git clone https://github.com/ADMICloud/examples.git
```

```
[bluedata@bluedata-6896 ~]$  
[bluedata@bluedata-6896 ~]$ git clone https://github.com/ADMICloud/examples.git  
Cloning into 'examples'...  
remote: Enumerating objects: 152, done.  
remote: Total 152 (delta 0), reused 0 (delta 0), pack-reused 152  
Receiving objects: 100% (152/152), 34.88 KiB | 0 bytes/s, done.  
Resolving deltas: 100% (34/34), done.  
[bluedata@bluedata-6896 ~]$  
[bluedata@bluedata-6896 ~]$
```

9.2 Change directory to the storm-example

```
cd examples/storm-example/
```

```
[bluedata@bluedata-6896 ~]$ ls  
examples  vagent.bin  
[bluedata@bluedata-6896 ~]$ cd examples/storm-example/  
[bluedata@bluedata-6896 storm-example]$  
[bluedata@bluedata-6896 storm-example]$  
[bluedata@bluedata-6896 storm-example]$
```

9.3 Building jar

Execute the below command to build jar for Storm WordCount example.

```
mvn clean install
```

```
[bluedata@bluedata-6896 storm-example]$  
[bluedata@bluedata-6896 storm-example]$ mvn clean install  
[INFO] Scanning for projects...  
[INFO]  
[INFO] -----  
[INFO] Building storm-example 1.0  
[INFO] -----  
[INFO]  
[INFO] --- maven-clean-plugin:2.4.1:clean (default-clean) @ storm-example ---  
[INFO]
```

Note: After executing this command a new directory target will be created where you can find Storm

WordCount example jar file. Use ls command to verify.

```
[bluedata@bluedata-6896 storm-example]$ ls
pom.xml  src  target
[bluedata@bluedata-6896 storm-example]$
[bluedata@bluedata-6896 storm-example]$
[bluedata@bluedata-6896 storm-example]$ ls target/
archive-tmp  classes  generated-sources  maven-archiver  maven-status  storm-example-1.0.jar  storm-example-1.0-jar-with-dependencies.jar  surefire
[bluedata@bluedata-6896 storm-example]$
```


9.4 Create topology for Storm WordCount example

Execute the below command to create topology for Storm WordCount example.

```
./bin/storm jar /home/bluedata/examples/storm-example/target/storm-
example-1.0-jar-with-dependencies.jar
admicloud.storm.WordCount.WordCountTopology WordCount
```

```
[bluedata@bluedata-6896 storm]$ ./bin/storm jar /home/bluedata/examples/storm-example/target/storm-example-1.0-jar-with-dependencies.jar admicloud.s
tom.wordcount.WordCountTopology WordCount
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/hdf/3.3.1.0-10/storm/lib/log4j-slf4j-impl-2.8.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/bluedata/examples/storm-example/target/storm-example-1.0-jar-with-dependencies.jar!/org/slf4j/impl/StaticLog
gerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Running: /usr/java/default/bin/java -Ddaemon.name= -Dstorm.options= -Dstorm.home=/usr/hdf/3.3.1.0-10/storm -Dstorm.log.dir=/var/log/storm -Djava.lib
rary.path=/usr/local/lib:/opt/local/lib:/usr/lib -Dstorm.conf.file= -cp /usr/hdf/3.3.1.0-10/storm/*:/usr/hdf/3.3.1.0-10/storm/lib/*:/usr/hdf/3.3.1.0
-10/storm/extlib/*:/home/bluedata/examples/storm-example/target/storm-example-1.0-jar-with-dependencies.jar:/usr/hdf/current/storm-supervisor/conf:/
usr/hdf/3.3.1.0-10/storm/bin -Dstorm.jar=/home/bluedata/examples/storm-example/target/storm-example-1.0-jar-with-dependencies.jar -Dstorm.dependency
.jars= -Dstorm.dependency.artifacts={} admicloud.storm.wordcount.WordCountTopology WordCount
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/hdf/3.3.1.0-10/storm/lib/log4j-slf4j-impl-2.8.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/bluedata/examples/storm-example/target/storm-example-1.0-jar-with-dependencies.jar!/org/slf4j/impl/StaticLog
gerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
741 [main] WARN o.a.s.u.Utils - STORM-VERSION new 1.2.1.3.3.1.0-10 old null
777 [main] INFO o.a.s.StormSubmitter - Generated ZooKeeper secret payload for MD5-digest: -6269264375147352907:-8793272926567057743
894 [main] INFO o.a.s.u.NimbusClient - Found leader nimbus : bluedata-6896.bdlocal:6627
918 [main] INFO o.a.s.s.a.AuthUtils - Got AutoCreds []
921 [main] INFO o.a.s.u.NimbusClient - Found leader nimbus : bluedata-6896.bdlocal:6627
941 [main] INFO o.a.s.StormSubmitter - Uploading dependencies - jars...
942 [main] INFO o.a.s.StormSubmitter - Uploading dependencies - artifacts...
942 [main] INFO o.a.s.StormSubmitter - Dependency Blob keys - jars : [] / artifacts : []
947 [main] INFO o.a.s.StormSubmitter - Uploading topology jar /home/bluedata/examples/storm-example/target/storm-example-1.0-jar-with-dependencies
.jar to assigned location: /hadoop/storm/nimbus/inbox/stormjar-a9a53b08-926d-4bb0-a096-08654df6cbad.jar
Start uploading file '/home/bluedata/examples/storm-example/target/storm-example-1.0-jar-with-dependencies.jar' to '/hadoop/storm/nimbus/inbox/storm
jar-a9a53b08-926d-4bb0-a096-08654df6cbad.jar' (165697 bytes)
[=====] 165697 / 165697
File '/home/bluedata/examples/storm-example/target/storm-example-1.0-jar-with-dependencies.jar' uploaded to '/hadoop/storm/nimbus/inbox/stormjar-a9a
53b08-926d-4bb0-a096-08654df6cbad.jar' (165697 bytes)
967 [main] INFO o.a.s.StormSubmitter - Successfully uploaded topology jar to assigned location: /hadoop/storm/nimbus/inbox/stormjar-a9a53b08-926d-
4bb0-a096-08654df6cbad.jar
967 [main] INFO o.a.s.StormSubmitter - Submitting topology WordCount in distributed mode with conf {"storm.zookeeper.topology.auth.scheme":"digest
","storm.zookeeper.topology.auth.payload":"-6269264375147352907:-8793272926567057743","topology.workers":3,"topology.debug":true}
967 [main] WARN o.a.s.u.Utils - STORM-VERSION new 1.2.1.3.3.1.0-10 old 1.2.1.3.3.1.0-10
1201 [main] INFO o.a.s.StormSubmitter - Finished submitting topology: WordCount
```

9.5 Verify WordCount topology from Storm UI



The screenshot shows the Storm UI interface with the following sections:

Storm UI

Cluster Summary

Version	Supervisors	Used slots	Free slots	Total slots	Executors	Tasks
1.2.1.3.3.1.0-10	1	2	0	2	29	29

Nimbus Summary

Search:

Host	Port	Status	Version	Up Time
bluedata-6896.bdlocal	6627	Leader	1.2.1.3.3.1.0-10	22h 21m 54s

Showing 1 to 1 of 1 entries

Topology Summary

Search:

Name	Owner	Status	Uptime	Num workers	Num executors	Num tasks	Replication count	Assigned Mem (MB)	Scheduler Info
WordCount	storm	ACTIVE	47s	2	29	29	1	1664	

Showing 1 to 1 of 1 entries

Note: Under Topology Summary WordCount topology is created.